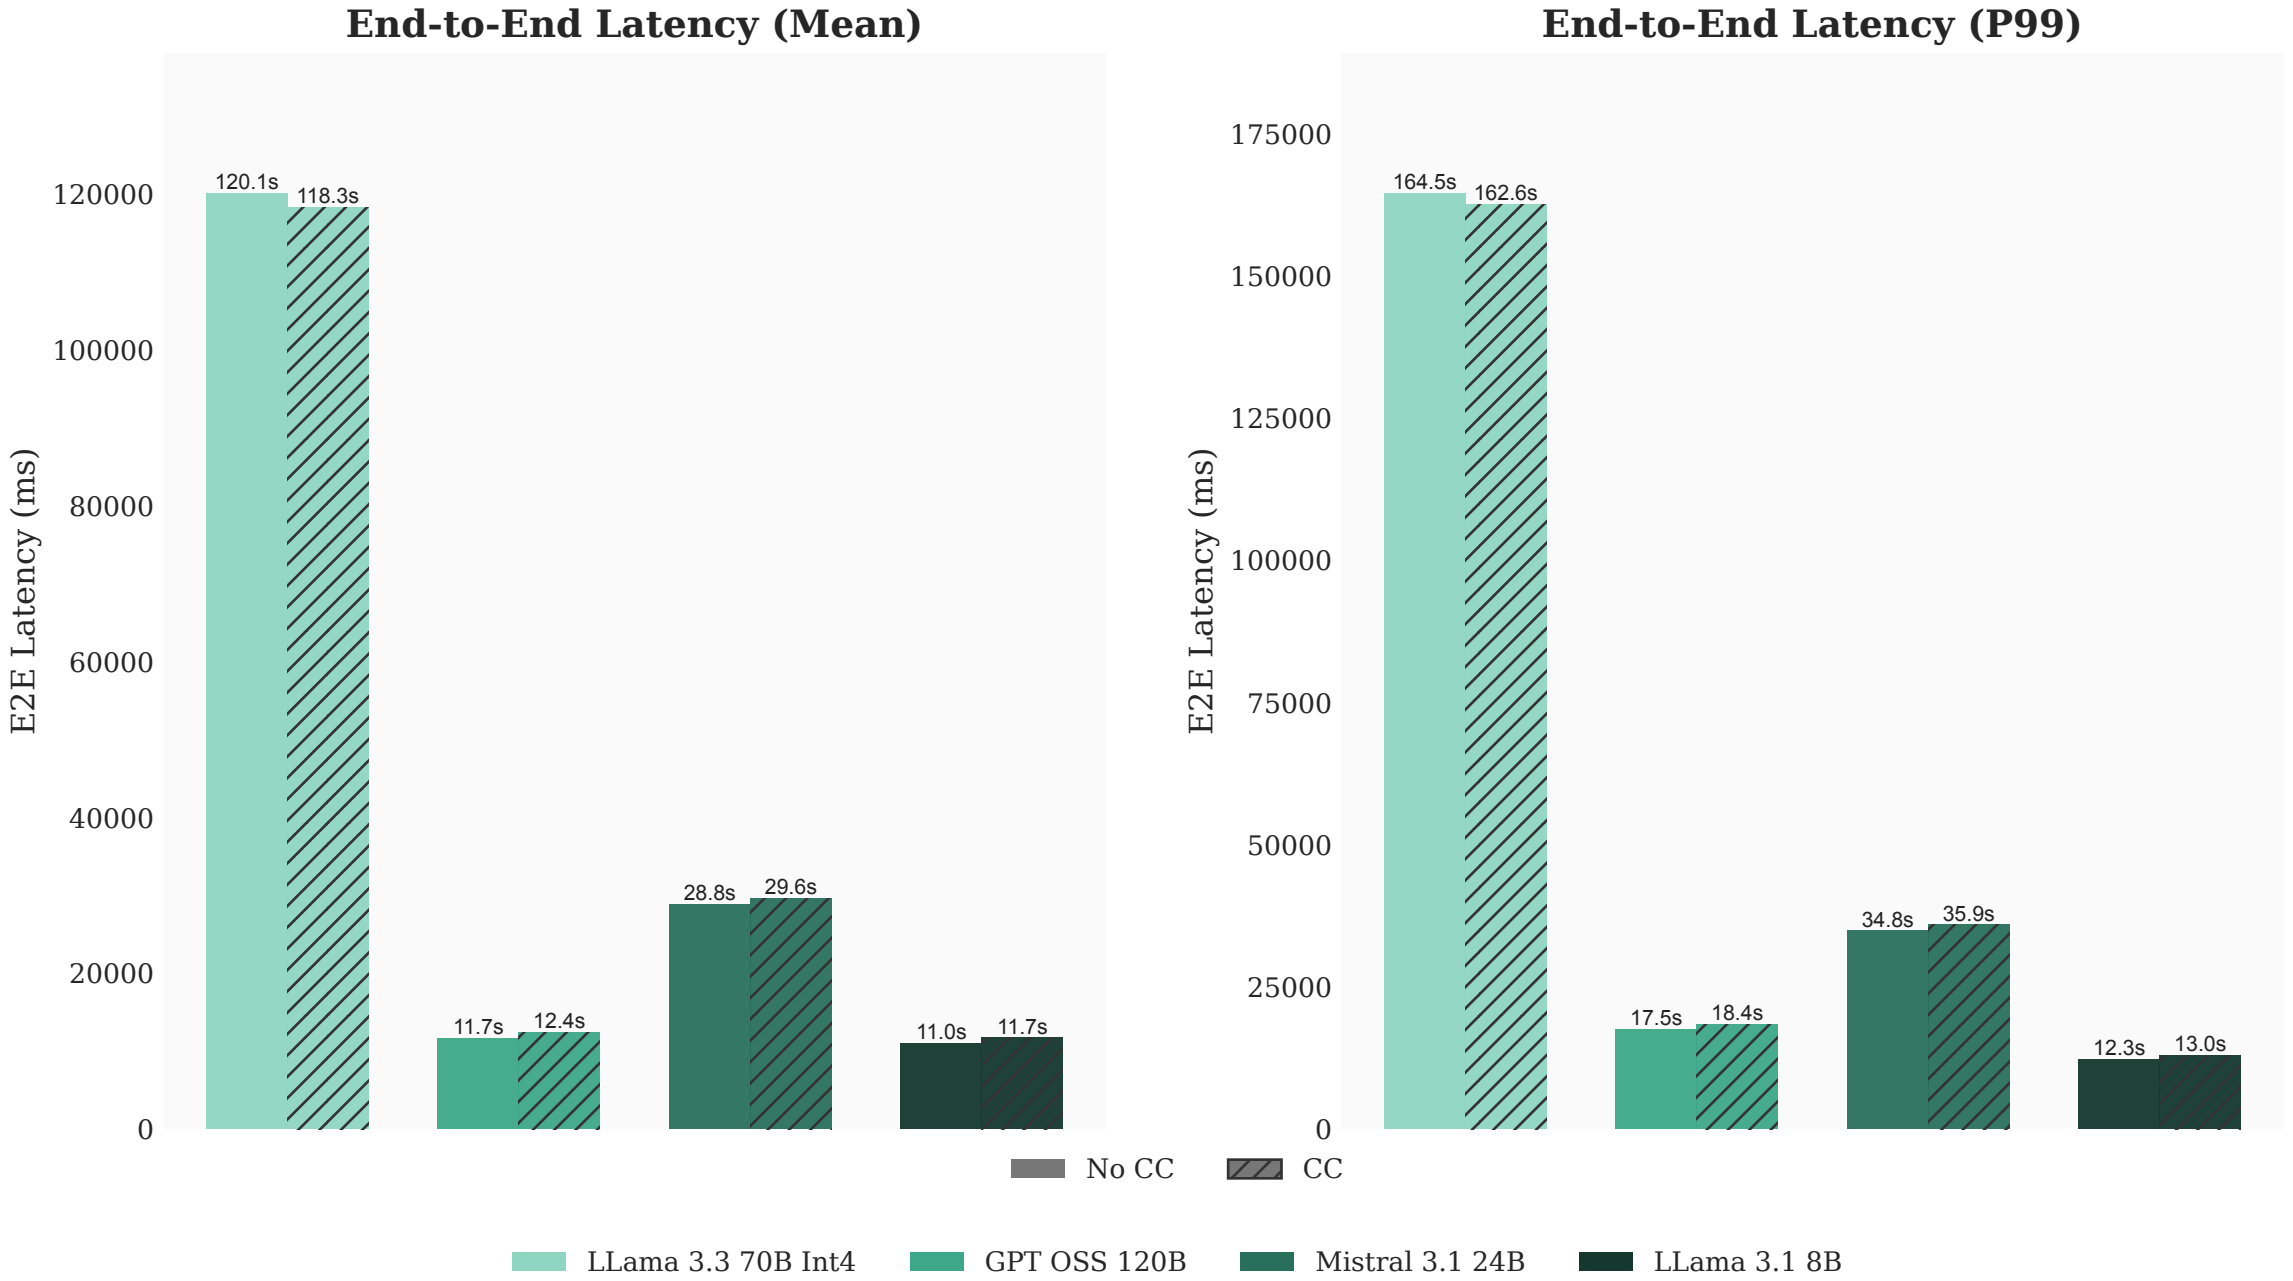
Random (1500 ⇒ 250) (Request Rate 100)
End-to-End Latency (Mean)
End-to-End Latency (P99)
E2E Latency (ms)
0
20000
40000
60000
80000
100000
120000
25000
50000
75000
100000
125000
150000
175000
120.1s
118.3s
11.7s
12.4s
28.8s
29.6s
11.0s
11.7s
164.5s
162.6s
17.5s
18.4s
34.8s
35.9s
12.3s
13.0s
No CC
CC
LLama 3.3 70B Int4
GPT OSS 120B
Mistral 3.1 24B
LLama 3.1 8B

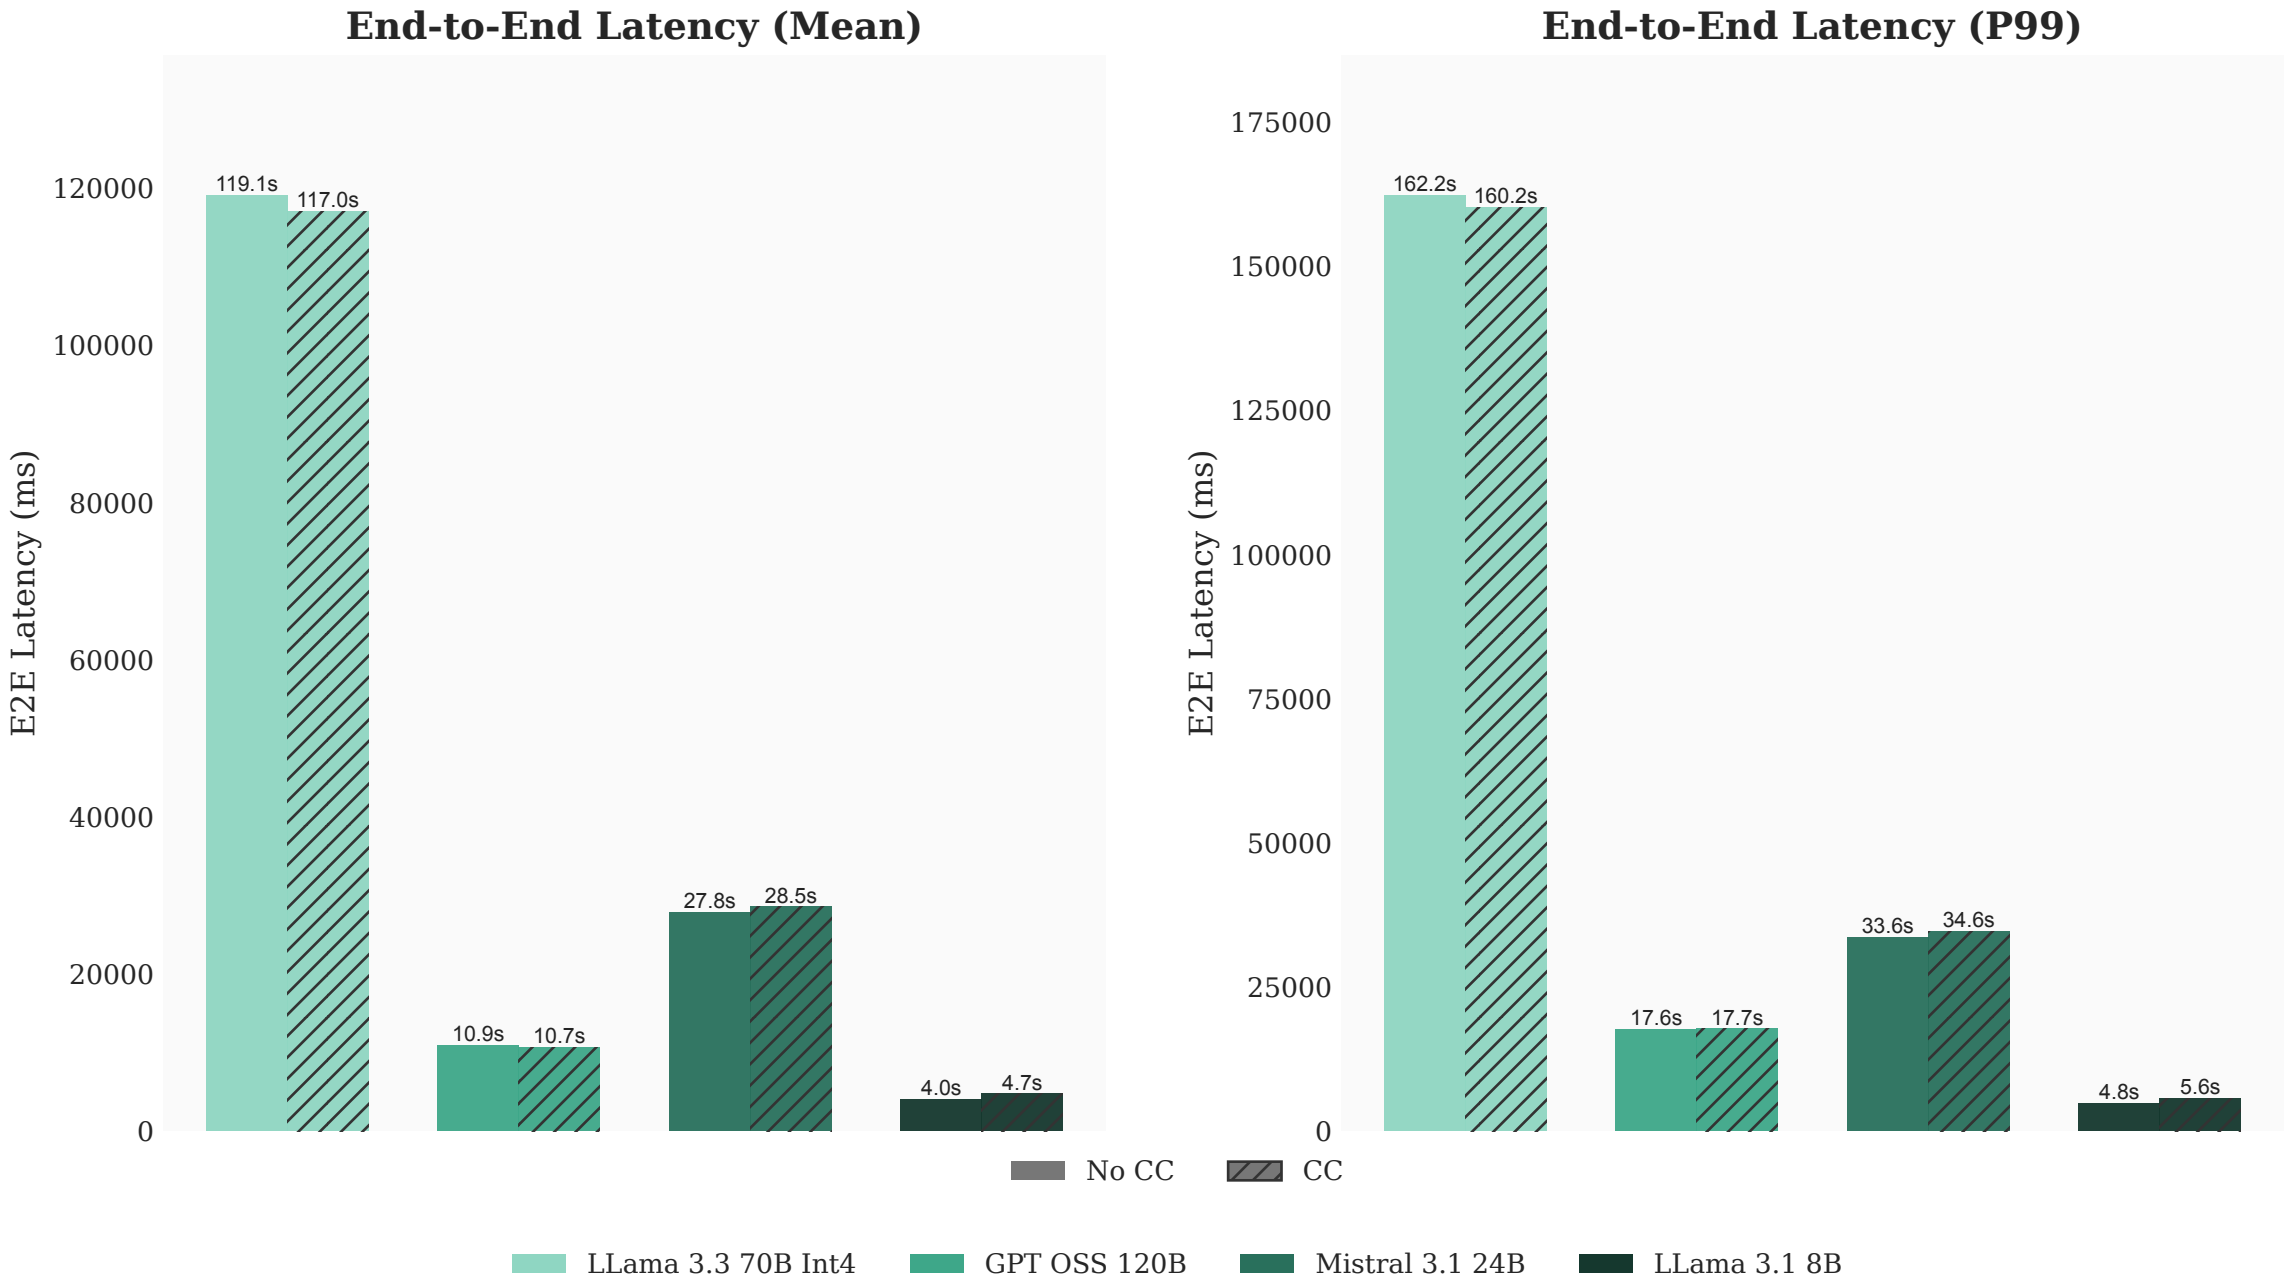
Random (1500 ⇒ 250) (Request Rate 50)
End-to-End Latency (Mean)
End-to-End Latency (P99)
E2E Latency (ms)
0
20000
40000
60000
80000
100000
120000
25000
50000
75000
100000
125000
150000
175000
119.1s
117.0s
10.9s
10.7s
27.8s
28.5s
4.0s
4.7s
162.2s
160.2s
17.6s
17.7s
33.6s
34.6s
4.8s
5.6s
No CC
CC
LLama 3.3 70B Int4
GPT OSS 120B
Mistral 3.1 24B
LLama 3.1 8B

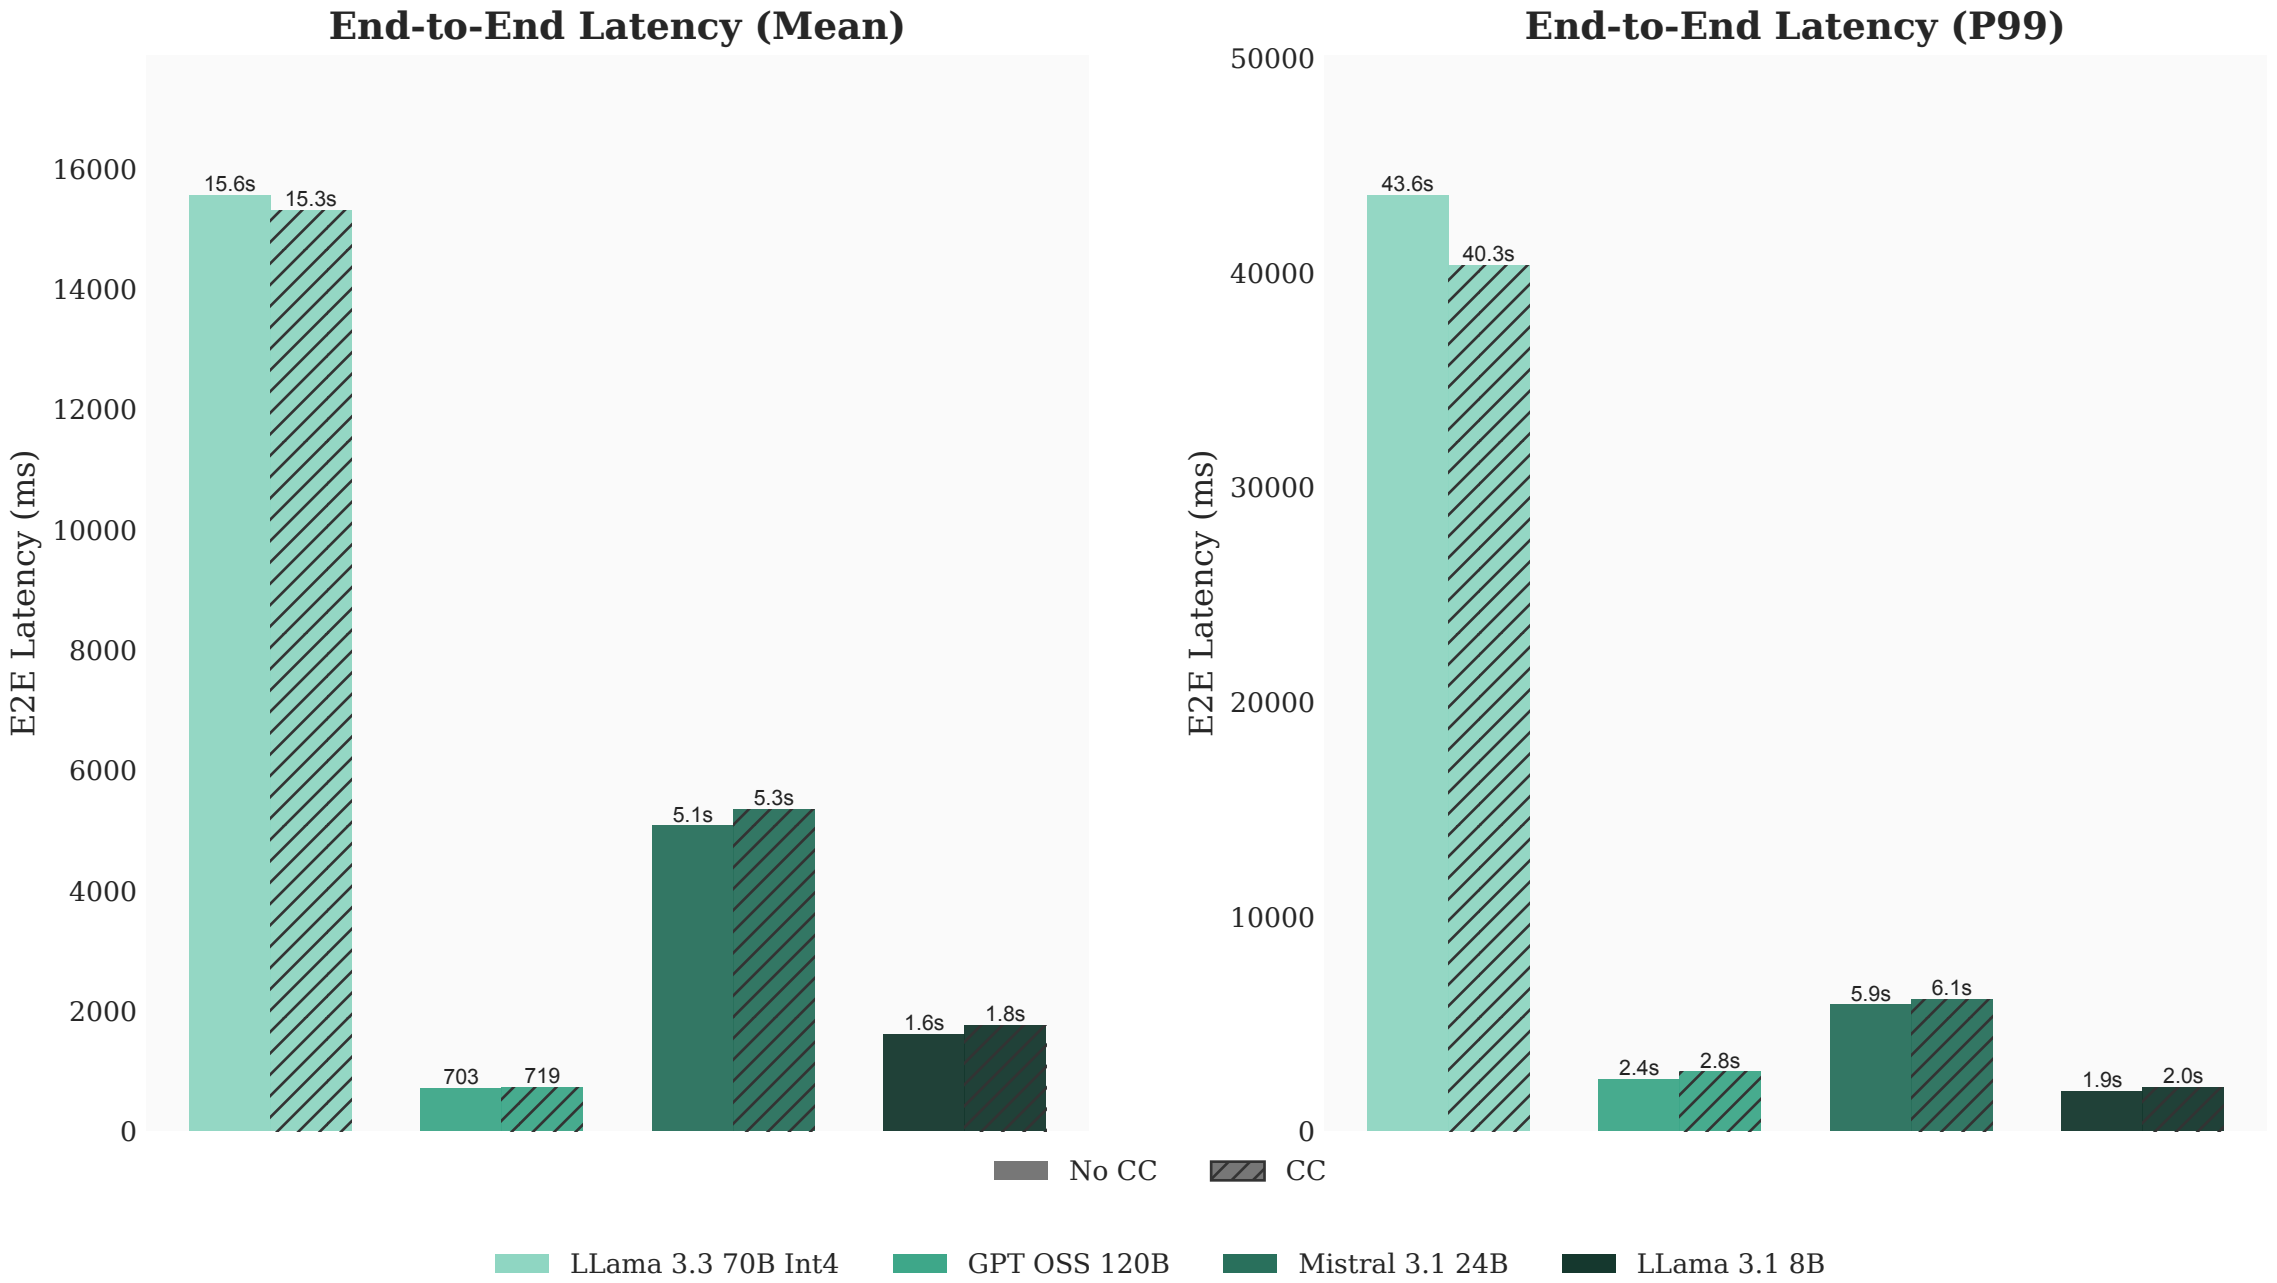
Random (1500 ⇒ 250) (Request Rate 1)
End-to-End Latency (Mean)
End-to-End Latency (P99)
E2E Latency (ms)
0
2000
4000
6000
8000
10000
12000
14000
16000
10000
20000
30000
40000
50000
15.6s
15.3s
703
719
5.1s
5.3s
1.6s
1.8s
43.6s
40.3s
2.4s
2.8s
5.9s
6.1s
1.9s
2.0s
No CC
CC
LLama 3.3 70B Int4
GPT OSS 120B
Mistral 3.1 24B
LLama 3.1 8B

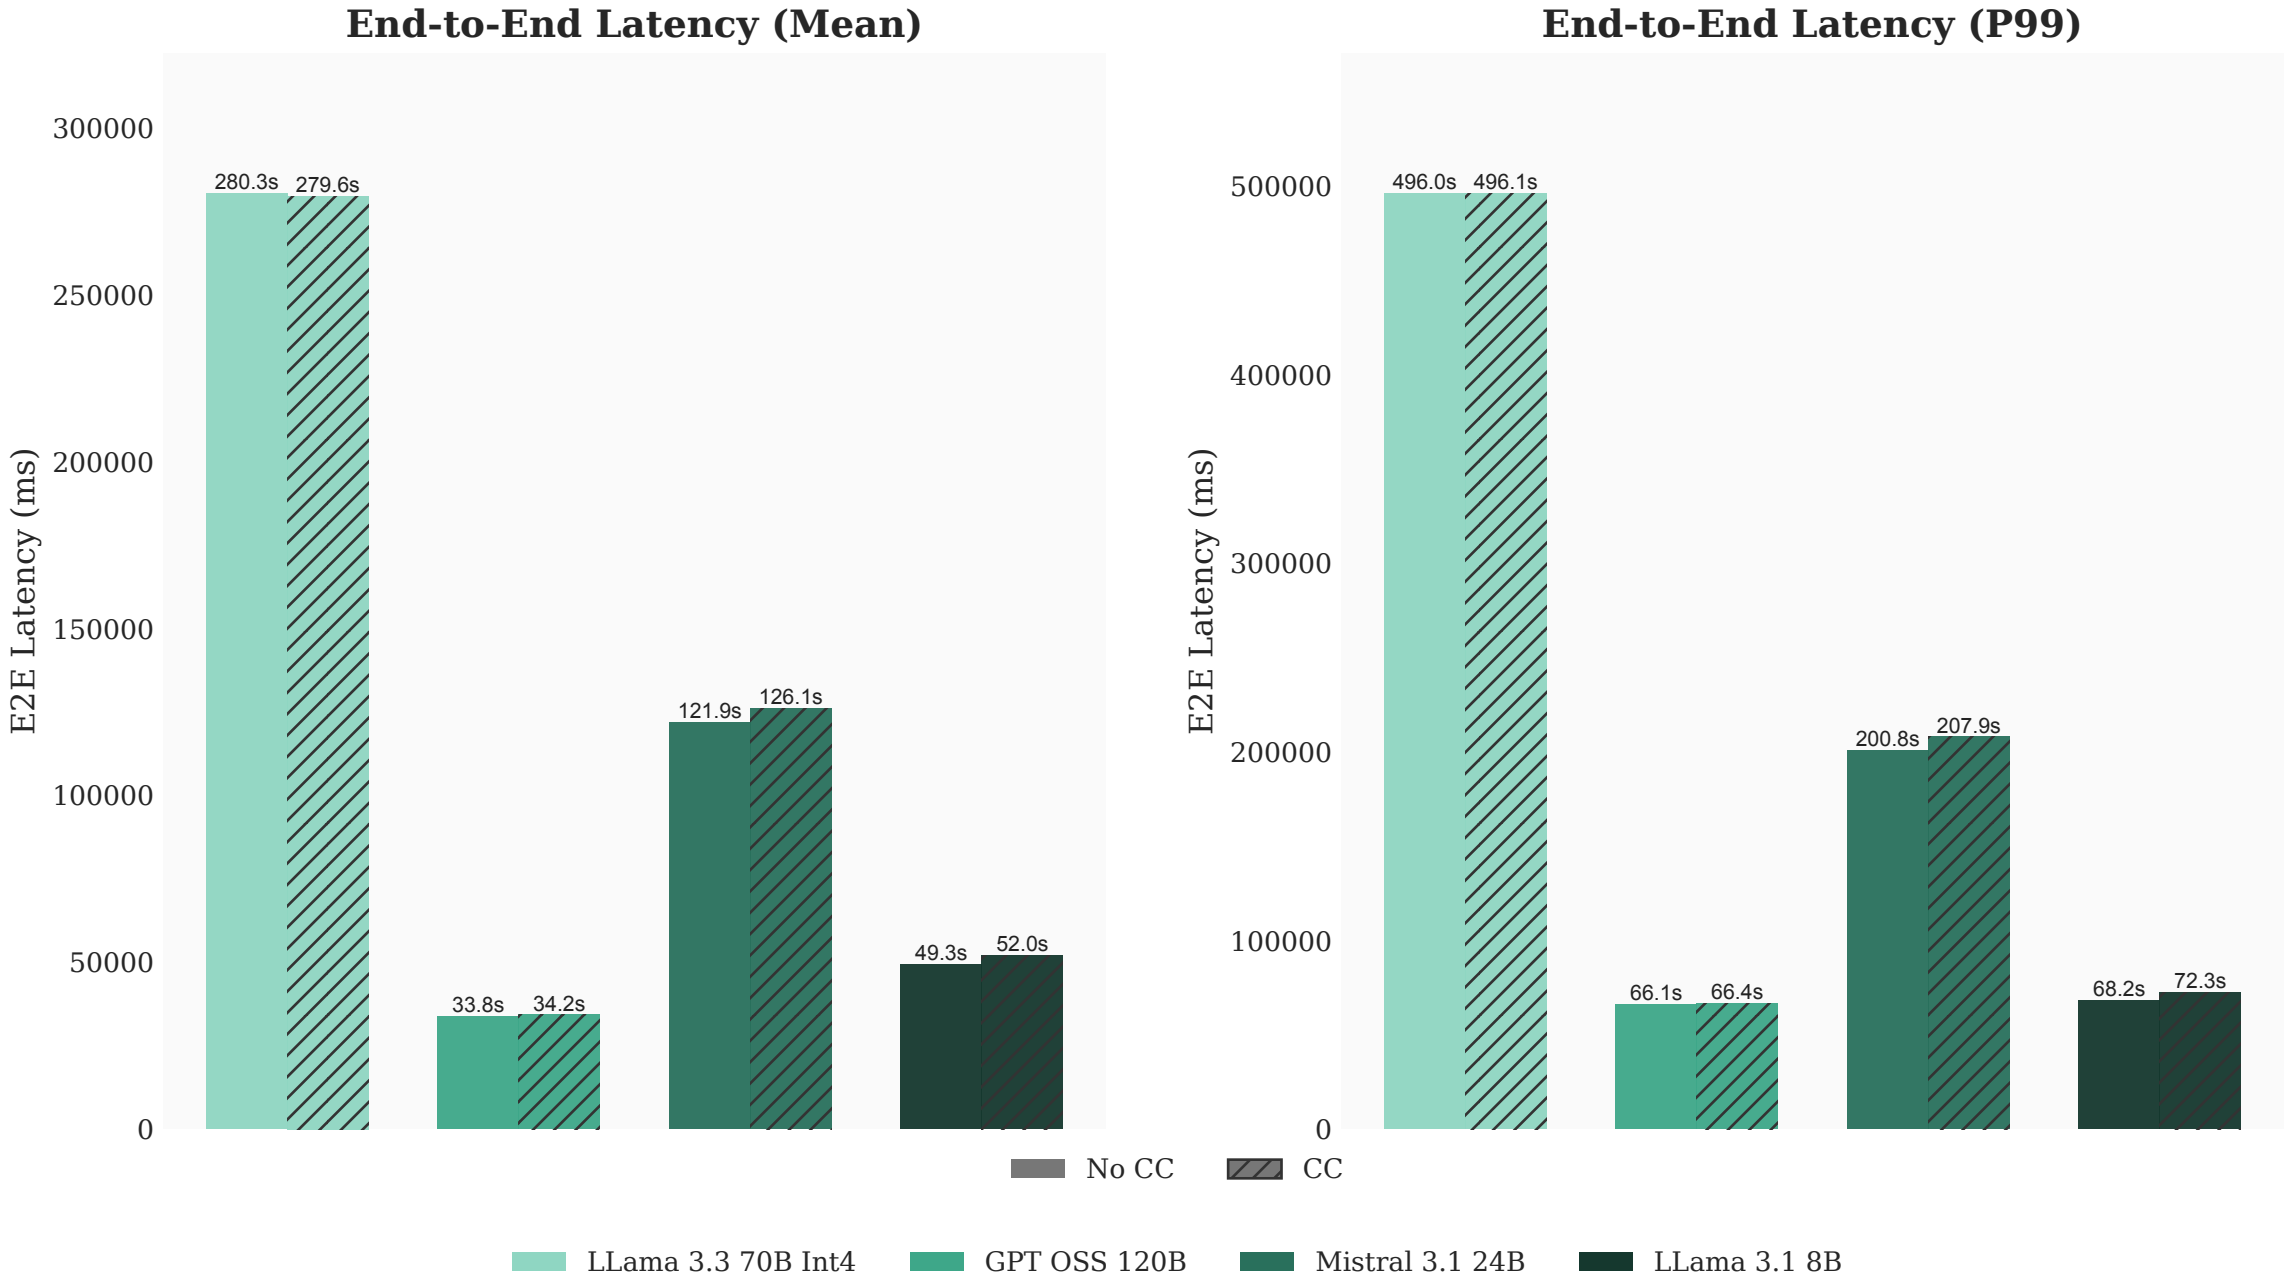
Random (4000 ⇒ 1000) (Request Rate 100)
End-to-End Latency (Mean)
End-to-End Latency (P99)
E2E Latency (ms)
300000
250000
200000
150000
100000
50000
0
500000
400000
300000
200000
100000
280.3s
279.6s
33.8s
34.2s
121.9s
126.1s
49.3s
52.0s
496.0s
496.1s
66.1s
66.4s
200.8s
207.9s
68.2s
72.3s
No CC
CC
LLama 3.3 70B Int4
GPT OSS 120B
Mistral 3.1 24B
LLama 3.1 8B

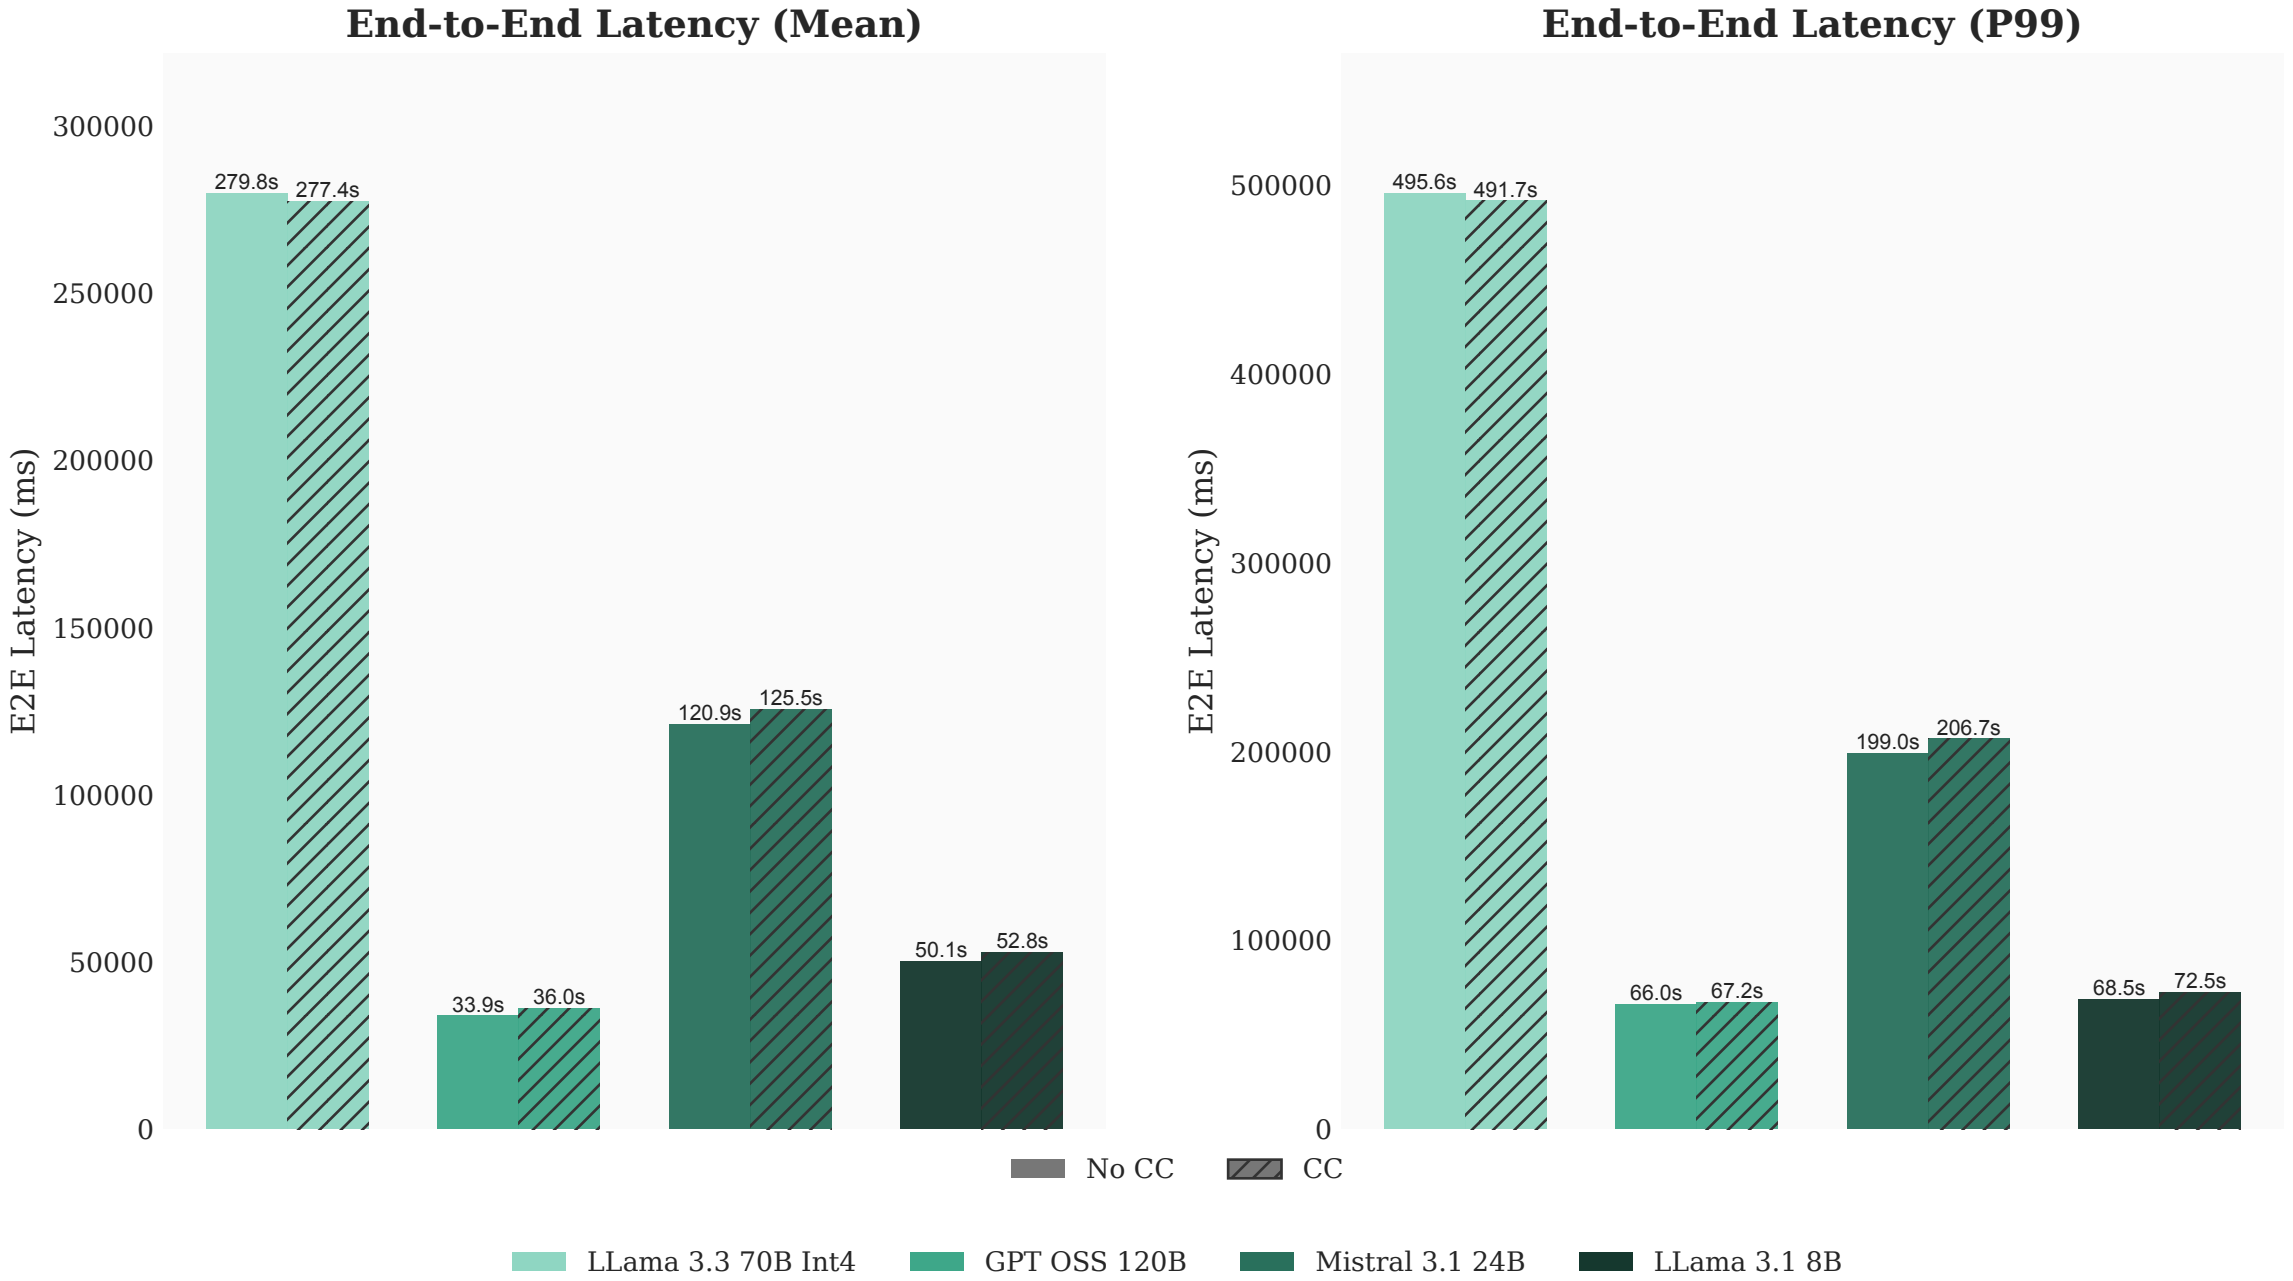
Random (4000 ⇒ 1000) (Request Rate 50)
End-to-End Latency (Mean)
End-to-End Latency (P99)
E2E Latency (ms)
E2E Latency (ms)
300000
250000
200000
150000
100000
50000
0
500000
400000
300000
200000
100000
0
279.8s
277.4s
33.9s
36.0s
120.9s
125.5s
50.1s
52.8s
495.6s
491.7s
66.0s
67.2s
199.0s
206.7s
68.5s
72.5s
No CC
CC
LLama 3.3 70B Int4
GPT OSS 120B
Mistral 3.1 24B
LLama 3.1 8B

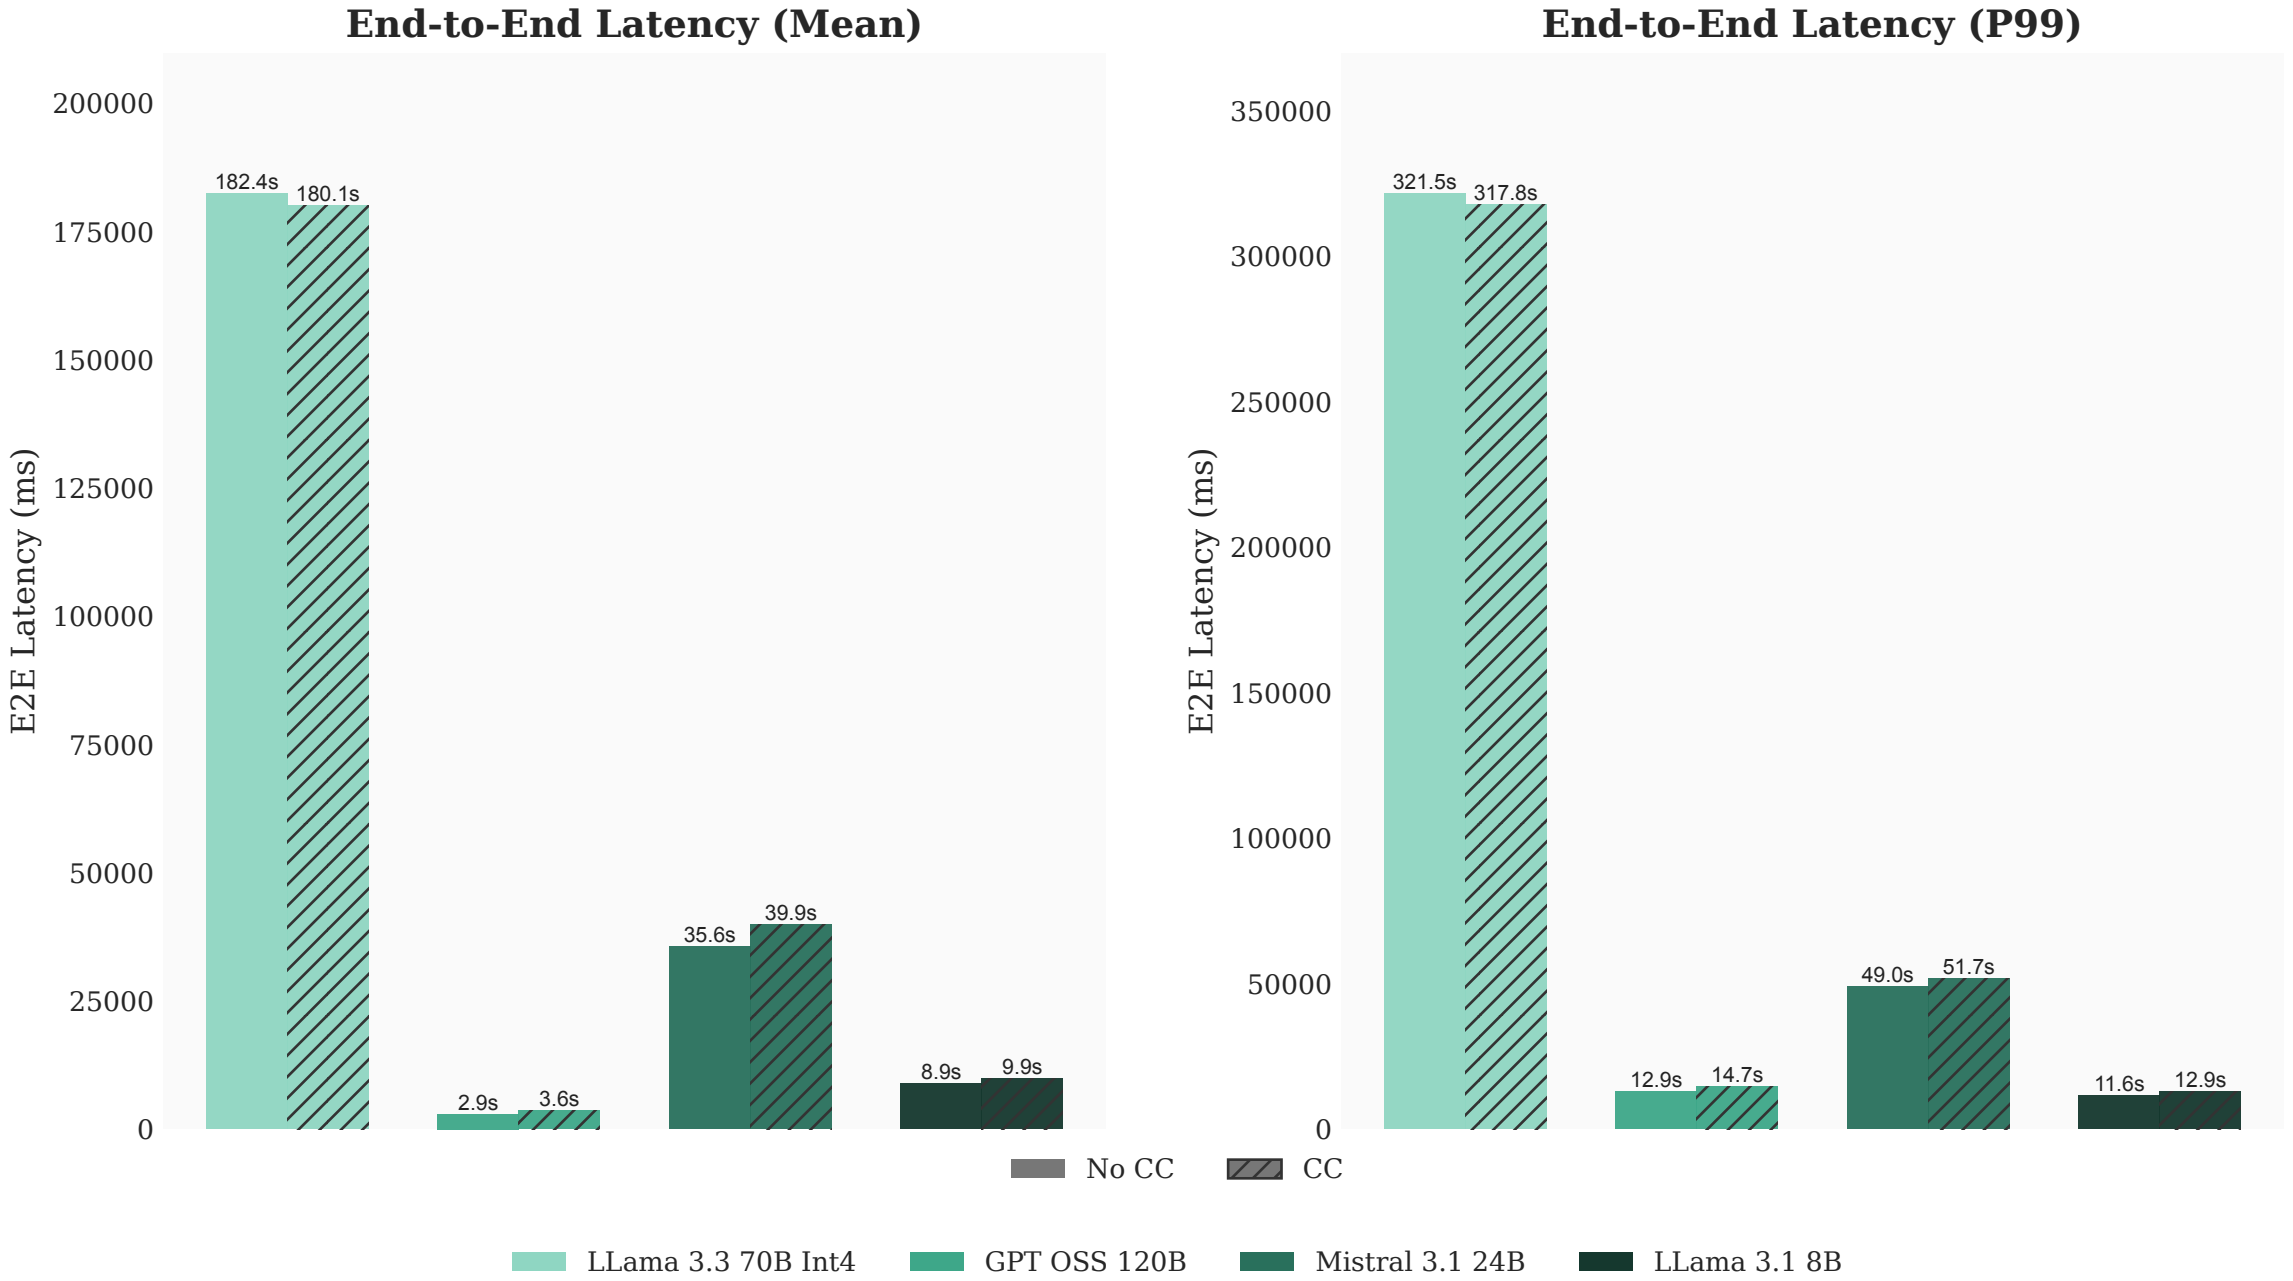
Random (4000 ⇒ 1000) (Request Rate 1)
End-to-End Latency (Mean)
End-to-End Latency (P99)
E2E Latency (ms)
200000
175000
150000
125000
100000
75000
50000
25000
0
350000
300000
250000
200000
150000
100000
50000
0
182.4s
180.1s
2.9s
3.6s
35.6s
39.9s
8.9s
9.9s
321.5s
317.8s
12.9s
14.7s
49.0s
51.7s
11.6s
12.9s
No CC
CC
LLama 3.3 70B Int4
GPT OSS 120B
Mistral 3.1 24B
LLama 3.1 8B

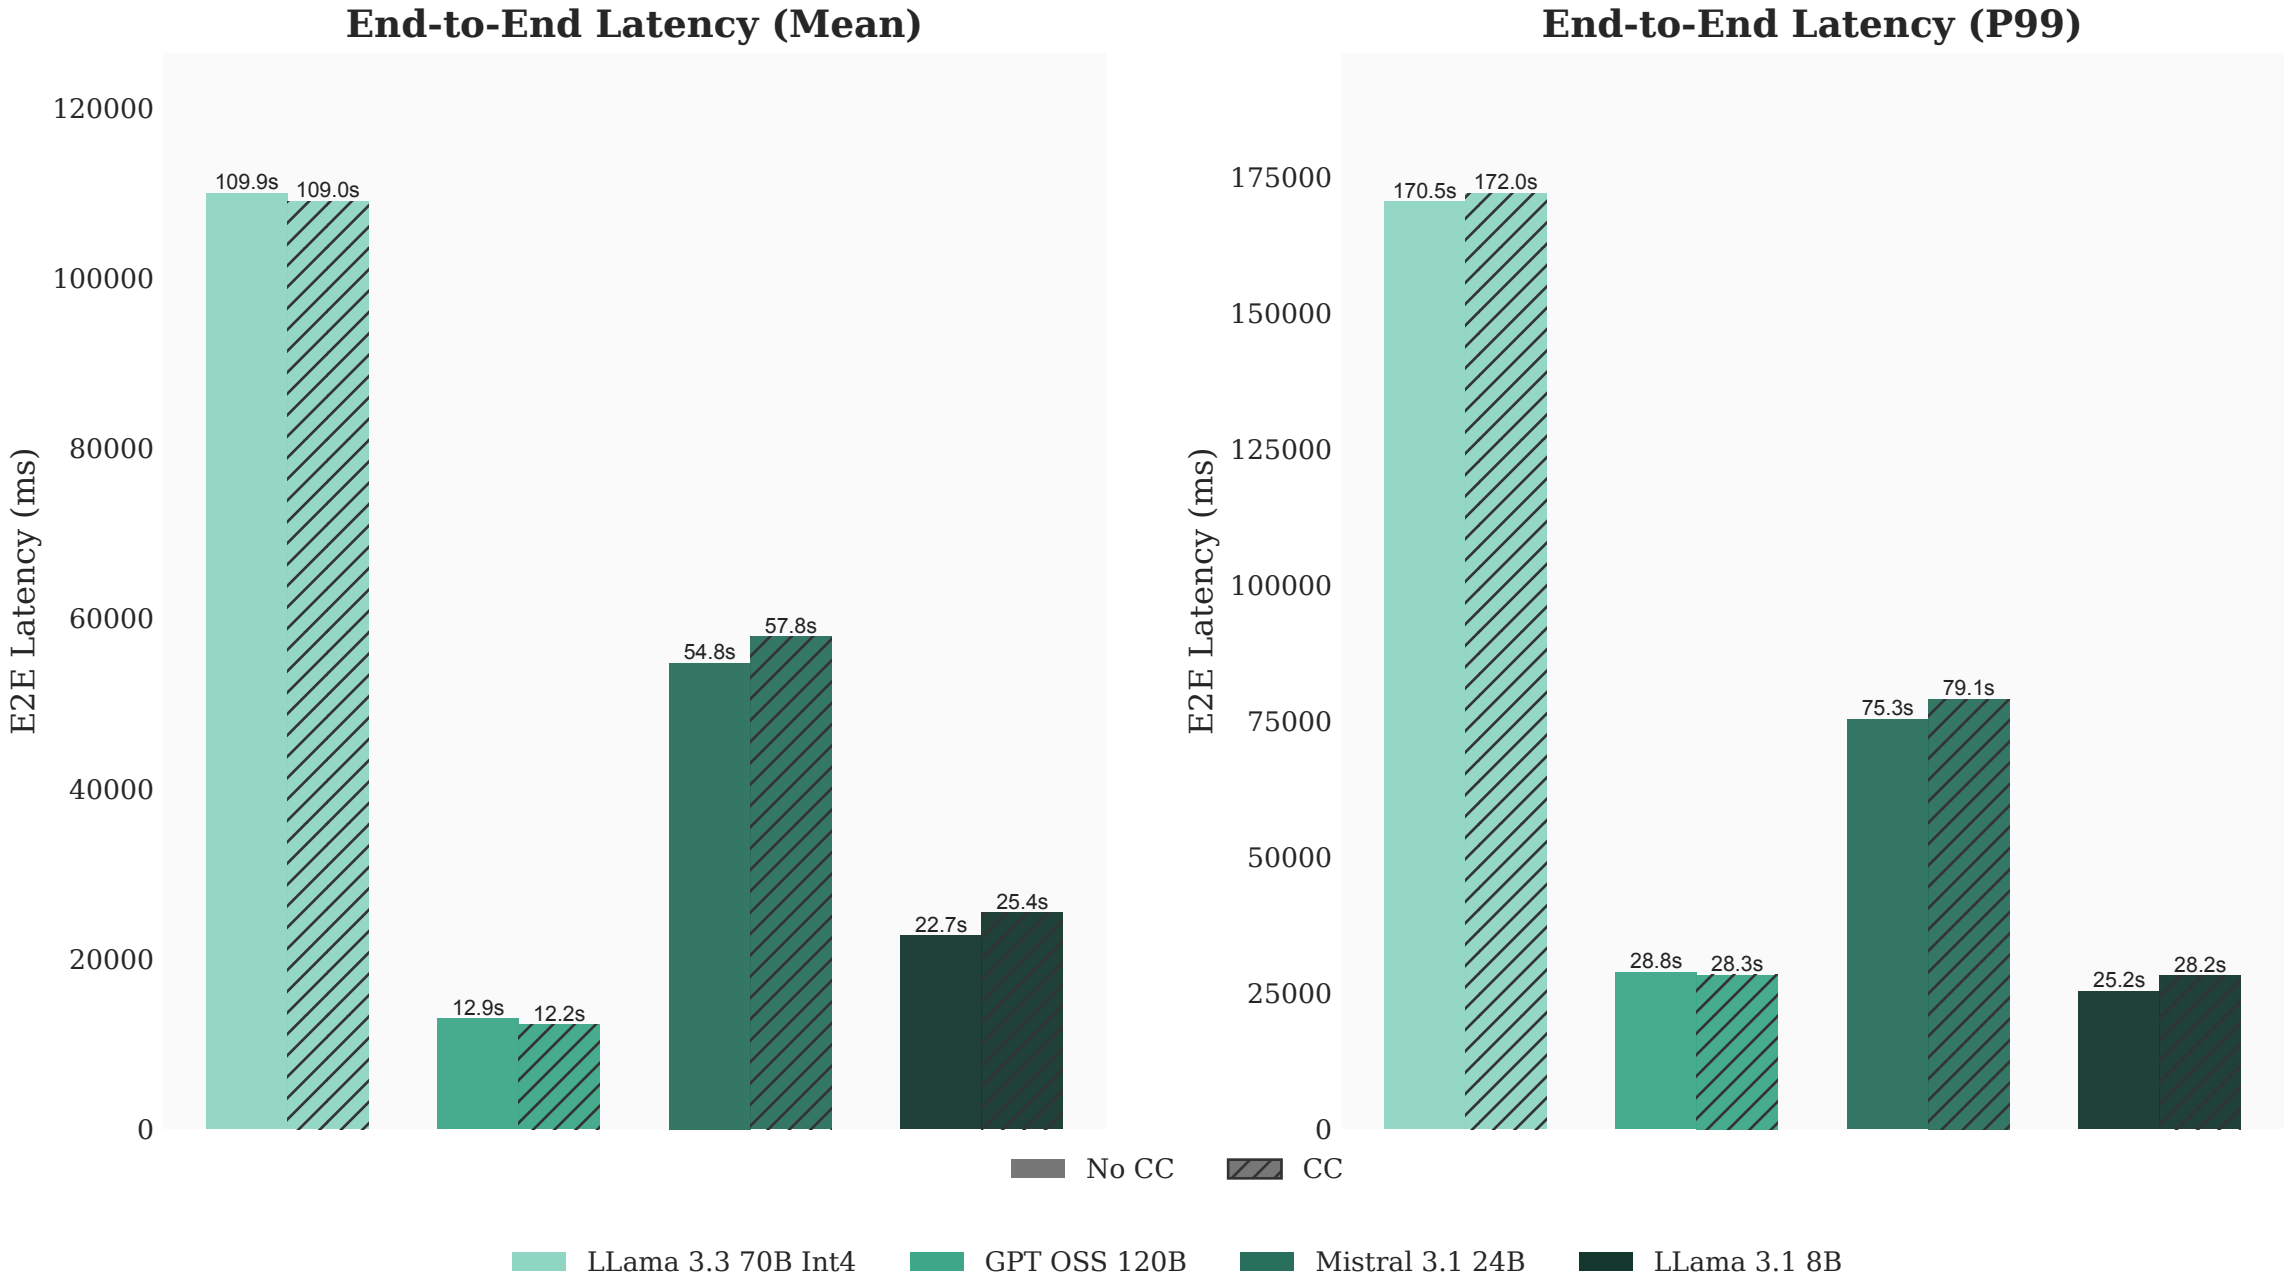
Random (1000 ⇒ 1000) (Request Rate 100)
End-to-End Latency (Mean)
End-to-End Latency (P99)
E2E Latency (ms)
E2E Latency (ms)
0
20000
40000
60000
80000
100000
120000
0
25000
50000
75000
100000
125000
150000
175000
109.9s
109.0s
12.9s
12.2s
54.8s
57.8s
22.7s
25.4s
170.5s
172.0s
28.8s
28.3s
75.3s
79.1s
25.2s
28.2s
No CC
CC
LLama 3.3 70B Int4
GPT OSS 120B
Mistral 3.1 24B
LLama 3.1 8B

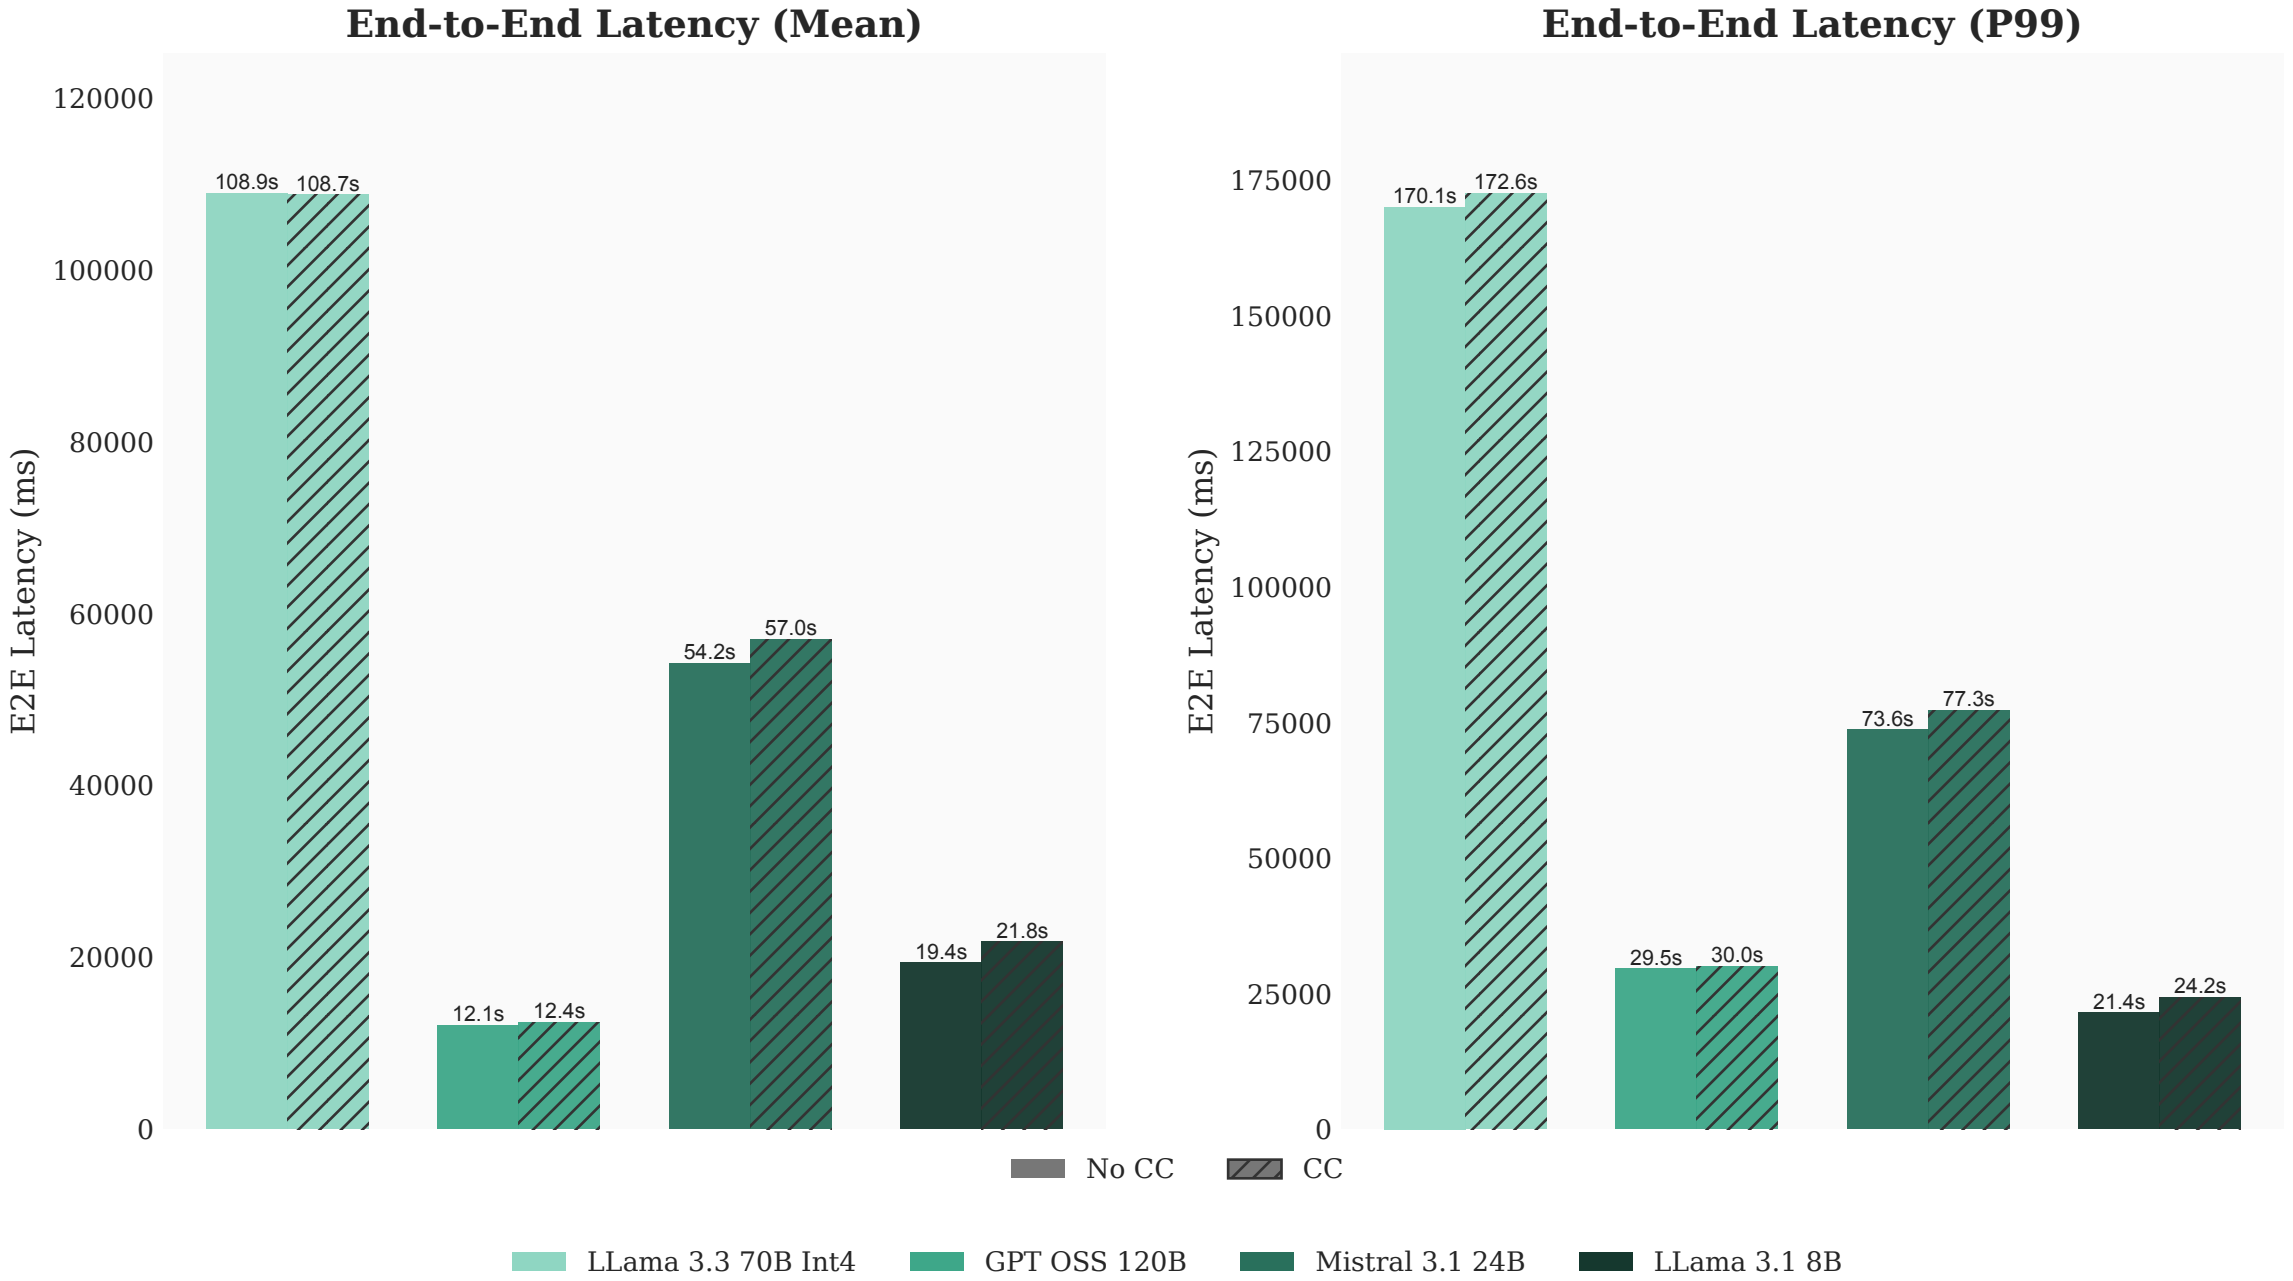
Random (1000 ⇒ 1000) (Request Rate 50)
End-to-End Latency (Mean)
End-to-End Latency (P99)
E2E Latency (ms)
0
20000
40000
60000
80000
100000
120000
25000
50000
75000
100000
125000
150000
175000
108.9s
108.7s
12.1s
12.4s
54.2s
57.0s
19.4s
21.8s
170.1s
172.6s
29.5s
30.0s
73.6s
77.3s
21.4s
24.2s
No CC
CC
LLama 3.3 70B Int4
GPT OSS 120B
Mistral 3.1 24B
LLama 3.1 8B

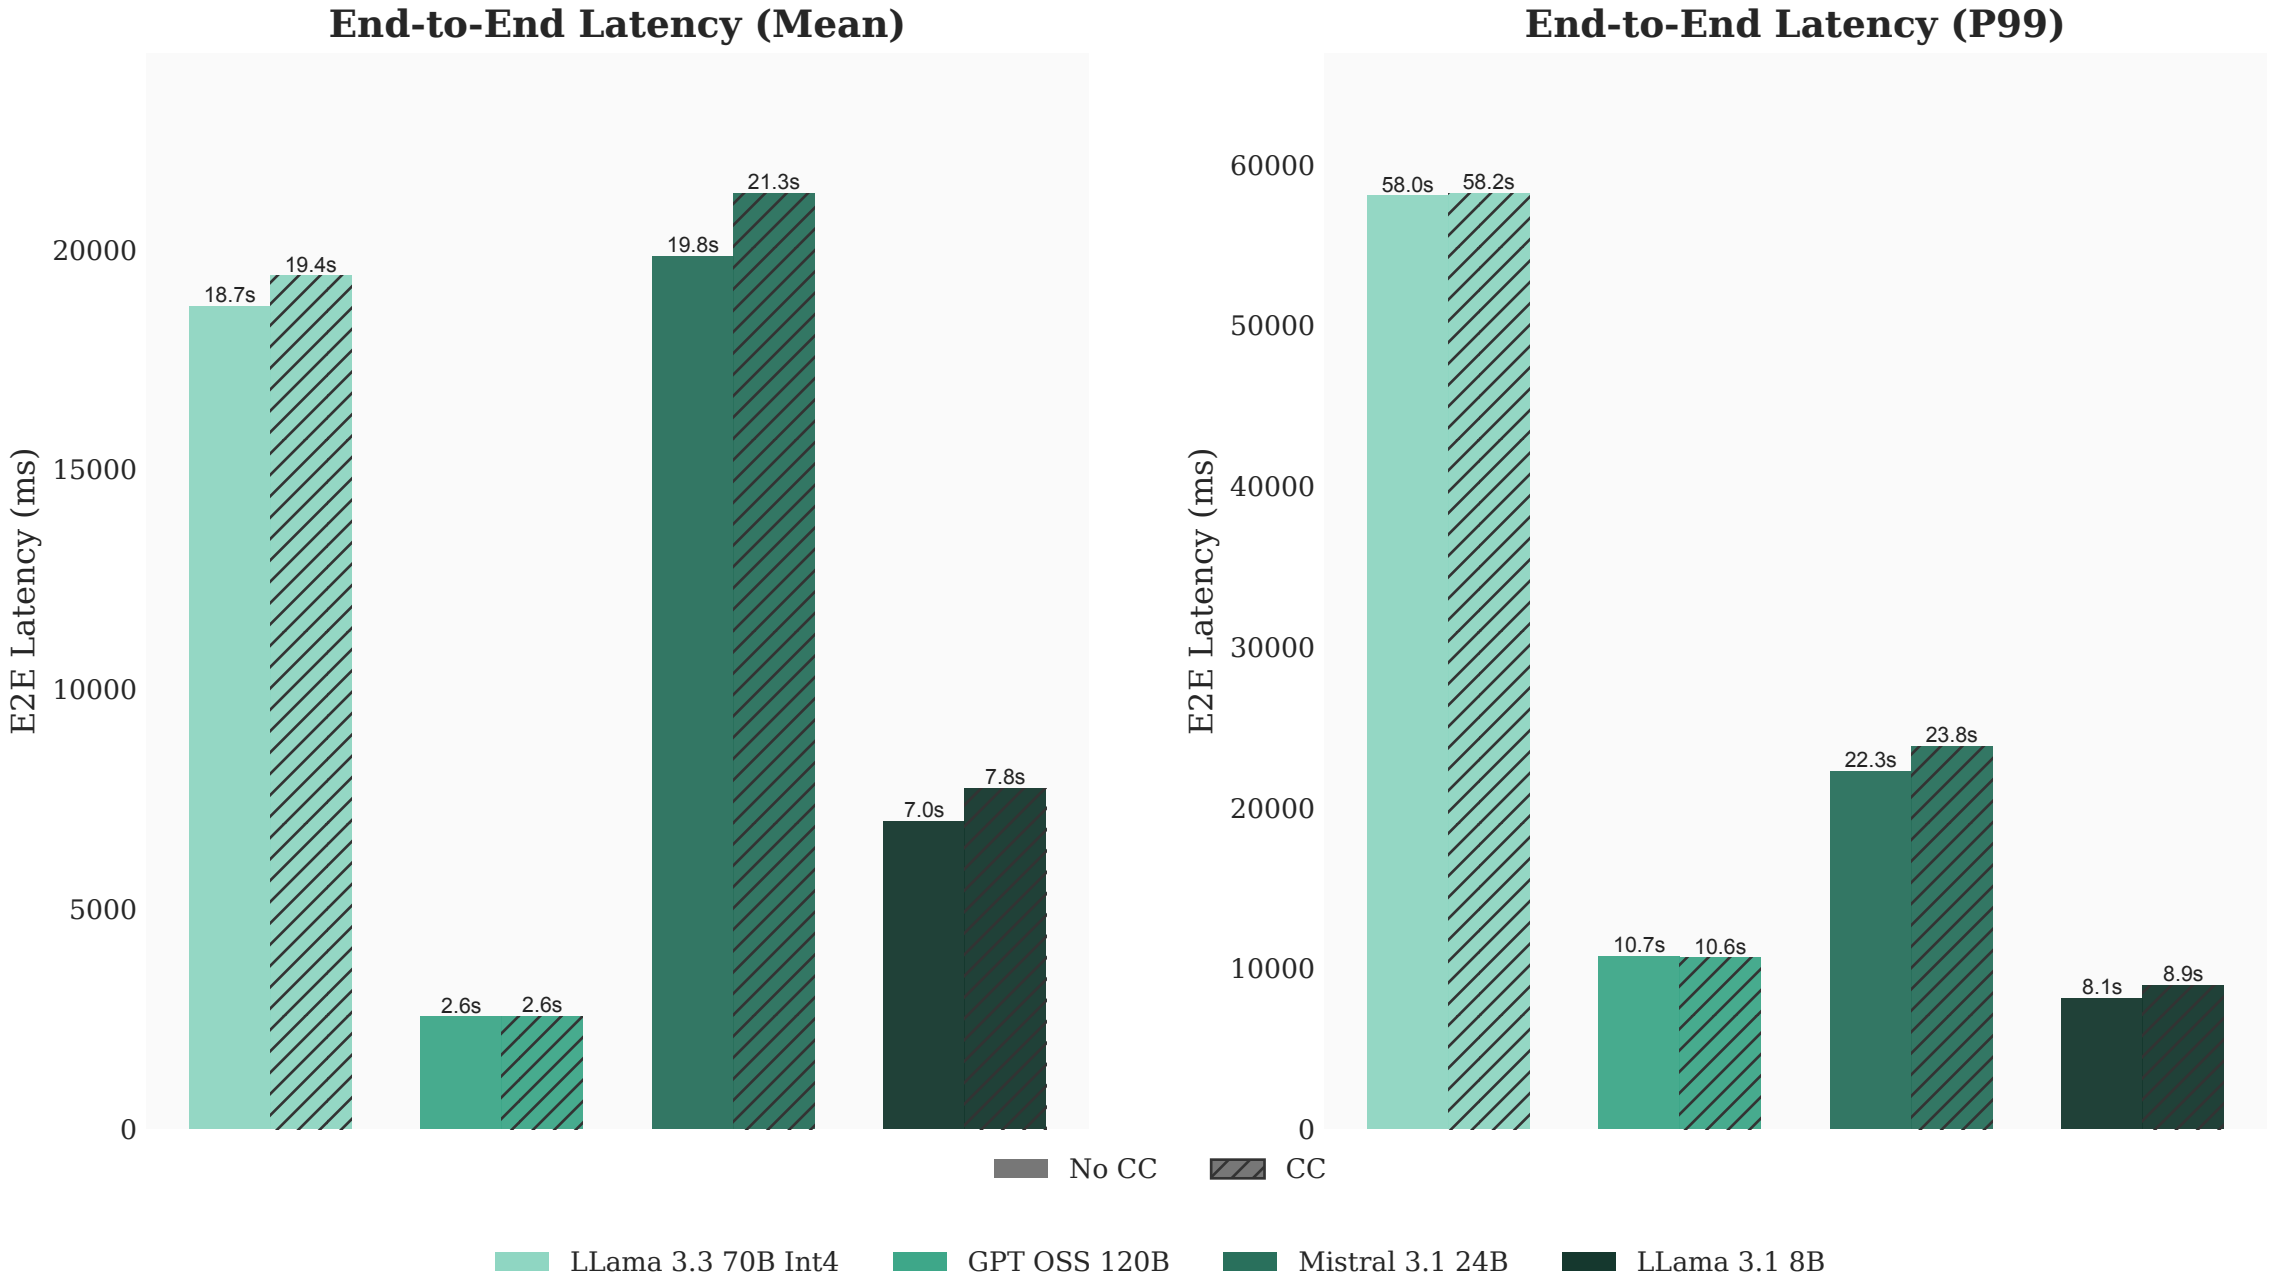
Random (1000 ⇒ 1000) (Request Rate 1)
End-to-End Latency (Mean)
End-to-End Latency (P99)
E2E Latency (ms)
0
5000
10000
15000
20000
30000
40000
50000
60000
18.7s
19.4s
2.6s
2.6s
19.8s
21.3s
7.0s
7.8s
58.0s
58.2s
10.7s
10.6s
22.3s
23.8s
8.1s
8.9s
No CC
CC
LLama 3.3 70B Int4
GPT OSS 120B
Mistral 3.1 24B
LLama 3.1 8B

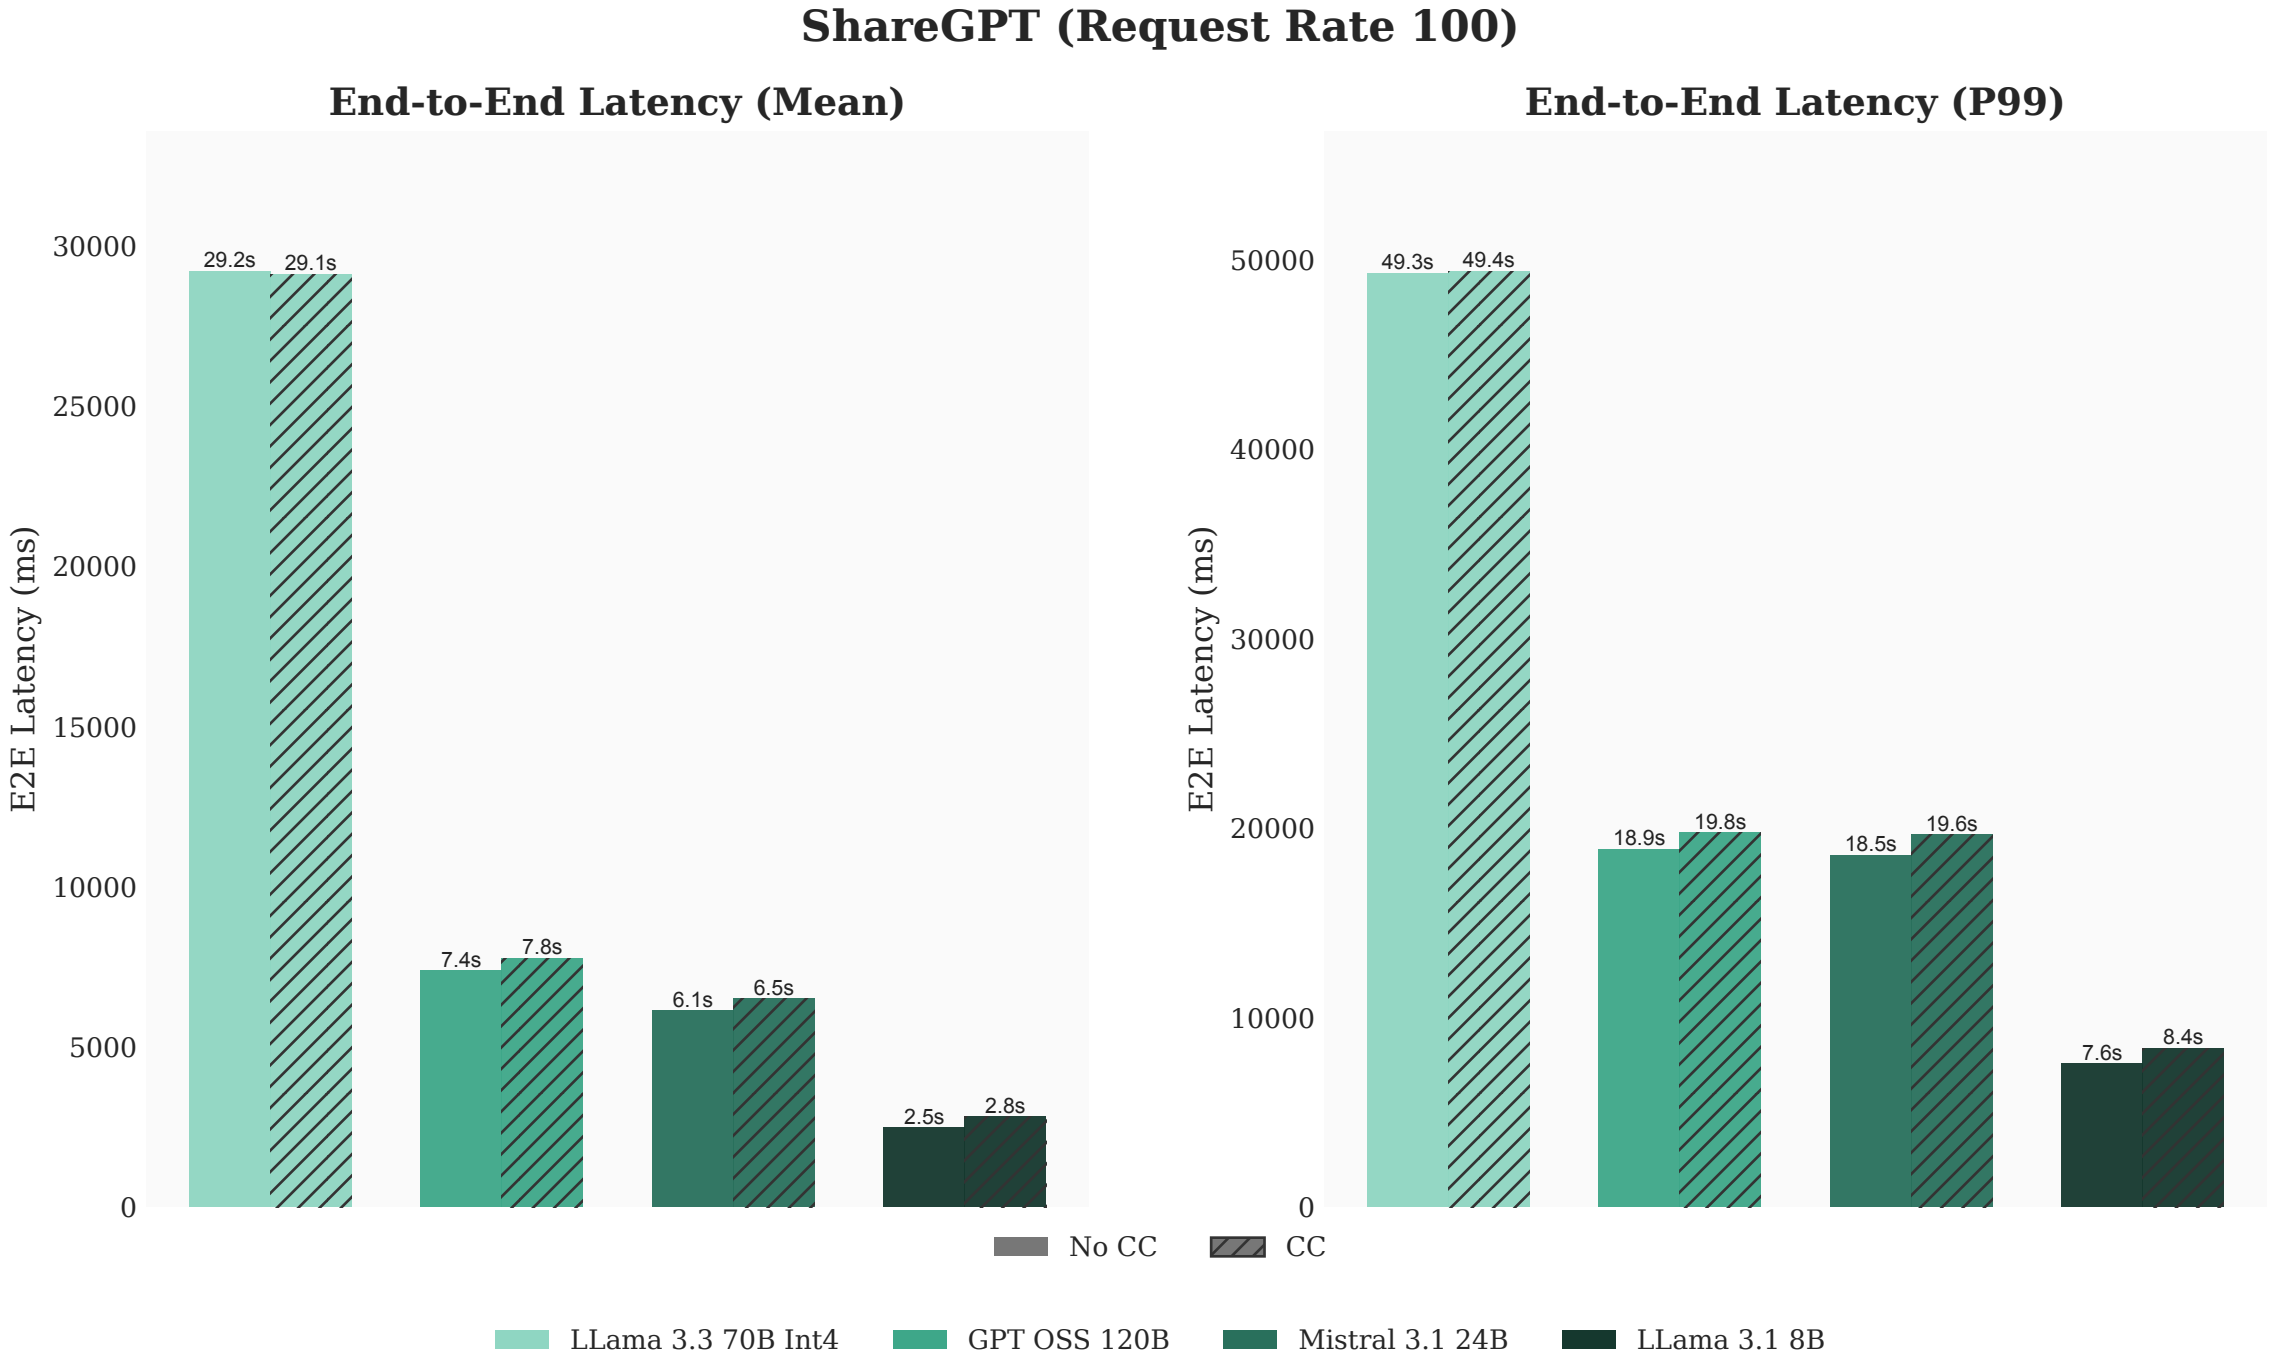
ShareGPT (Request Rate 100)
End-to-End Latency (Mean)
End-to-End Latency (P99)
E2E Latency (ms)
0
5000
10000
15000
20000
25000
30000
40000
50000
29.2s
29.1s
7.4s
7.8s
6.1s
6.5s
2.5s
2.8s
49.3s
49.4s
18.9s
19.8s
18.5s
19.6s
7.6s
8.4s
No CC
CC
LLama 3.3 70B Int4
GPT OSS 120B
Mistral 3.1 24B
LLama 3.1 8B

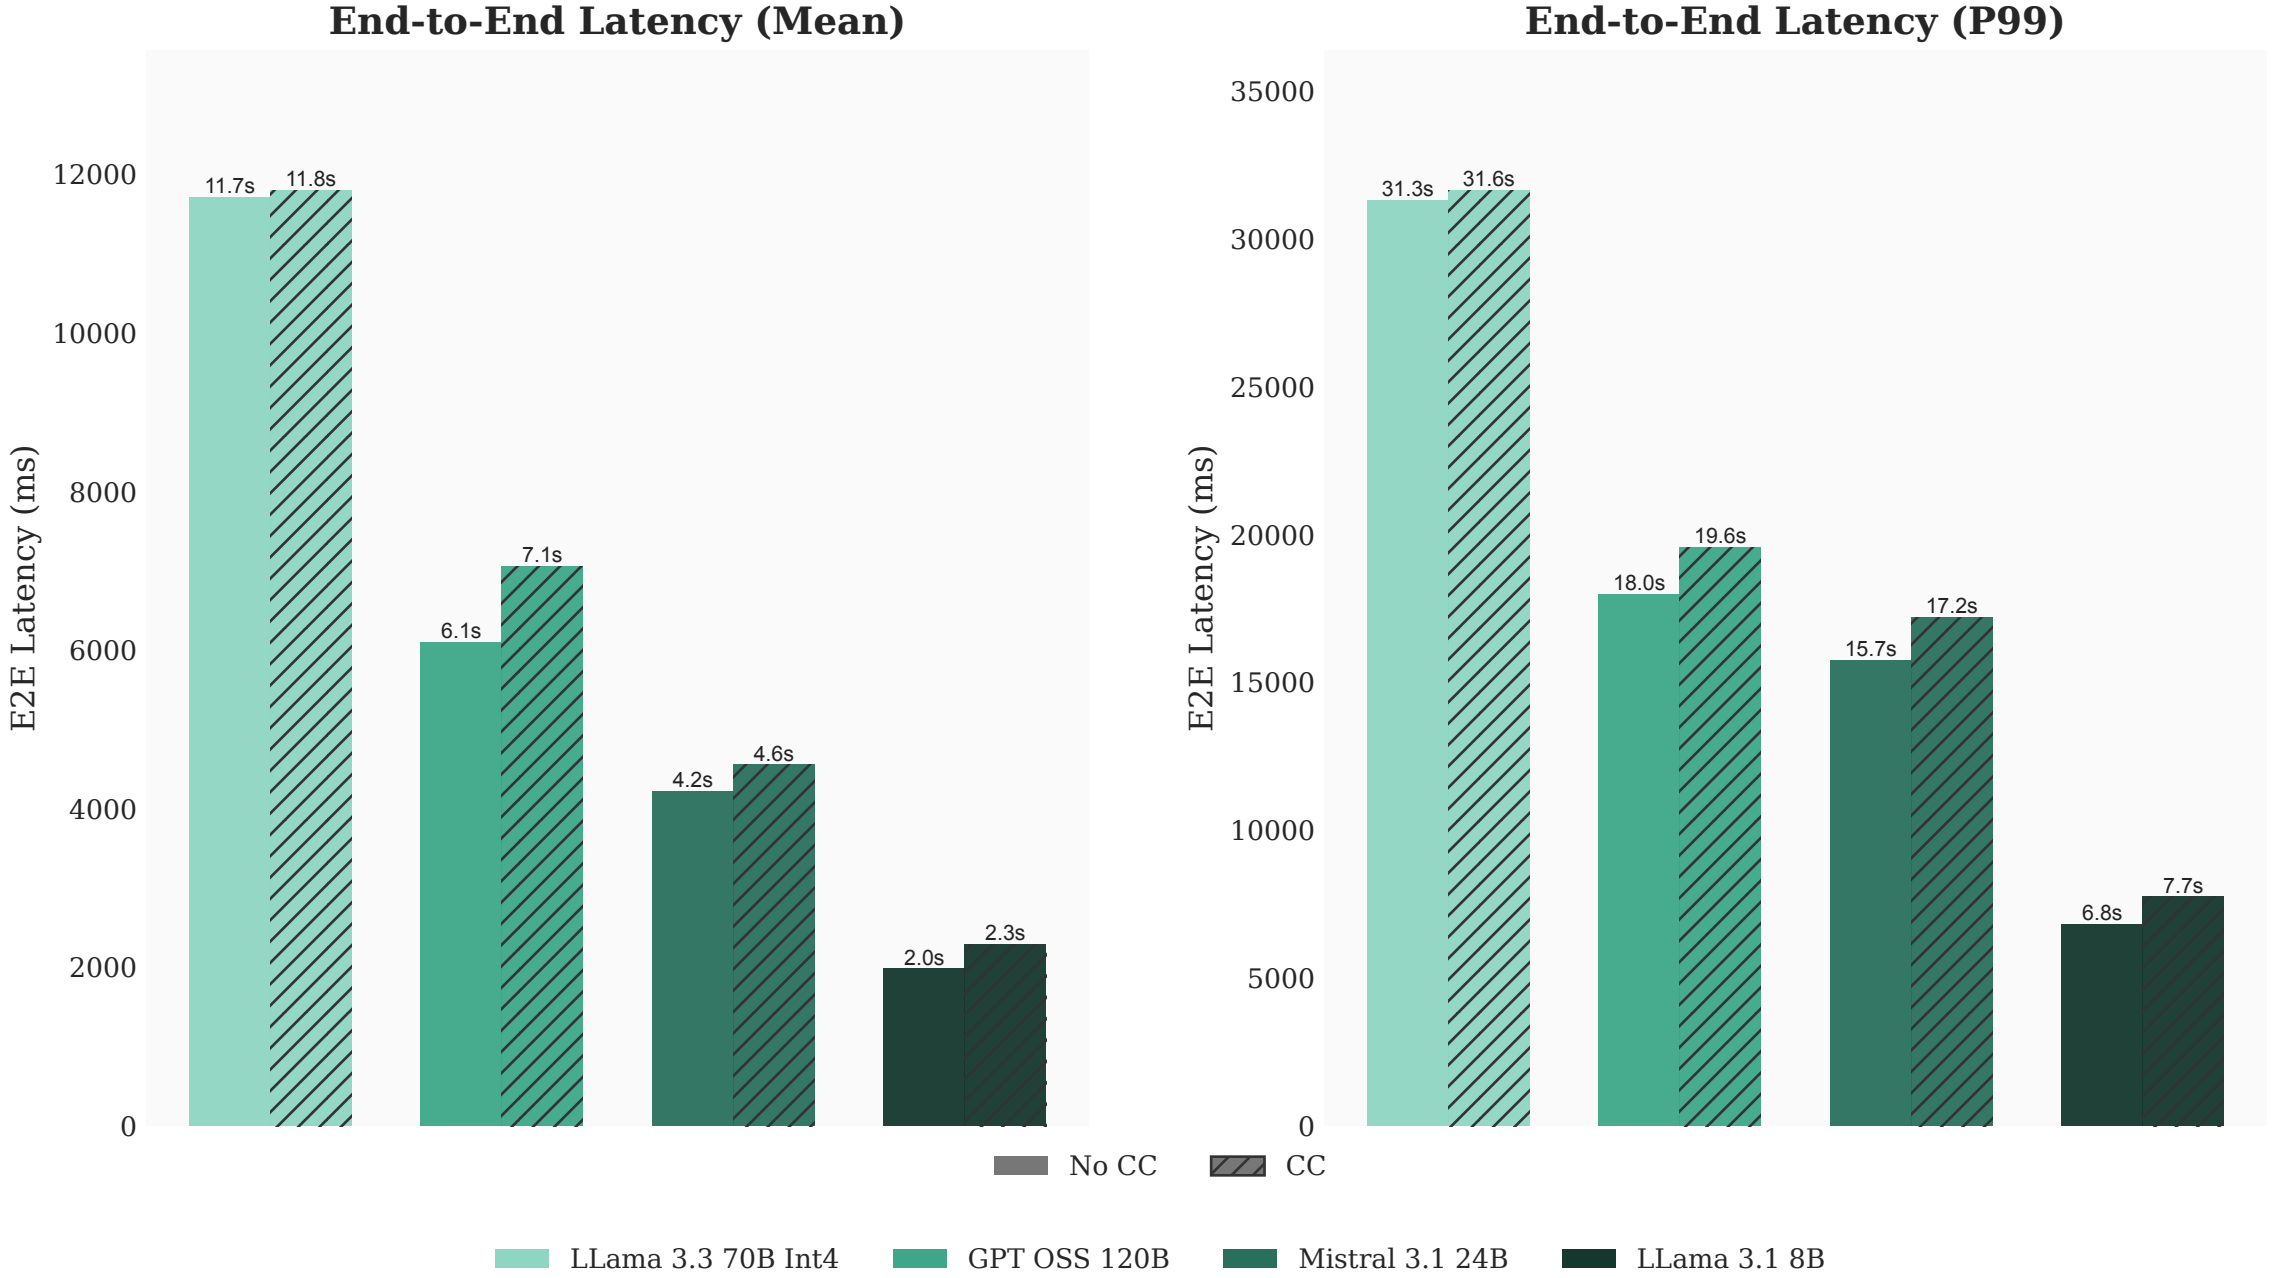
ShareGPT (Request Rate 50)
End-to-End Latency (Mean)
End-to-End Latency (P99)
E2E Latency (ms)
0
2000
4000
6000
8000
10000
12000
5000
15000
20000
25000
30000
35000
11.7s
11.8s
6.1s
7.1s
4.2s
4.6s
2.0s
2.3s
31.3s
31.6s
18.0s
19.6s
15.7s
17.2s
6.8s
7.7s
No CC
CC
LLama 3.3 70B Int4
GPT OSS 120B
Mistral 3.1 24B
LLama 3.1 8B

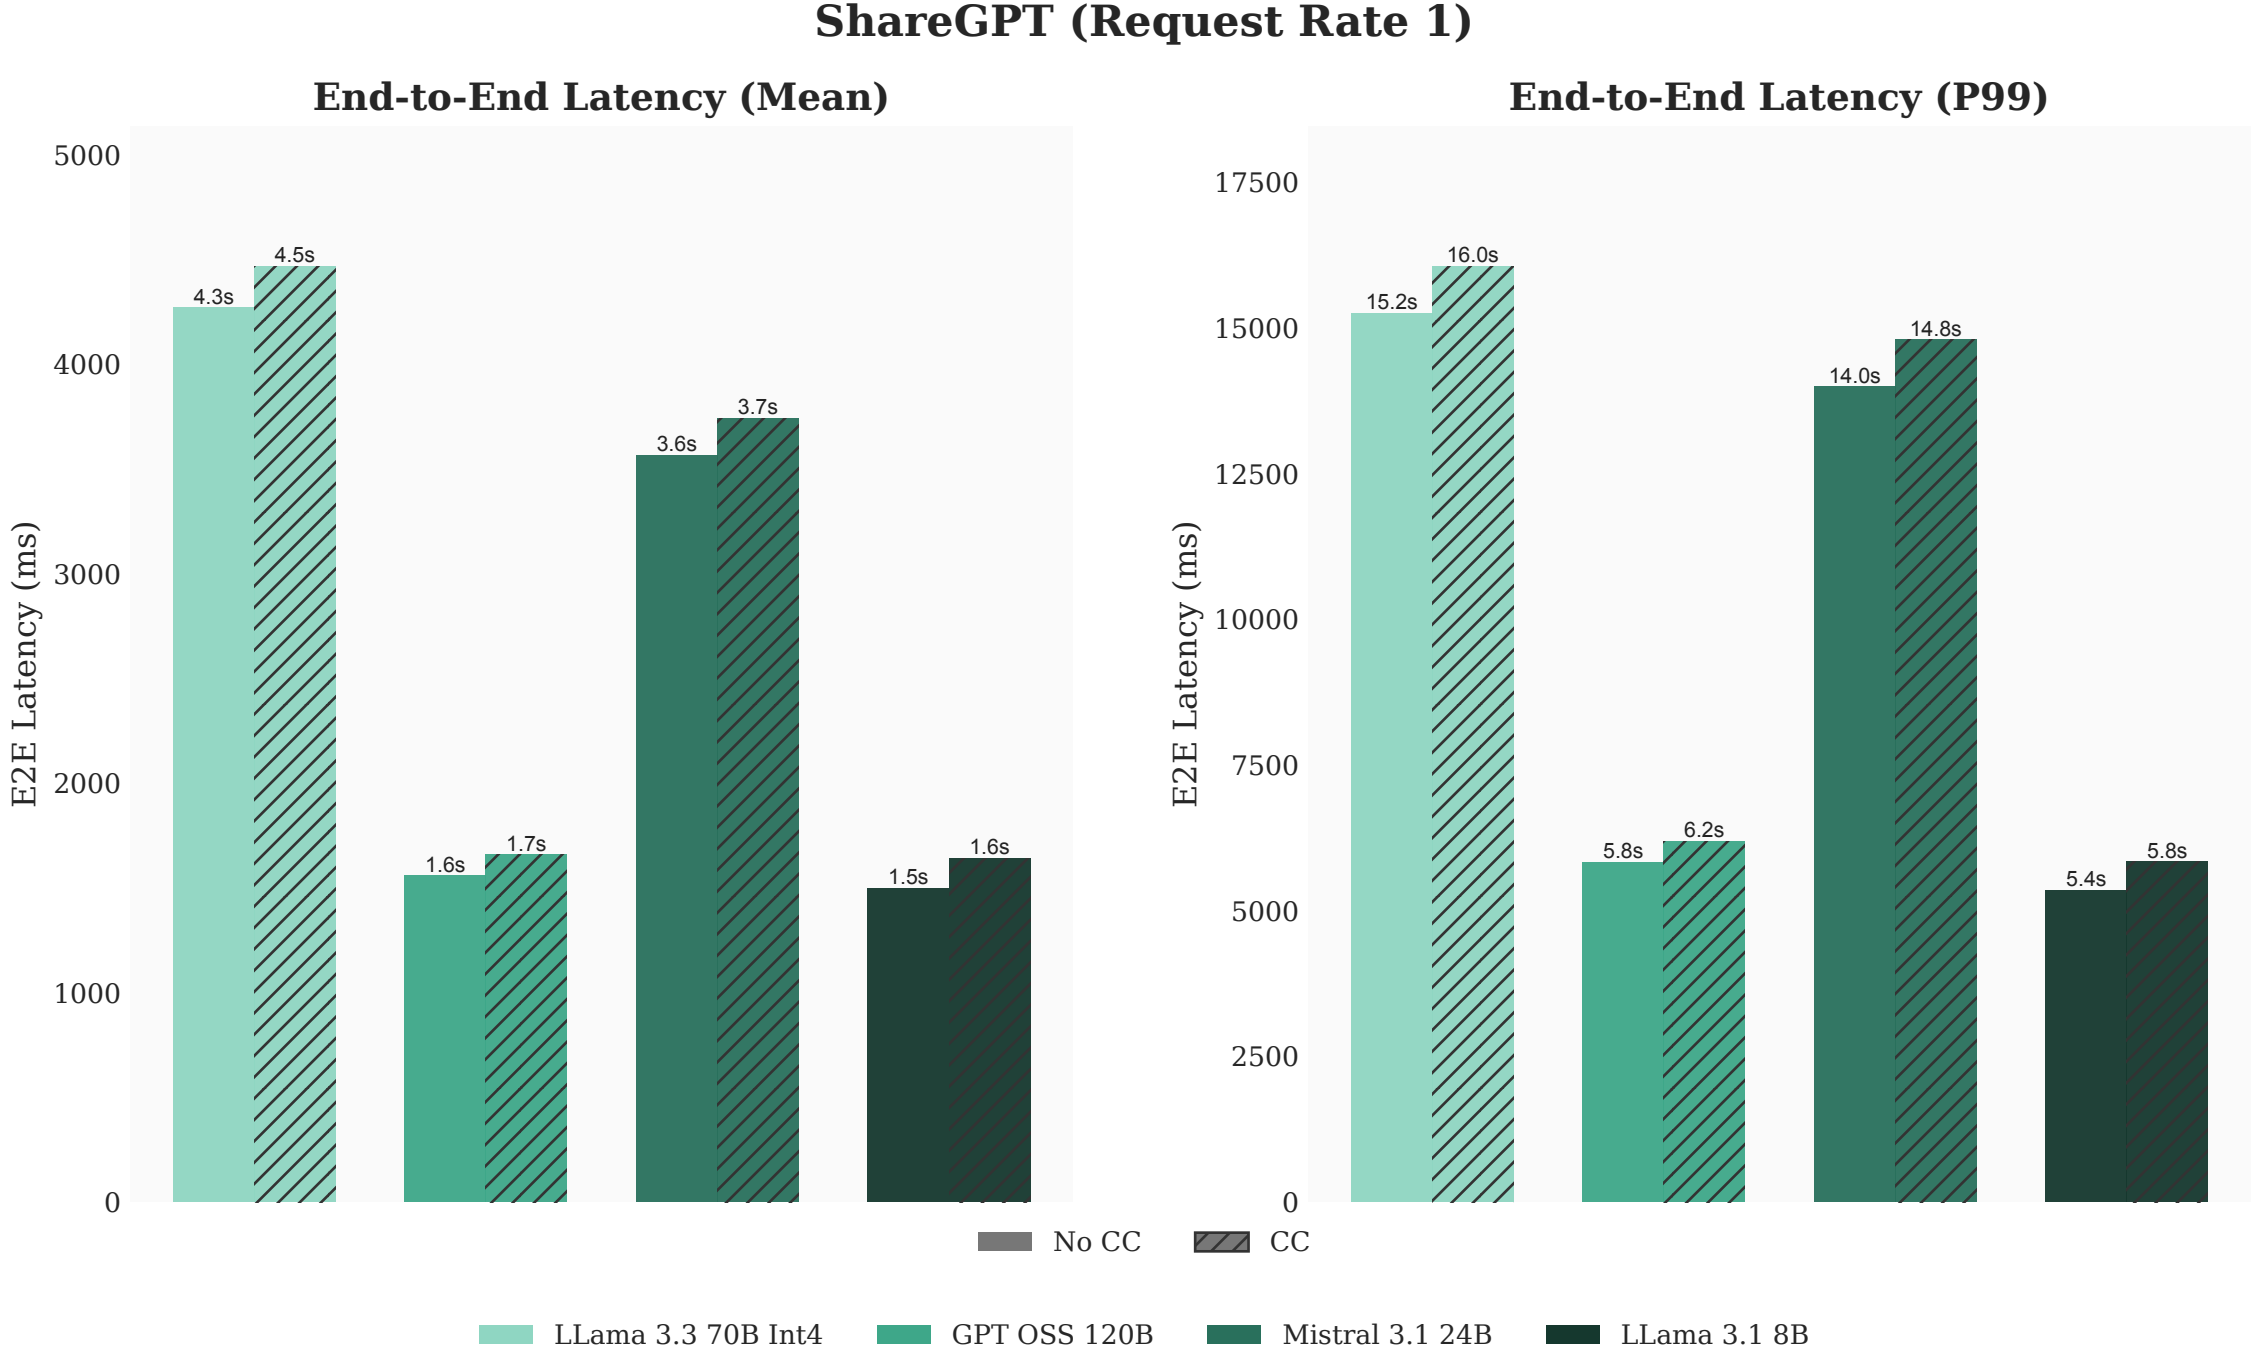
ShareGPT (Request Rate 1)
End-to-End Latency (Mean)
End-to-End Latency (P99)
E2E Latency (ms)
5000
4000
3000
2000
1000
0
17500
15000
12500
10000
7500
5000
2500
4.3s
4.5s
1.6s
1.7s
3.6s
3.7s
1.5s
1.6s
15.2s
16.0s
5.8s
6.2s
14.0s
14.8s
5.4s
5.8s
No CC
CC
LLama 3.3 70B Int4
GPT OSS 120B
Mistral 3.1 24B
LLama 3.1 8B

Edit 10K Characters (Request Rate 100)

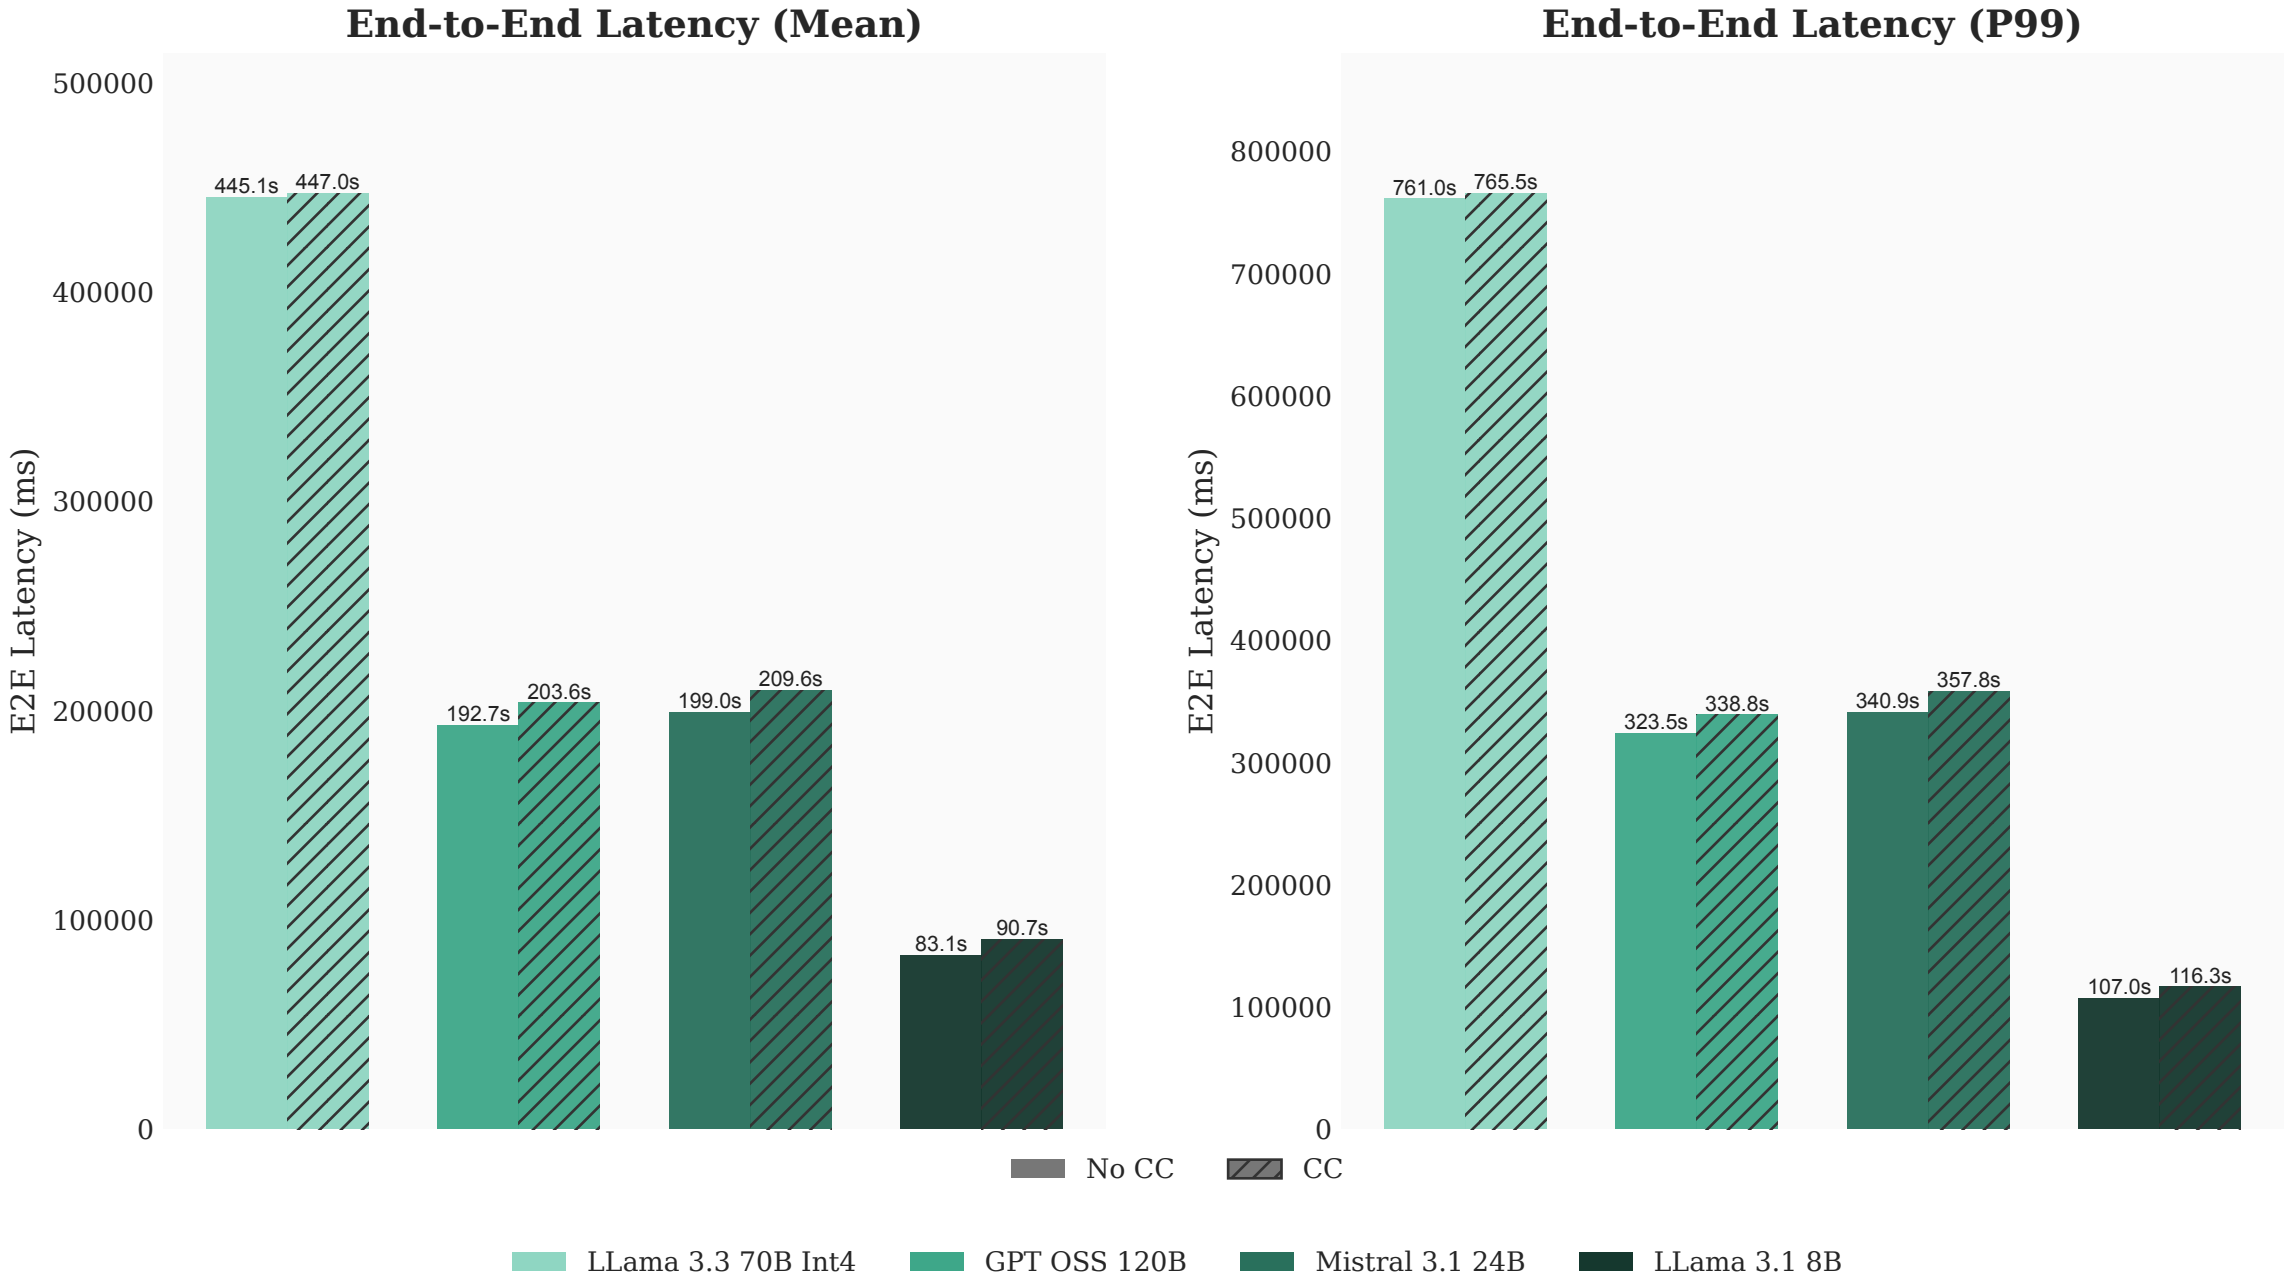

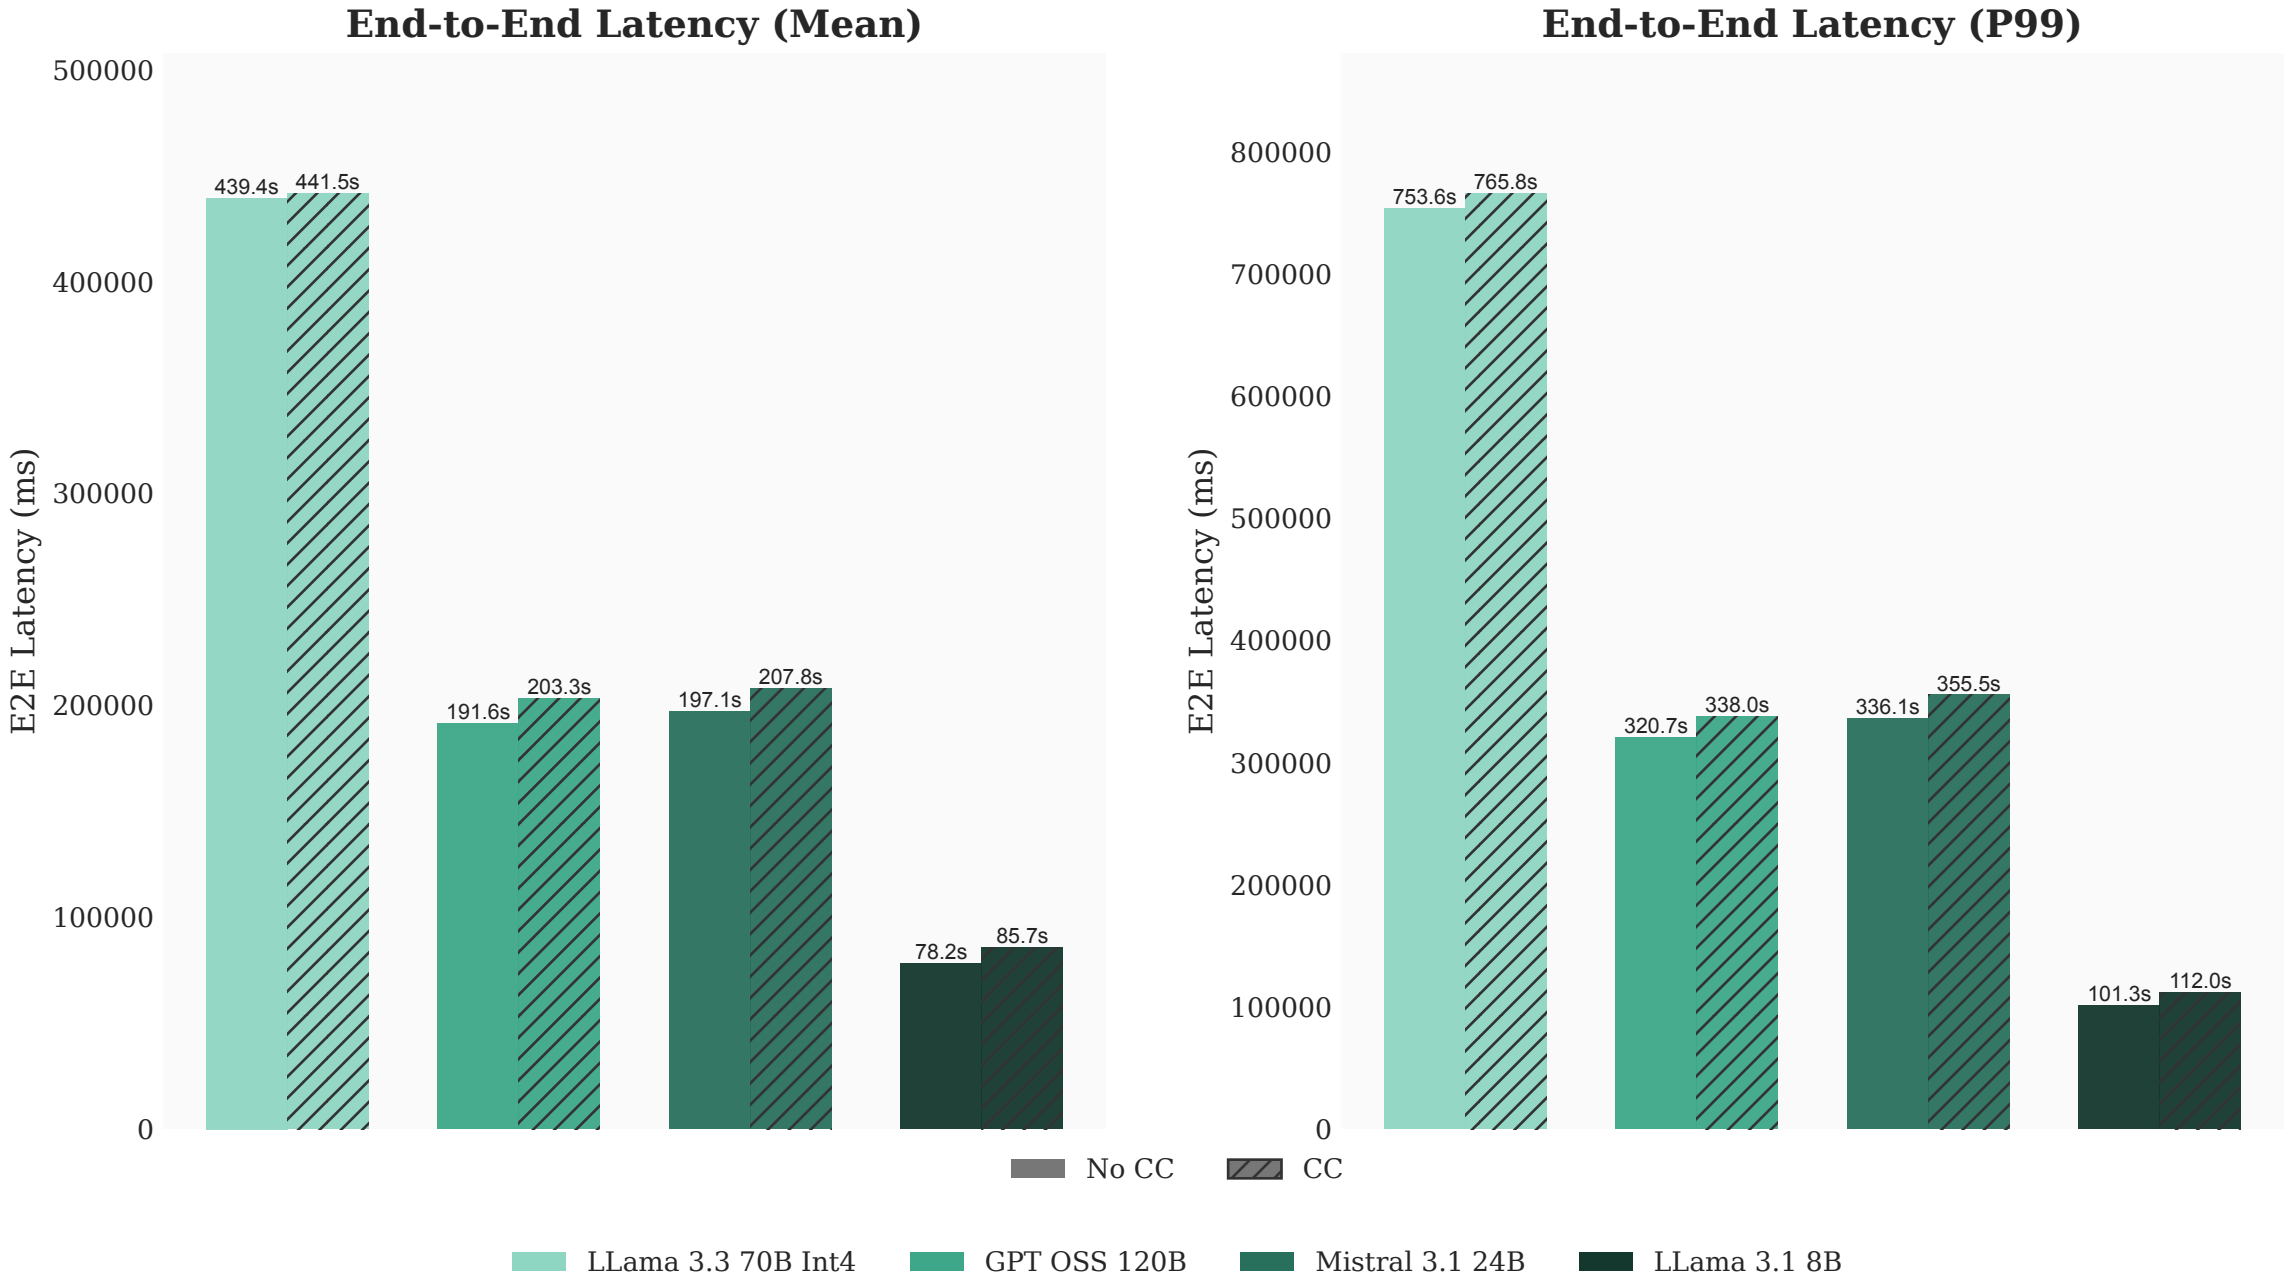
Edit 10K Characters (Request Rate 50)
End-to-End Latency (Mean)
End-to-End Latency (P99)
E2E Latency (ms)
500000
400000
300000
200000
100000
0
800000
700000
600000
500000
439.4s
441.5s
191.6s
203.3s
197.1s
207.8s
78.2s
85.7s
753.6s
765.8s
320.7s
338.0s
336.1s
355.5s
101.3s
112.0s
No CC
CC
LLama 3.3 70B Int4
GPT OSS 120B
Mistral 3.1 24B
LLama 3.1 8B

Edit 10K Characters (Request Rate 1)

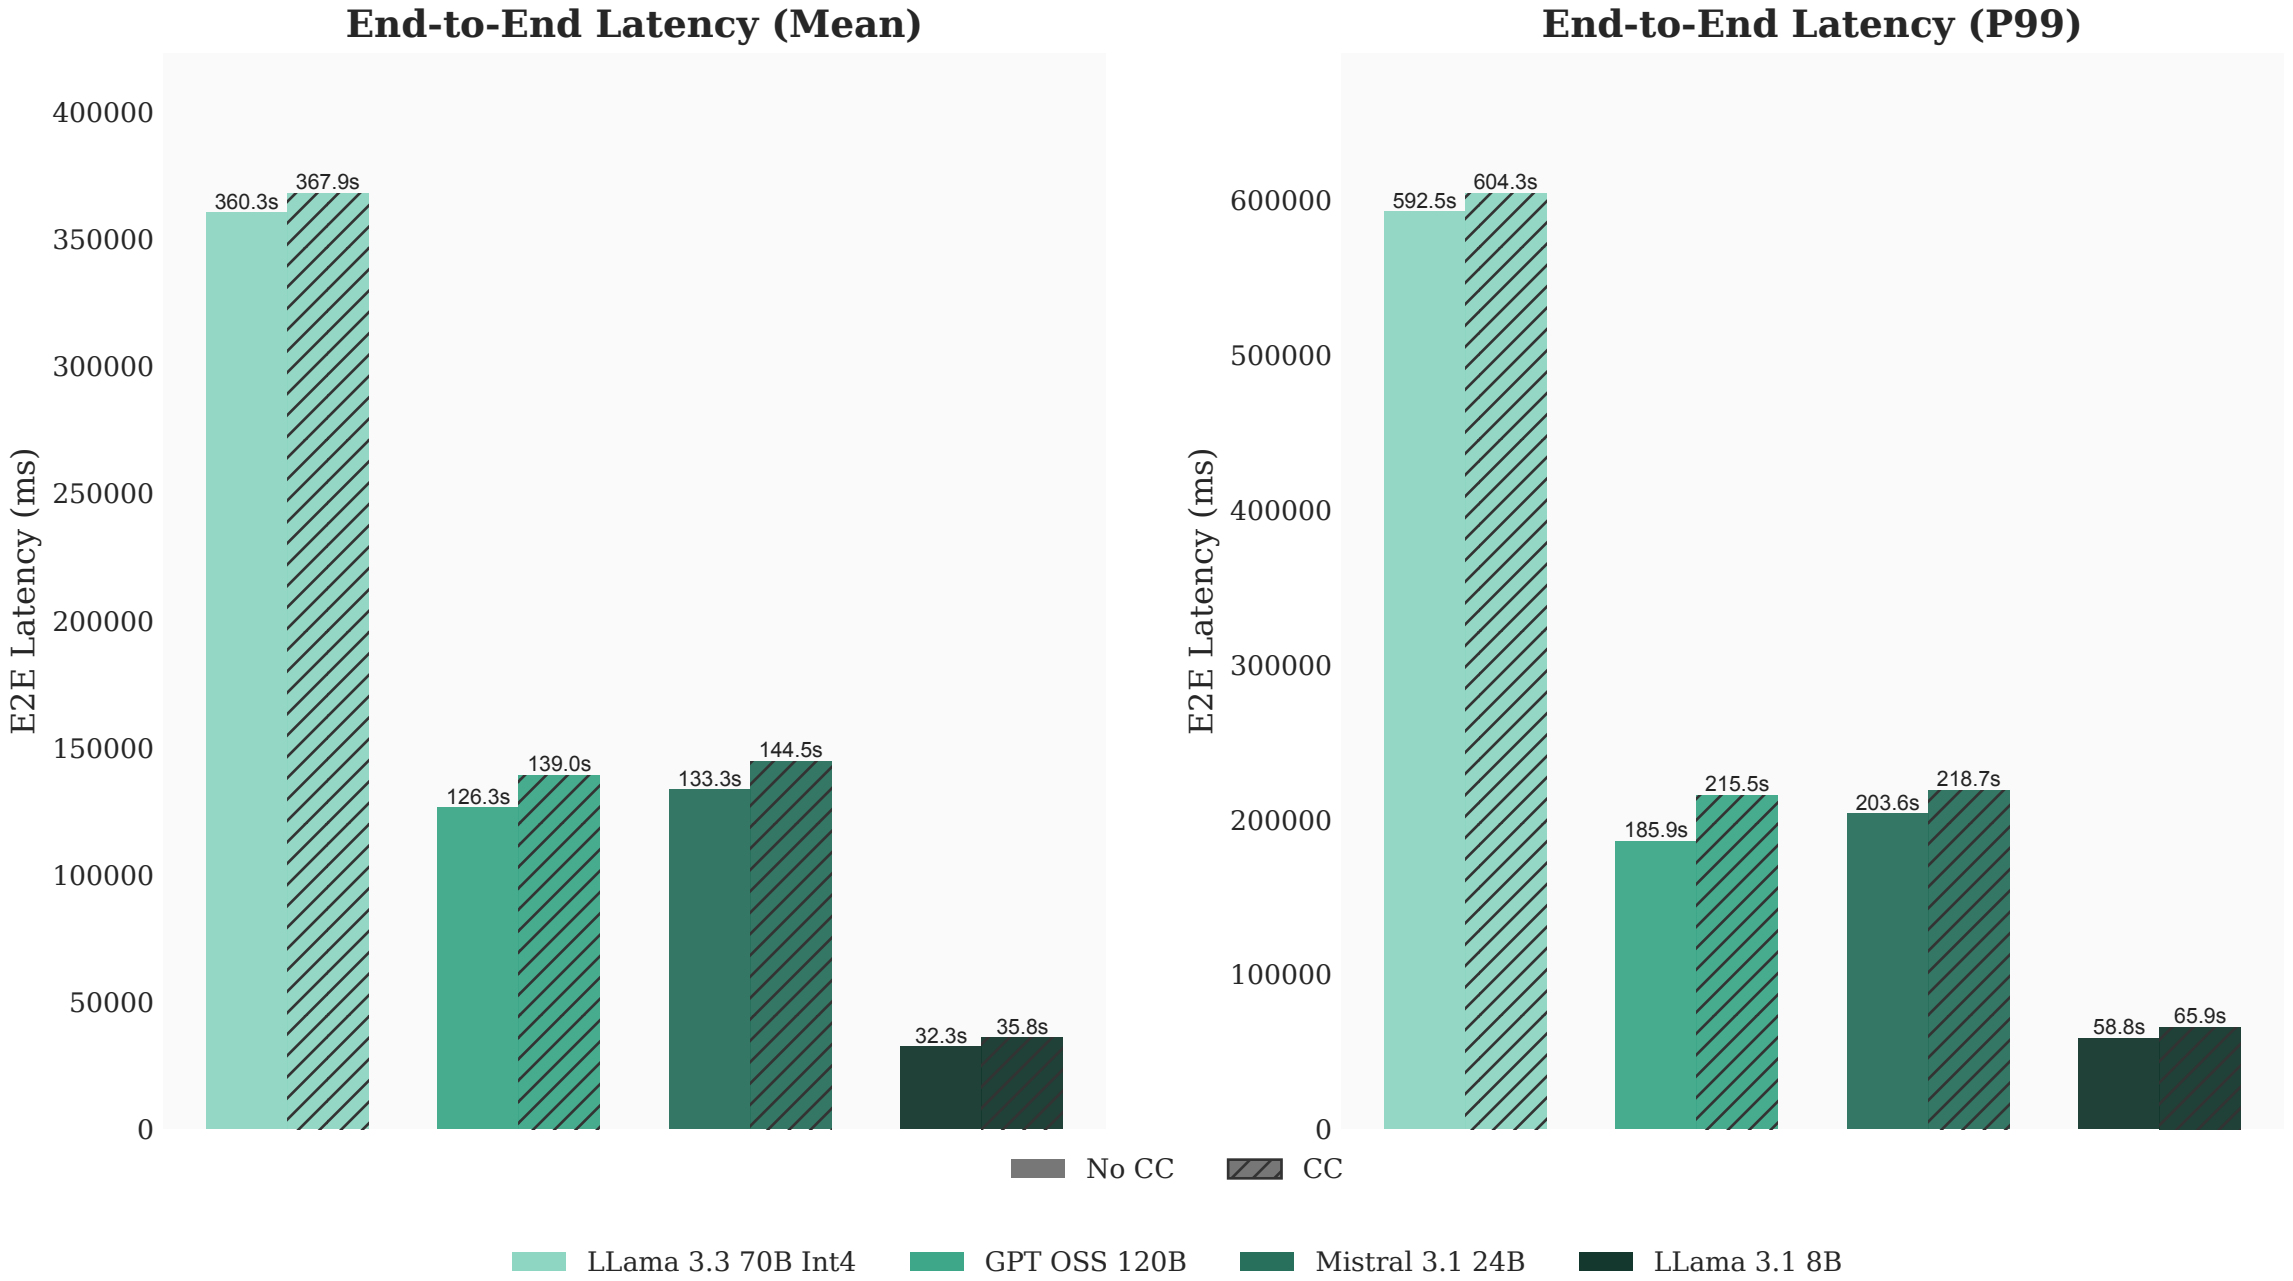

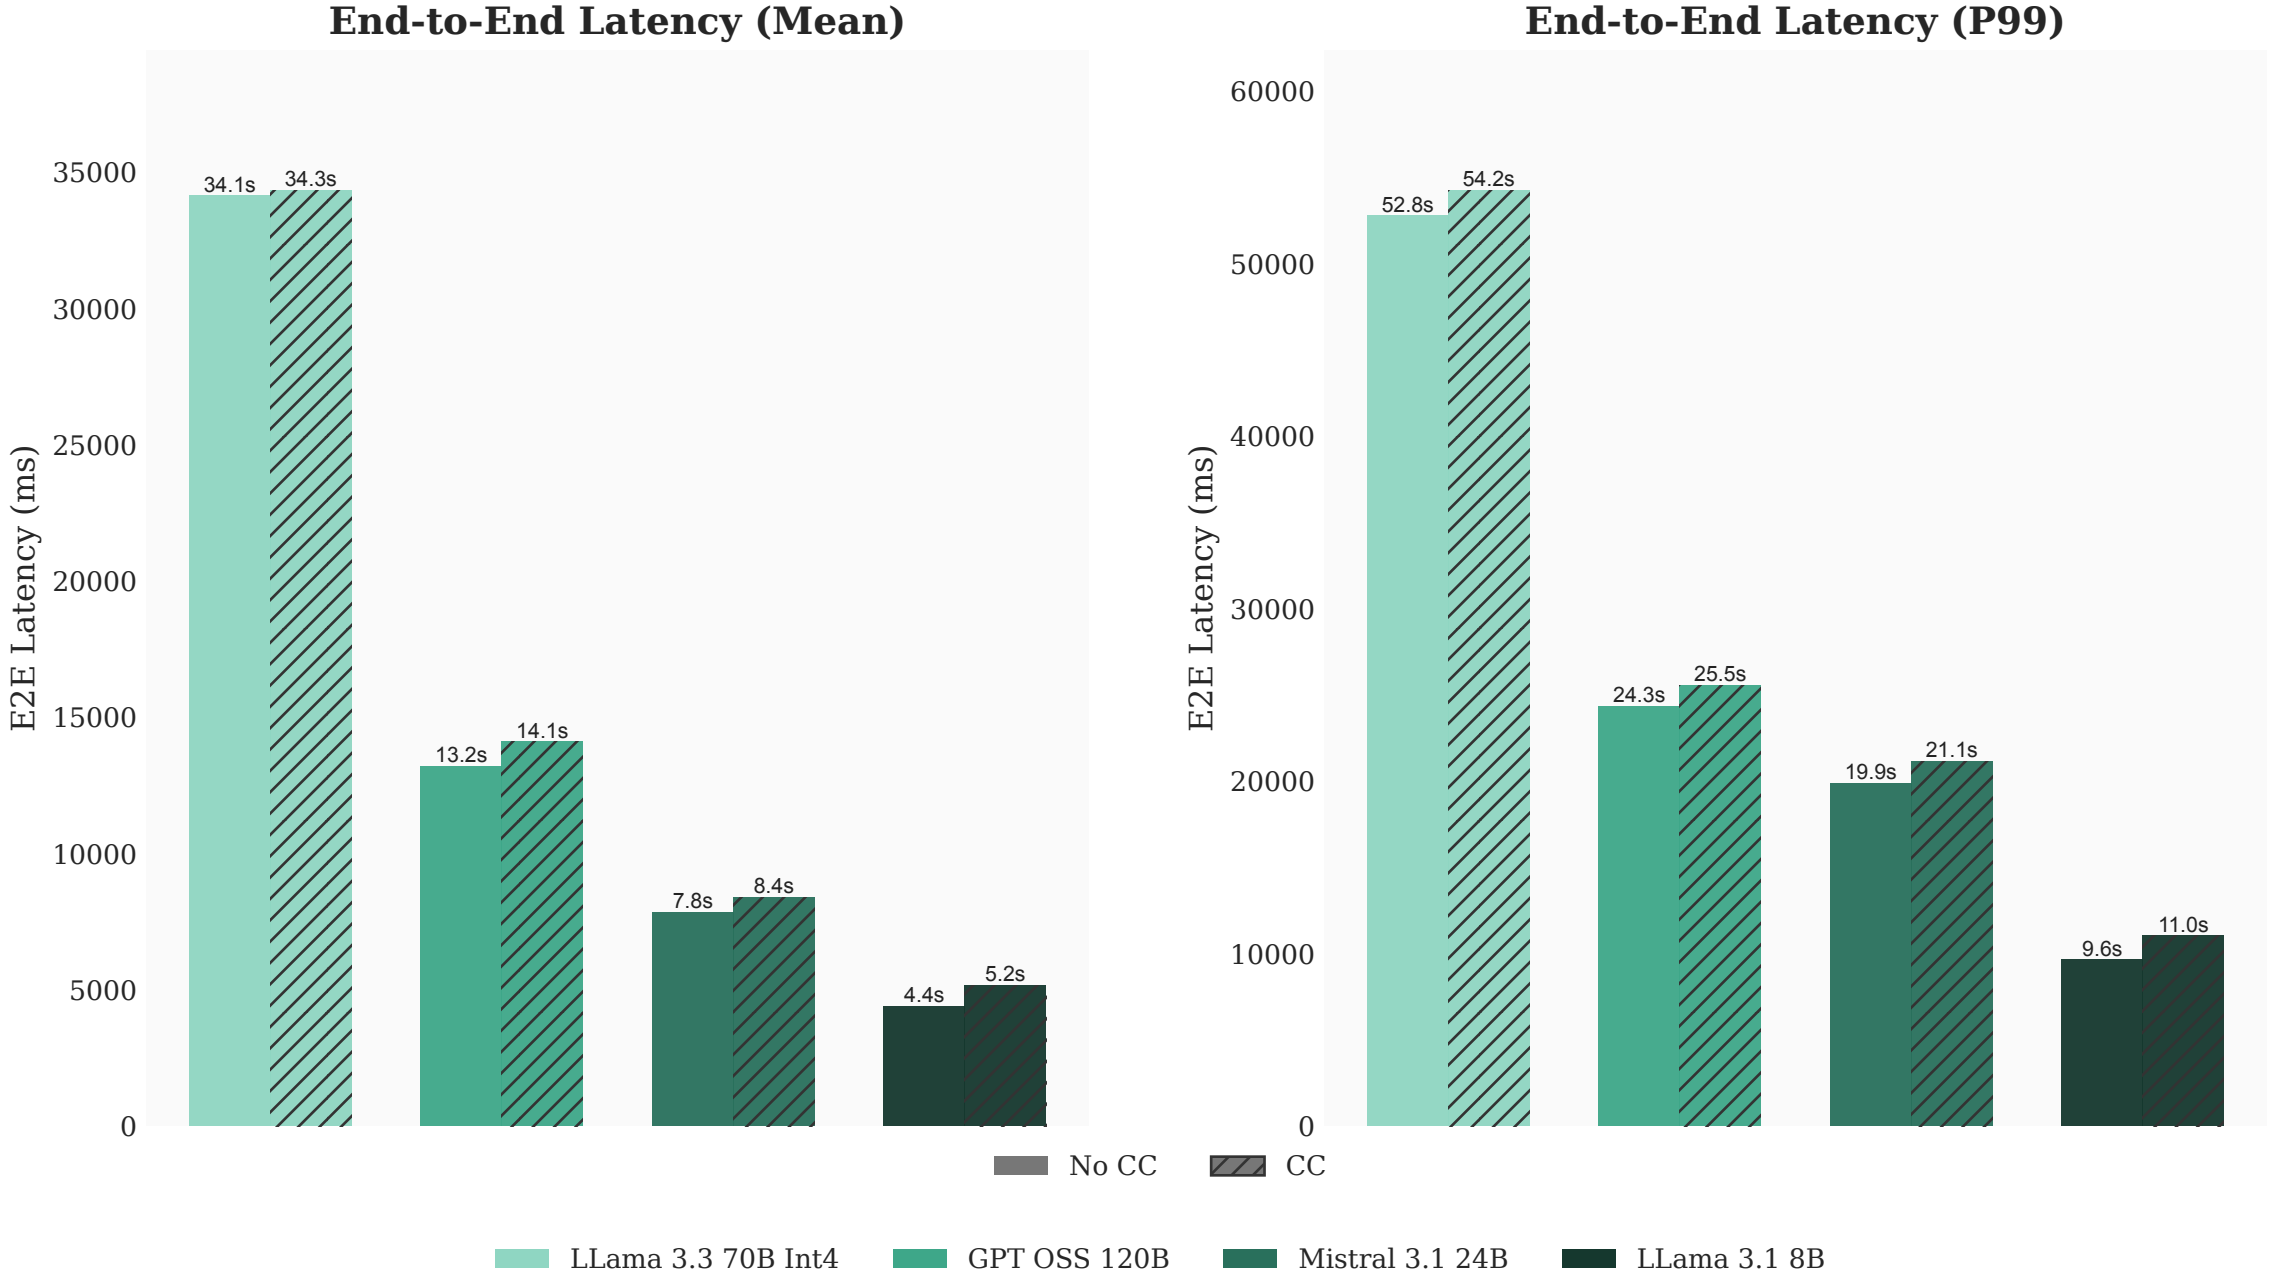
Numina Math (Request Rate 100)
End-to-End Latency (Mean)
End-to-End Latency (P99)
E2E Latency (ms)
0
5000
10000
15000
20000
25000
30000
35000
40000
50000
60000
34.1s
34.3s
13.2s
14.1s
7.8s
8.4s
4.4s
5.2s
52.8s
54.2s
24.3s
25.5s
19.9s
21.1s
9.6s
11.0s
No CC
CC
LLama 3.3 70B Int4
GPT OSS 120B
Mistral 3.1 24B
LLama 3.1 8B

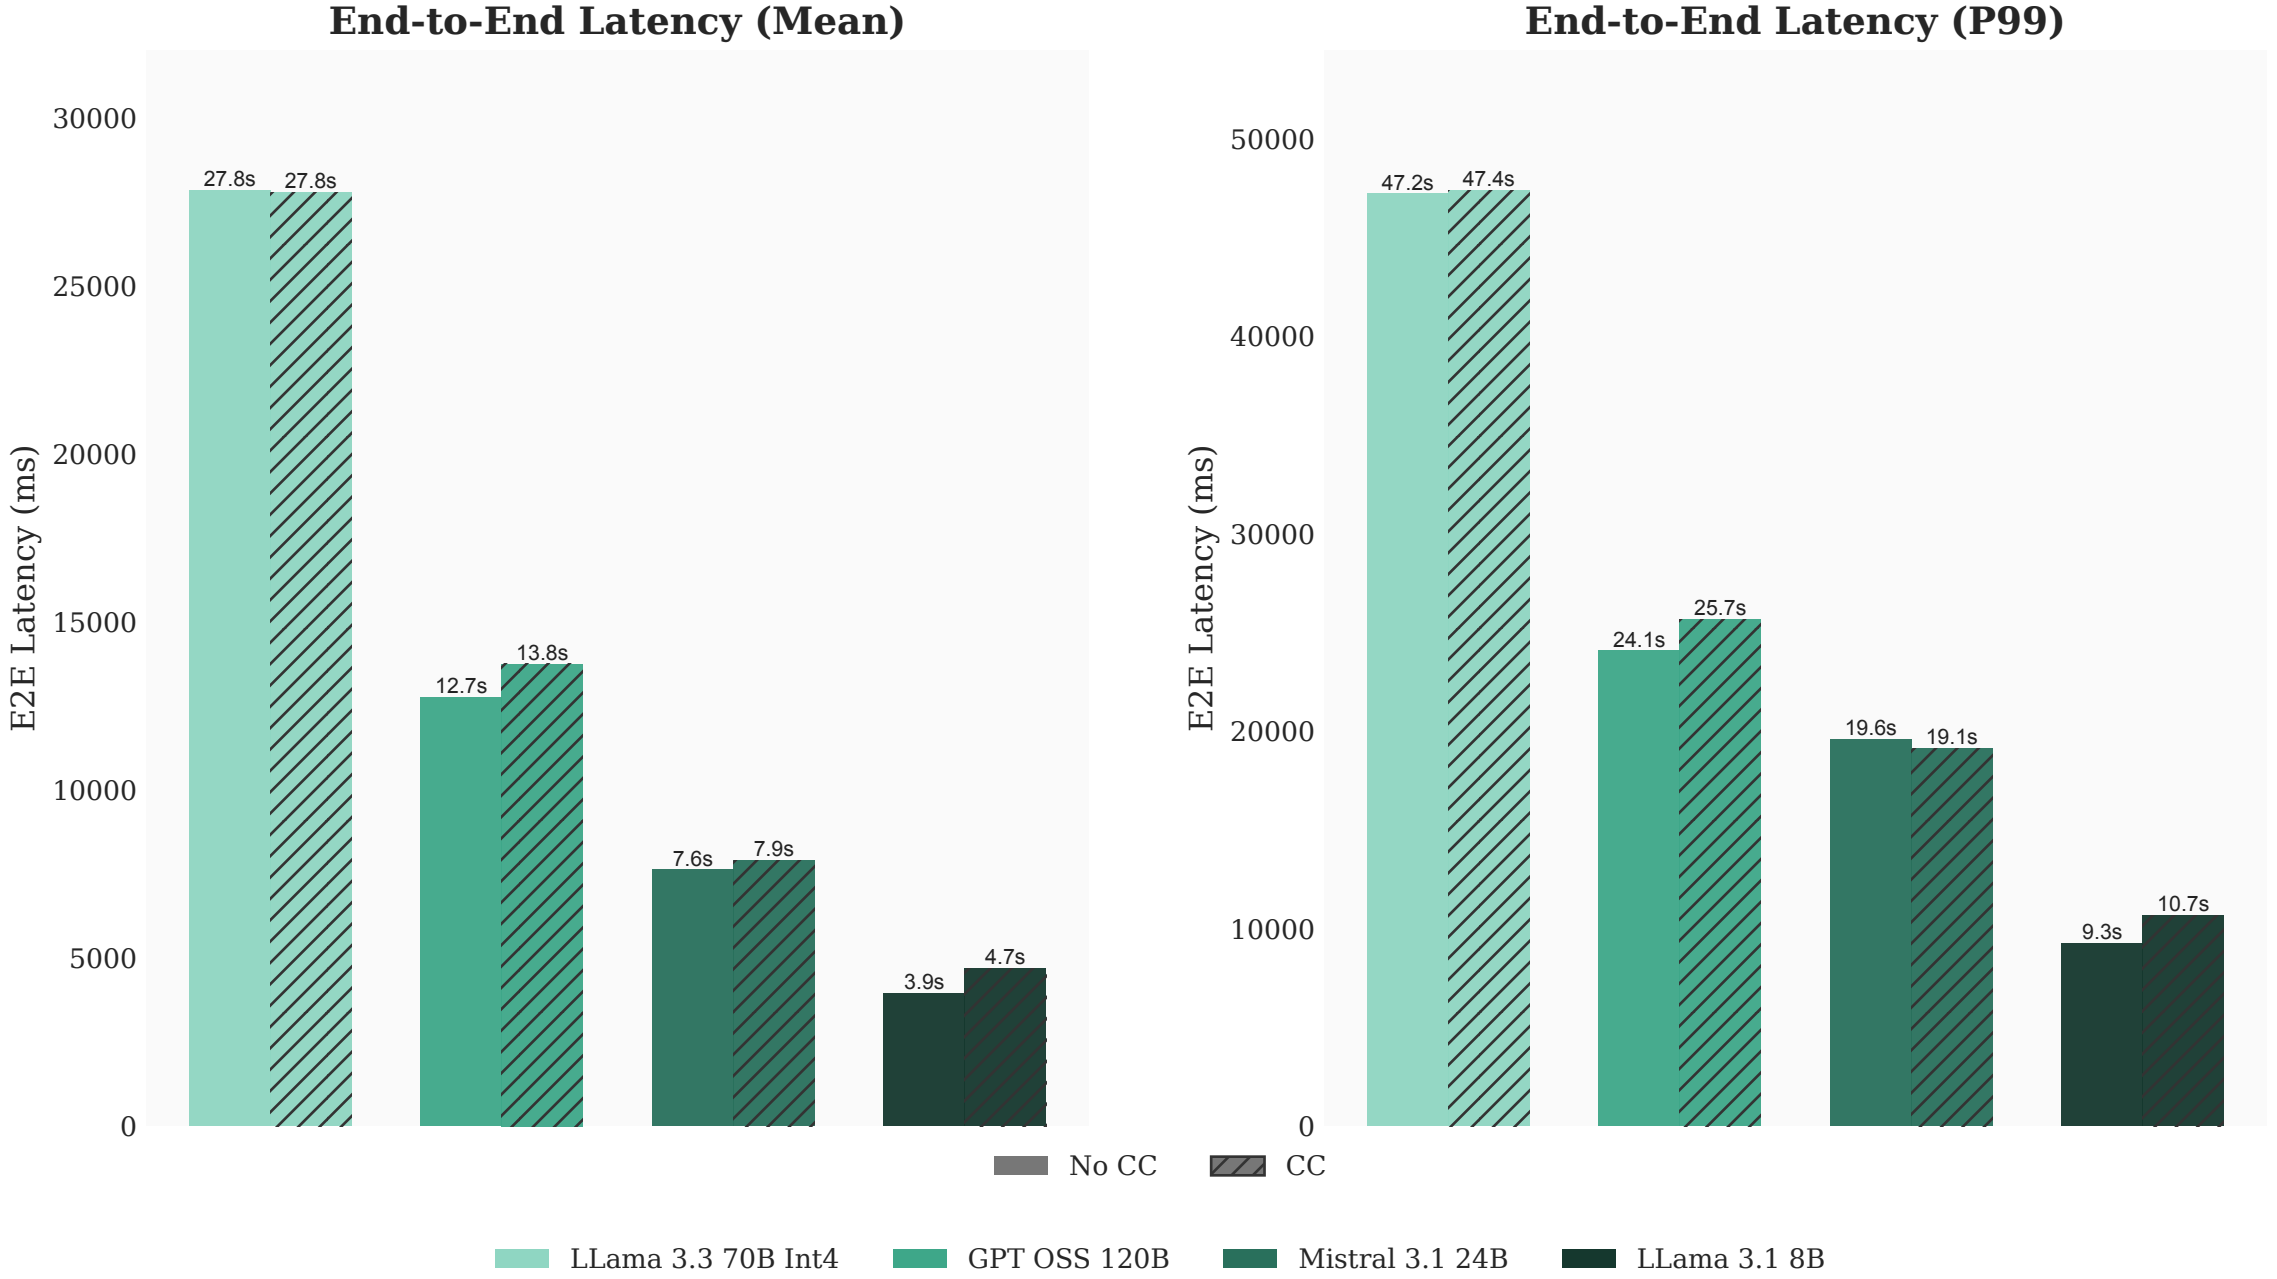
Numina Math (Request Rate 50)
End-to-End Latency (Mean)
End-to-End Latency (P99)
E2E Latency (ms)
30000
25000
20000
15000
10000
5000
0
50000
40000
30000
20000
10000
0
27.8s
27.8s
12.7s
13.8s
7.6s
7.9s
3.9s
4.7s
47.2s
47.4s
24.1s
25.7s
19.6s
19.1s
9.3s
10.7s
No CC
CC
LLama 3.3 70B Int4
GPT OSS 120B
Mistral 3.1 24B
LLama 3.1 8B

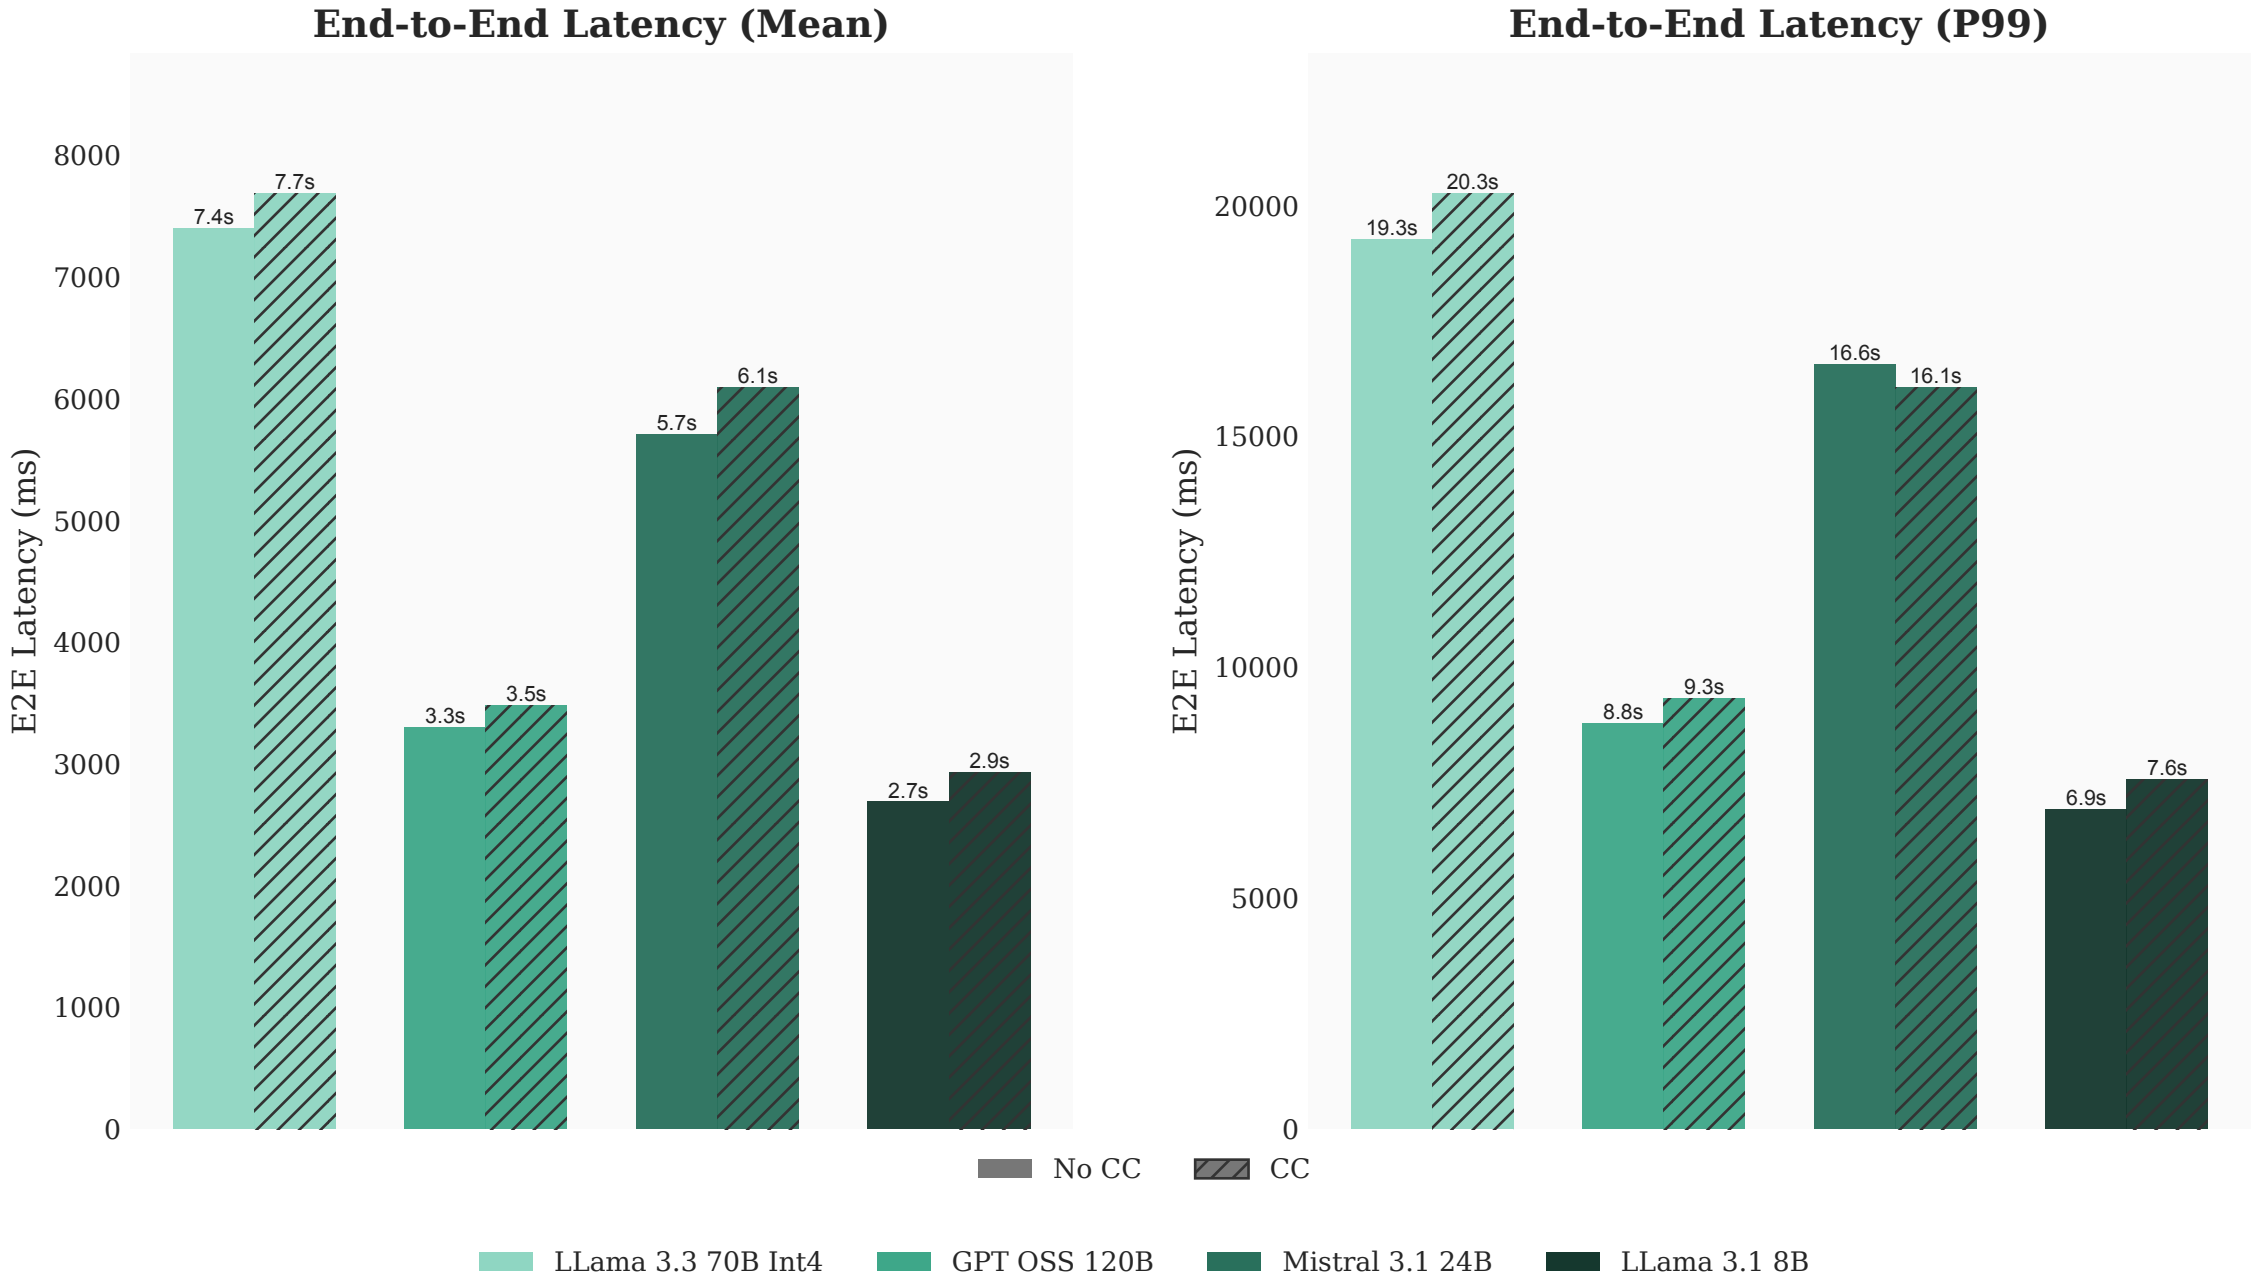
Numina Math (Request Rate 1)
End-to-End Latency (Mean)
End-to-End Latency (P99)
E2E Latency (ms)
0
1000
2000
3000
4000
5000
6000
7000
8000
10000
15000
20000
7.4s
7.7s
3.3s
3.5s
5.7s
6.1s
2.7s
2.9s
19.3s
20.3s
8.8s
9.3s
16.6s
16.1s
6.9s
7.6s
No CC
CC
LLama 3.3 70B Int4
GPT OSS 120B
Mistral 3.1 24B
LLama 3.1 8B