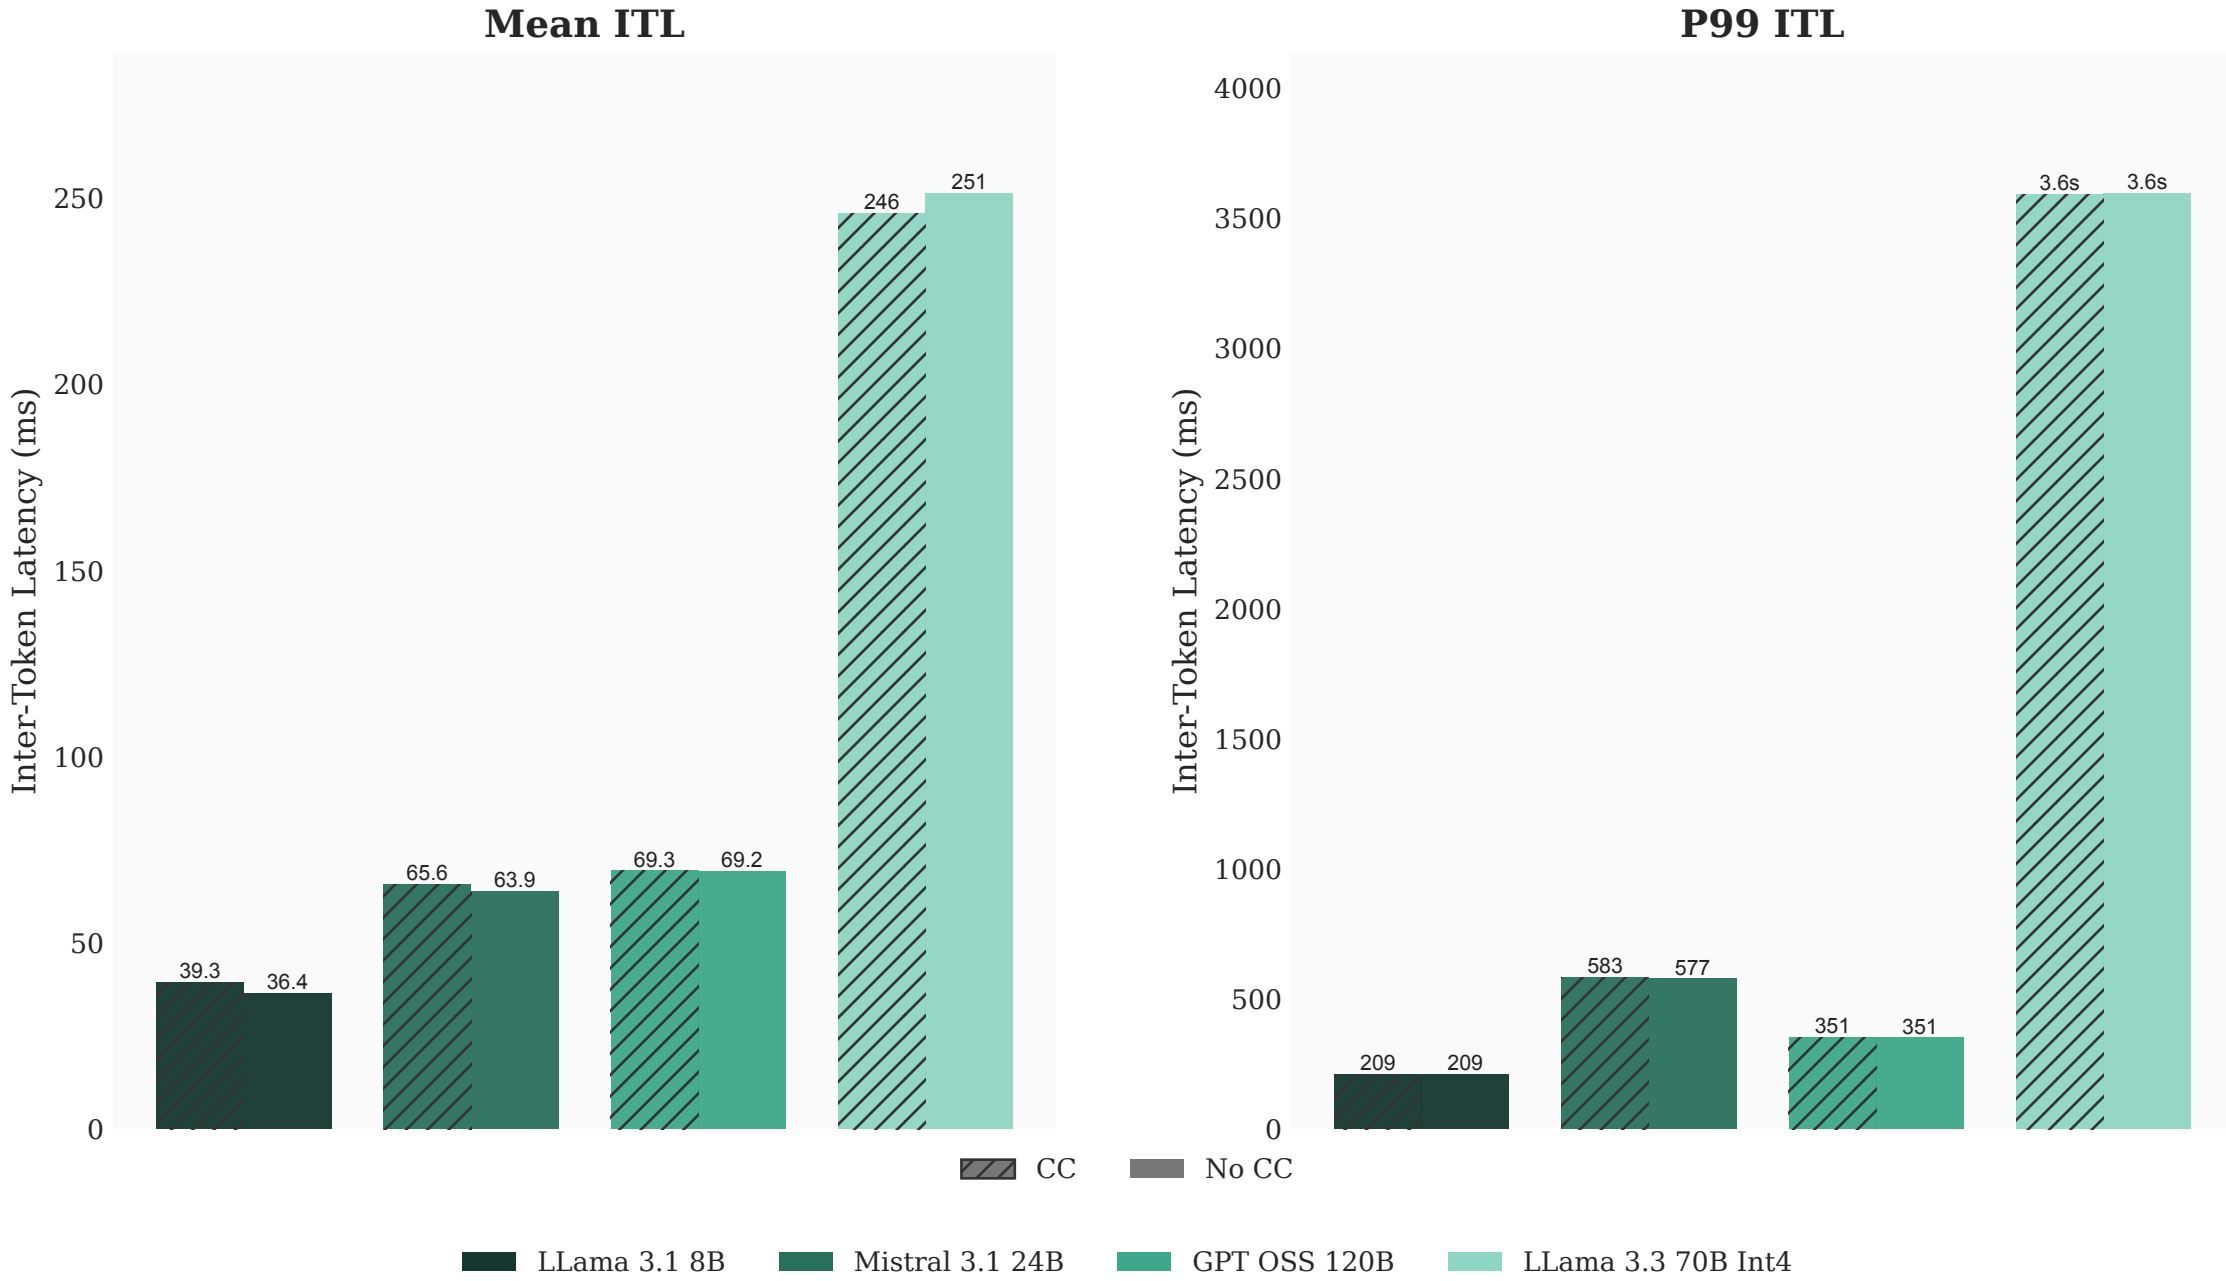
Random (1500 ⇒ 250) (100 Request Rate)
Mean ITL
P99 ITL
Inter-Token Latency (ms)
Inter-Token Latency (ms)
0
50
100
150
200
250
0
500
1000
1500
2000
2500
3000
3500
4000
39.3
36.4
65.6
63.9
69.3
69.2
246
251
209
209
583
577
351
351
3.6s
3.6s
CC
No CC
LLama 3.1 8B
Mistral 3.1 24B
GPT OSS 120B
LLama 3.3 70B Int4

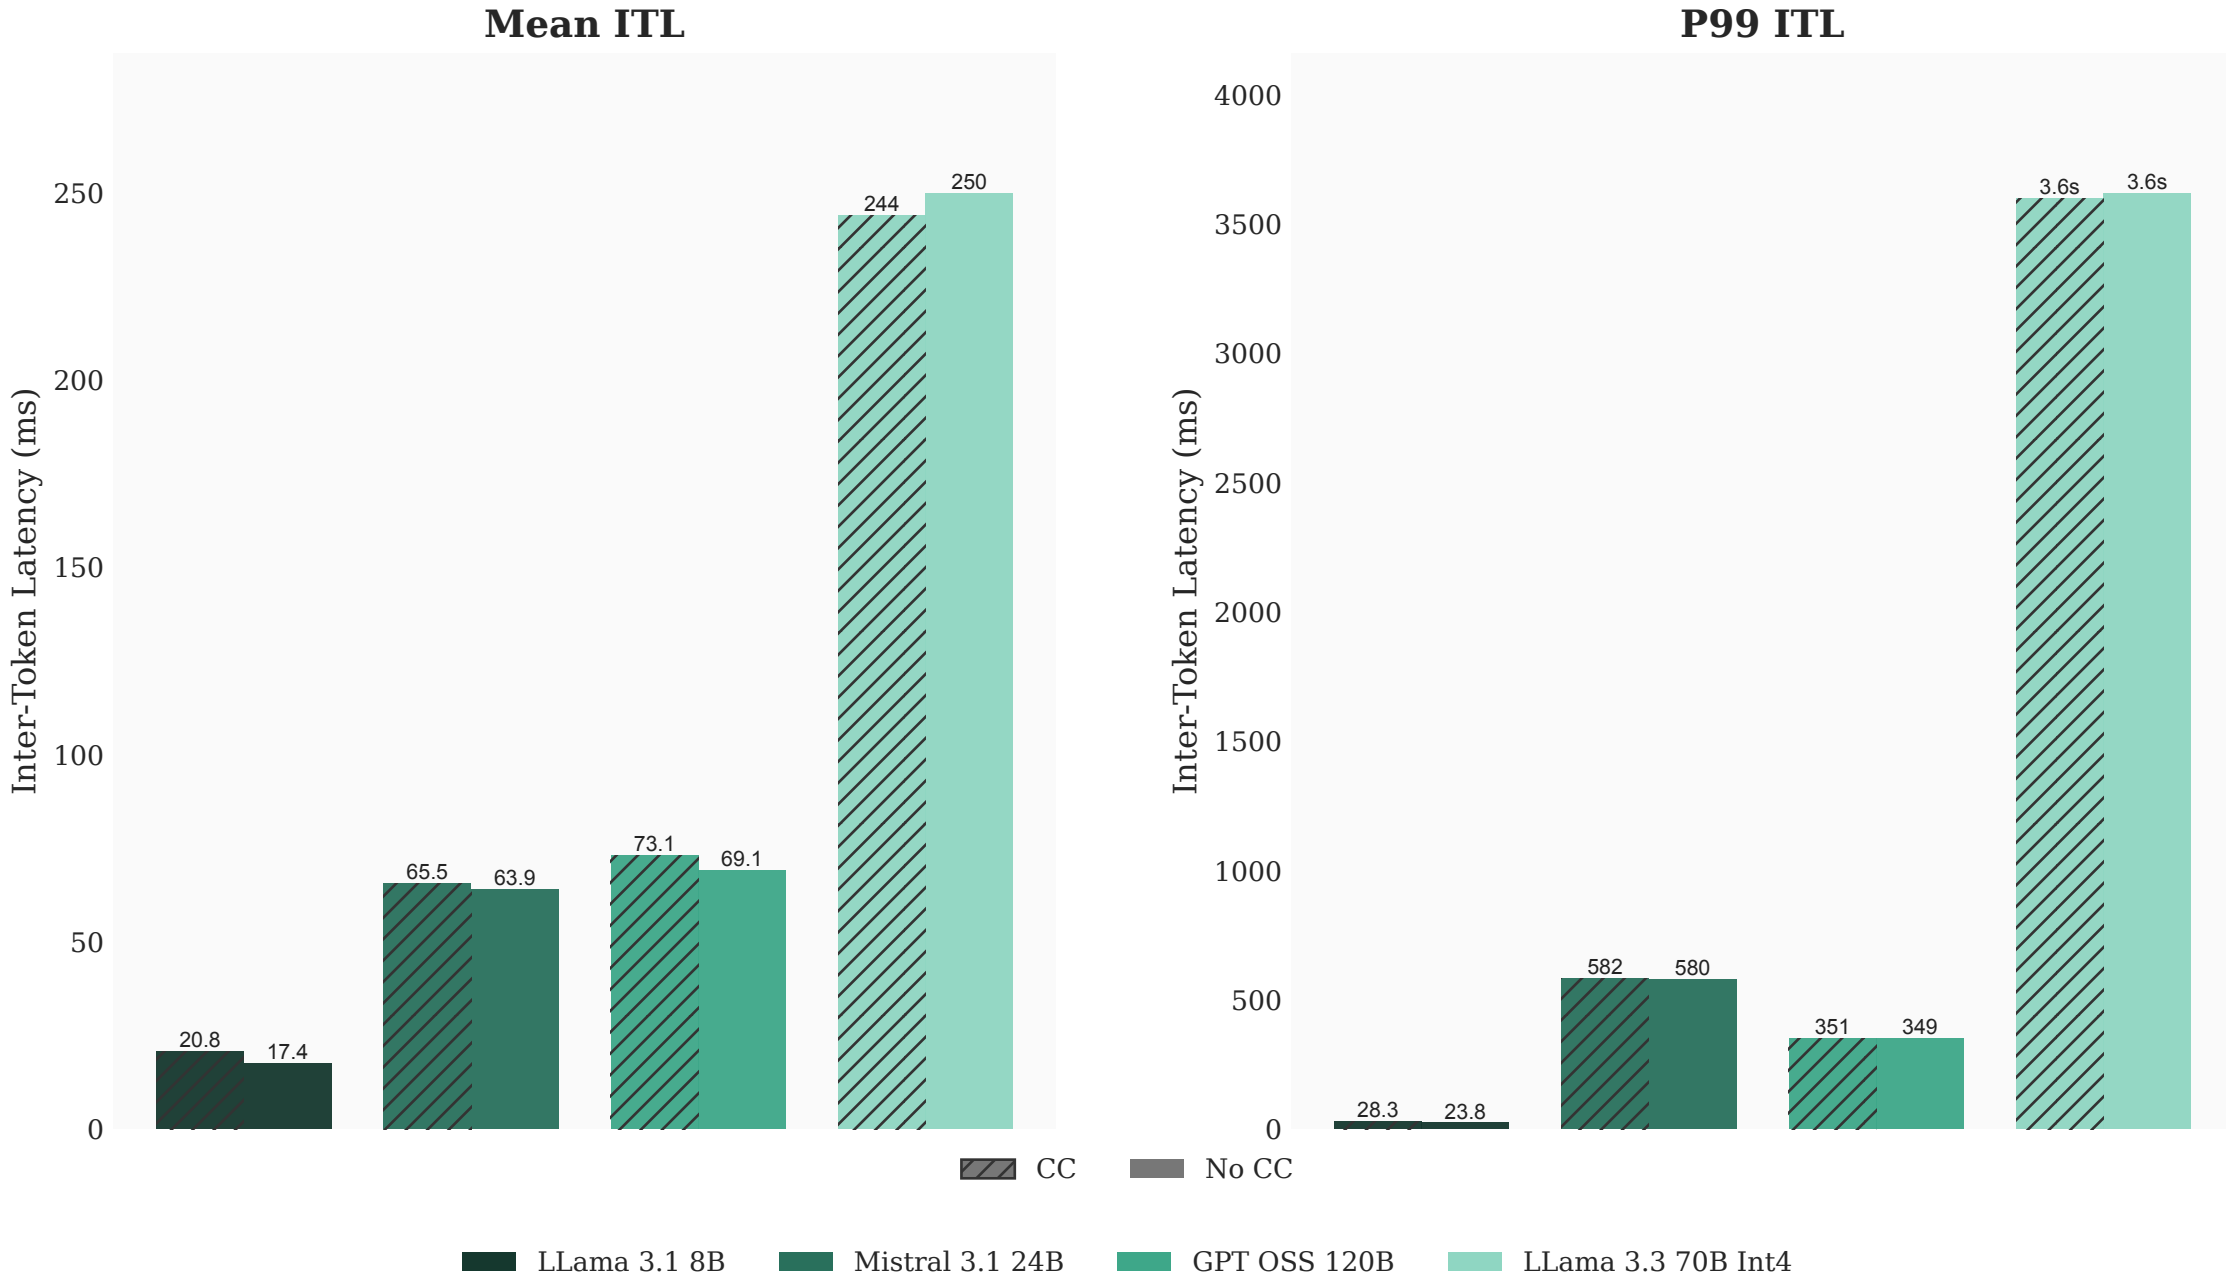
Random (1500 ⇒ 250) (50 Request Rate)
Mean ITL
P99 ITL
Inter-Token Latency (ms)
Inter-Token Latency (ms)
0
50
100
150
200
250
0
500
1000
1500
2000
2500
3000
3500
4000
20.8
17.4
65.5
63.9
73.1
69.1
244
250
28.3
23.8
582
580
351
349
3.6s
3.6s
CC
No CC
LLama 3.1 8B
Mistral 3.1 24B
GPT OSS 120B
LLama 3.3 70B Int4

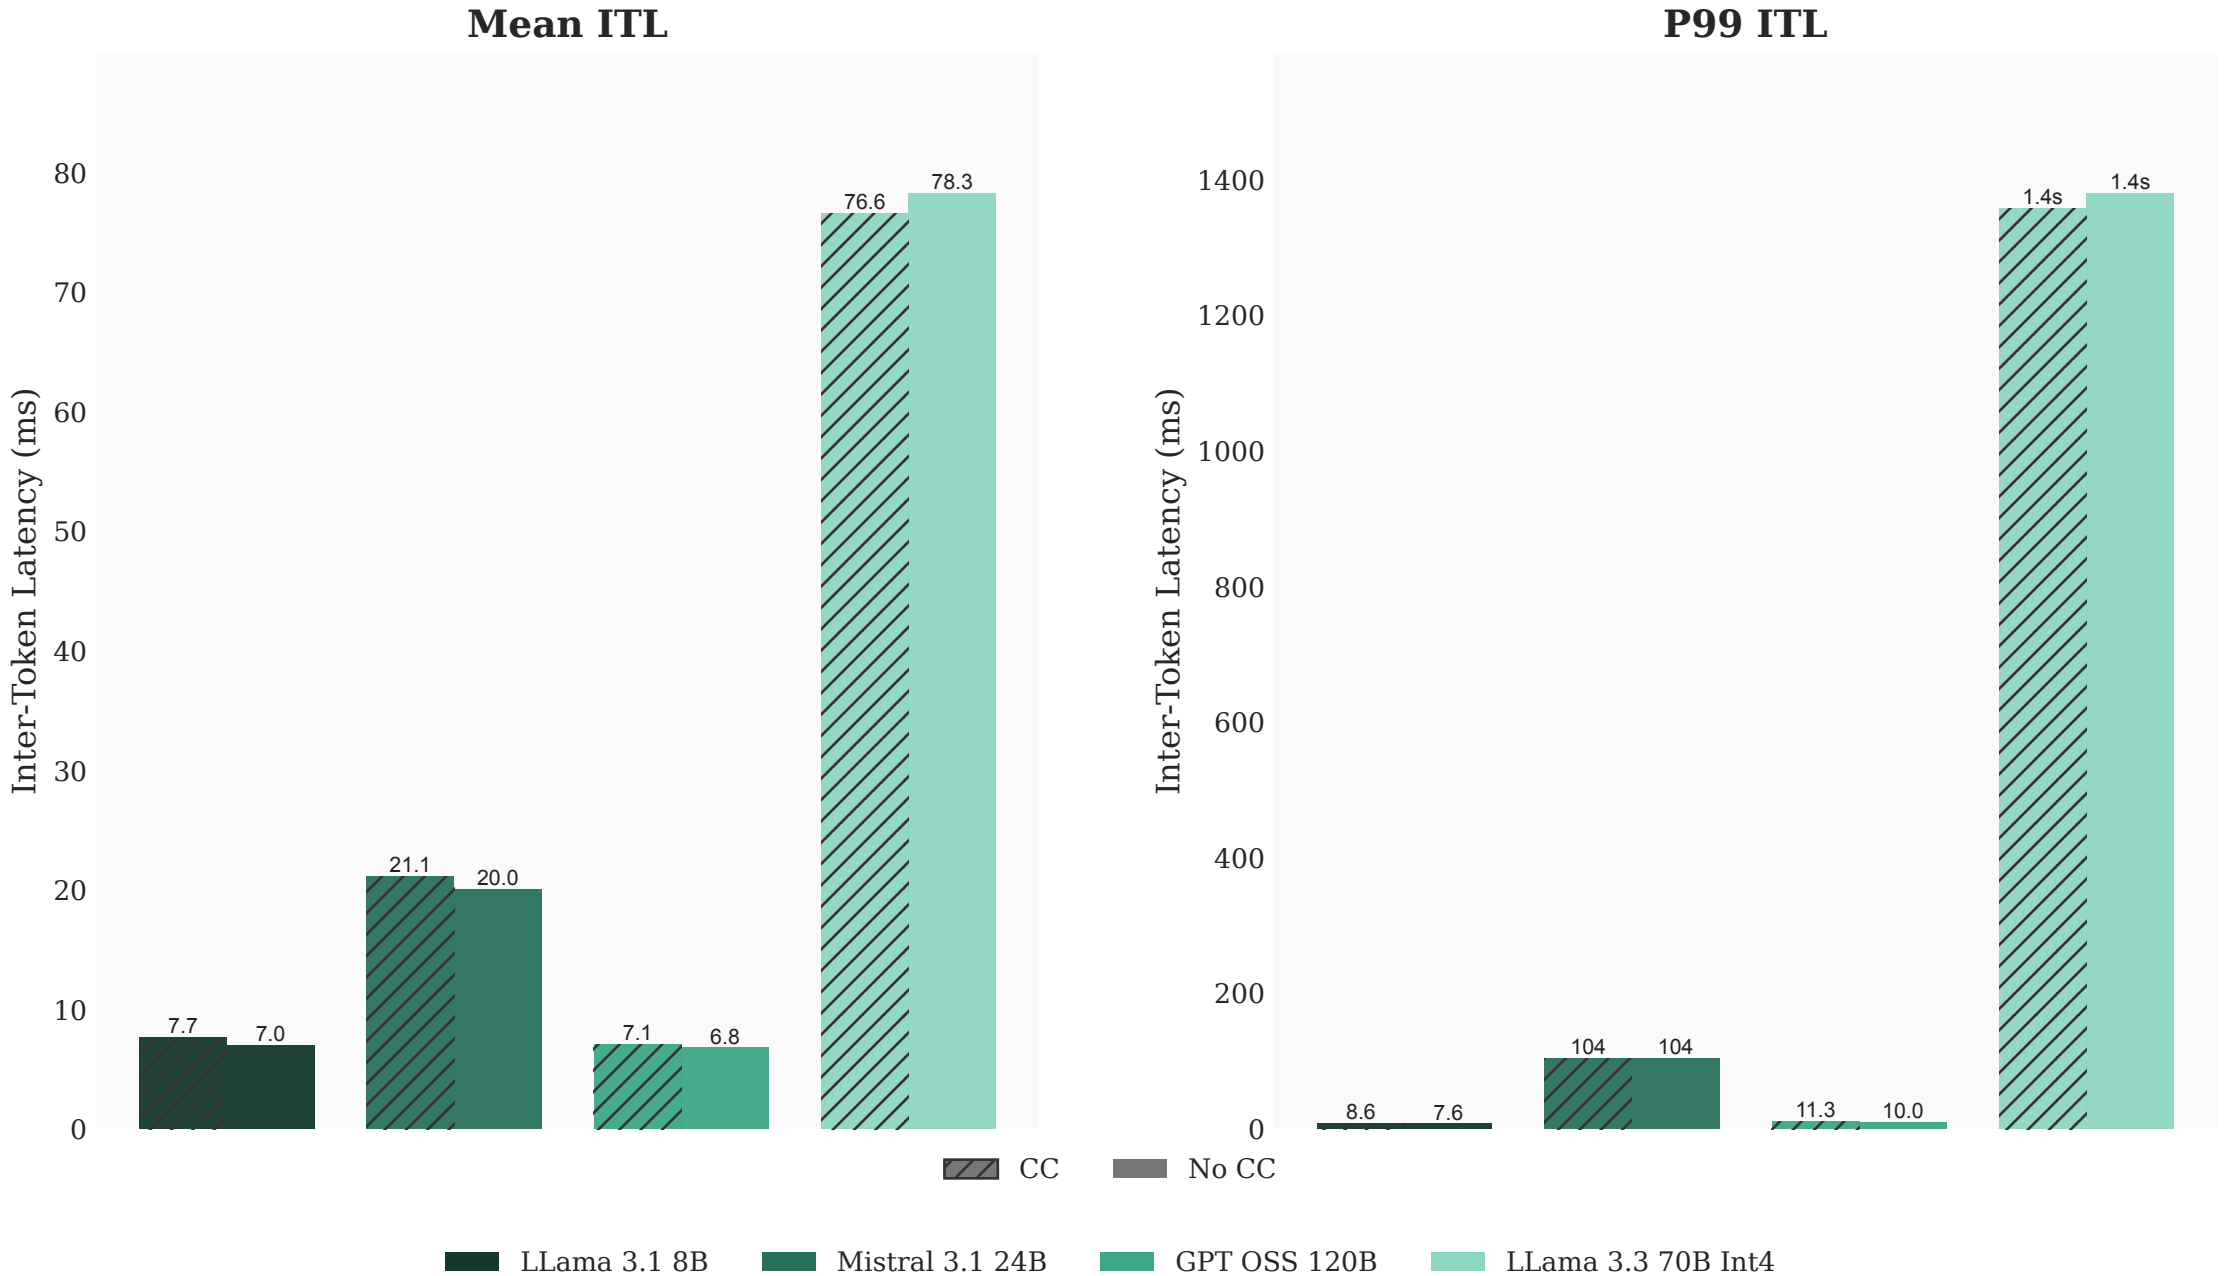
Random (1500 ⇒ 250) (Single Request)
Mean ITL
P99 ITL
Inter-Token Latency (ms)
0
10
20
30
40
50
60
70
80
200
400
600
800
1000
1200
1400
7.7
7.0
21.1
20.0
7.1
6.8
76.6
78.3
8.6
7.6
104
104
11.3
10.0
1.4s
1.4s
CC
No CC
LLama 3.1 8B
Mistral 3.1 24B
GPT OSS 120B
LLama 3.3 70B Int4

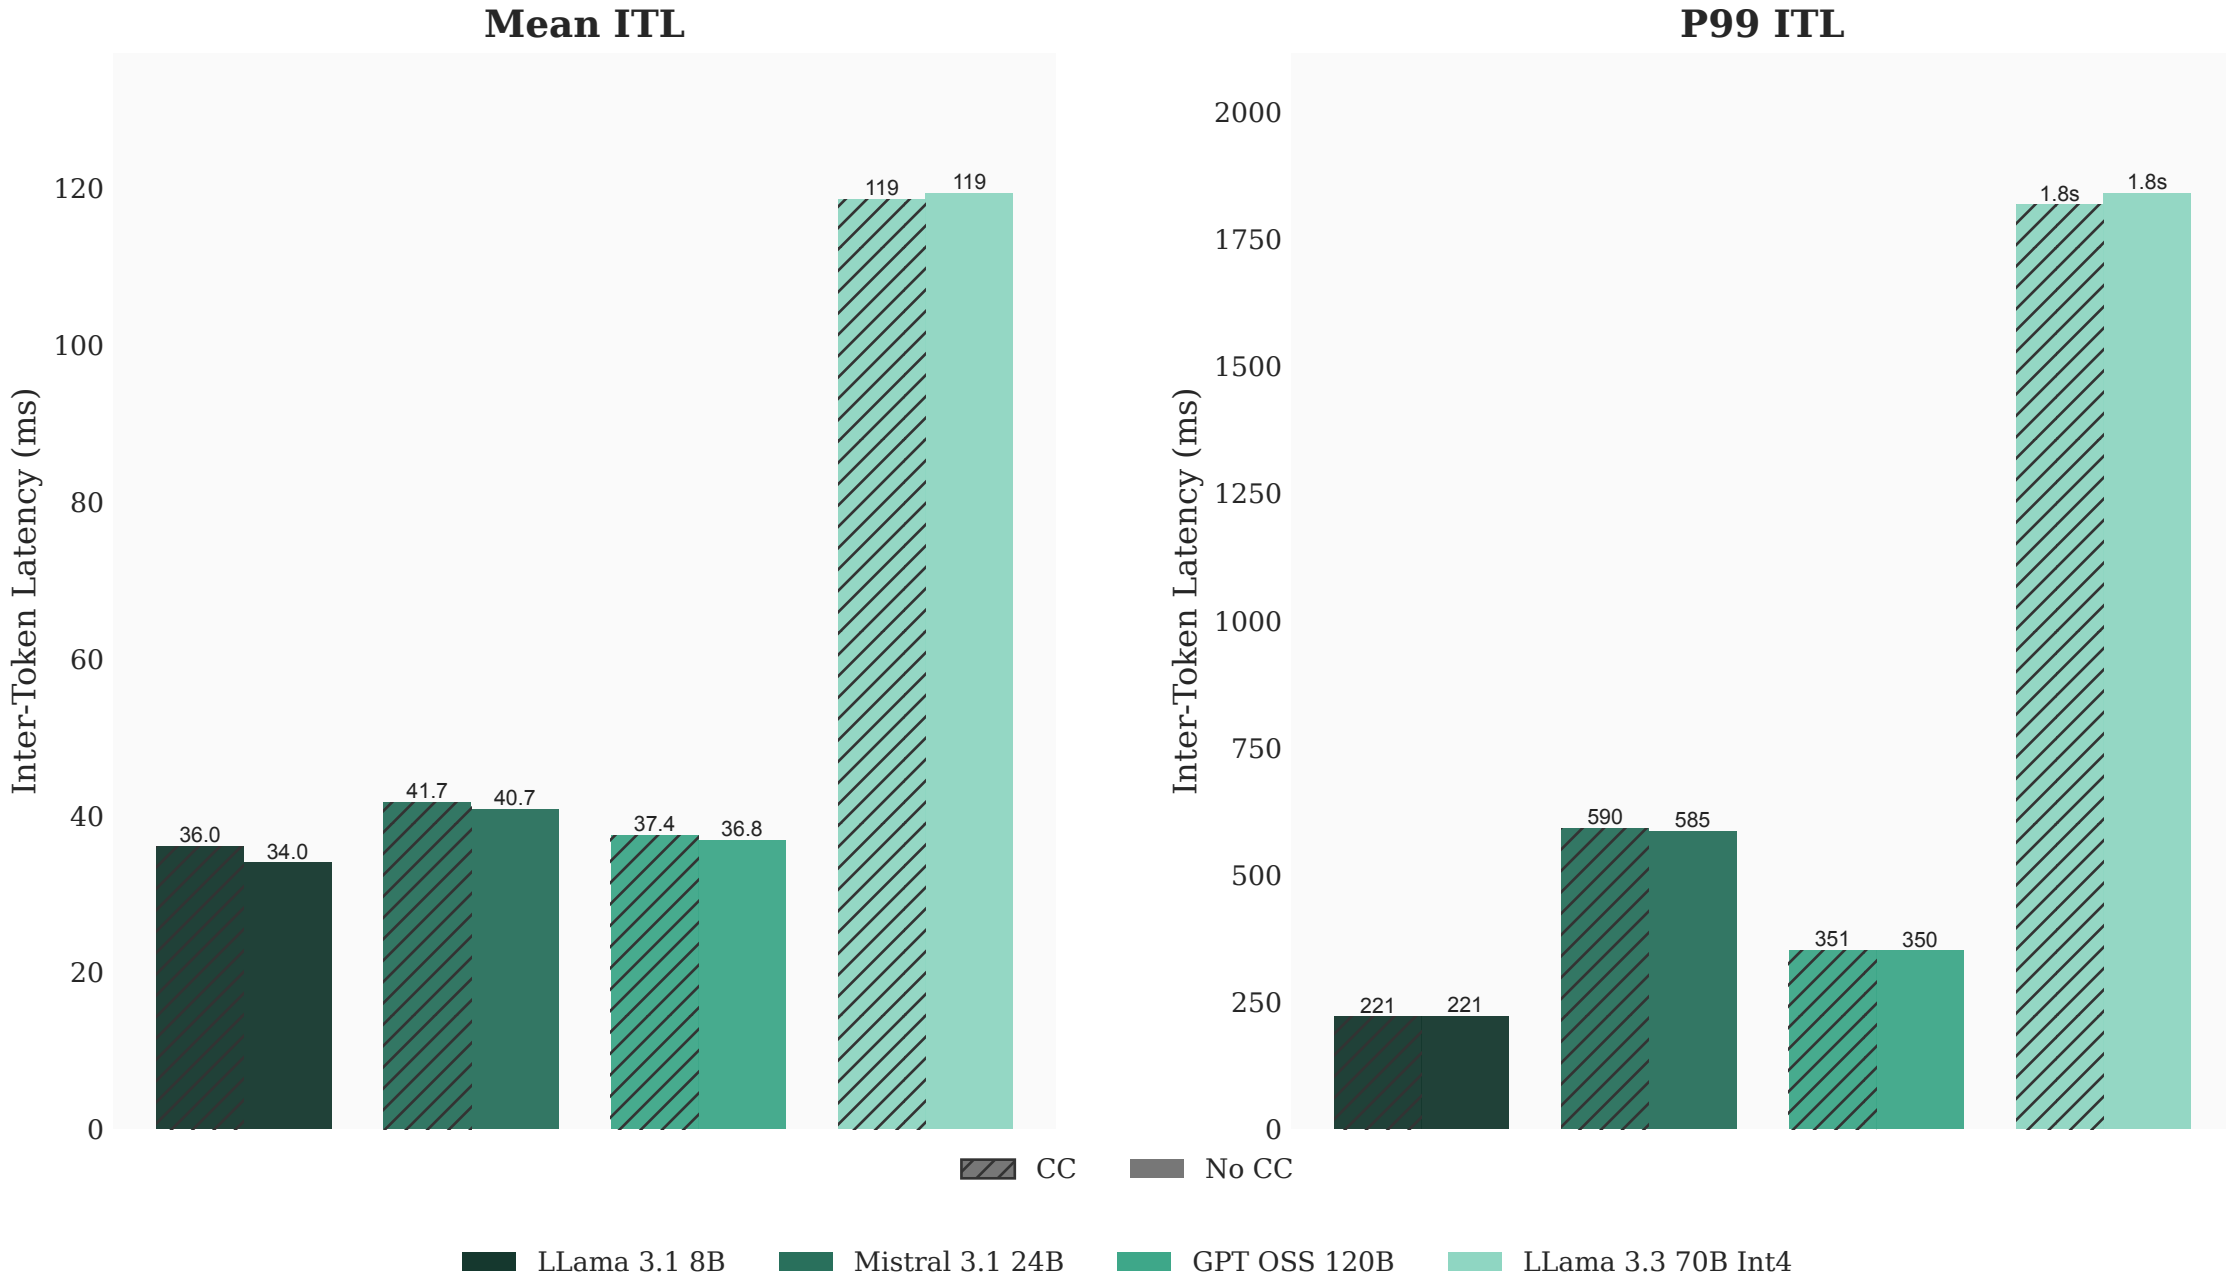
Random (4000 ⇒ 1000) (100 Request Rate)
Mean ITL
P99 ITL
Inter-Token Latency (ms)
Inter-Token Latency (ms)
0
20
40
60
80
100
120
0
250
500
750
1000
1250
1500
1750
2000
36.0
34.0
41.7
40.7
37.4
36.8
119
119
221
221
590
585
351
350
1.8s
1.8s
CC
No CC
LLama 3.1 8B
Mistral 3.1 24B
GPT OSS 120B
LLama 3.3 70B Int4

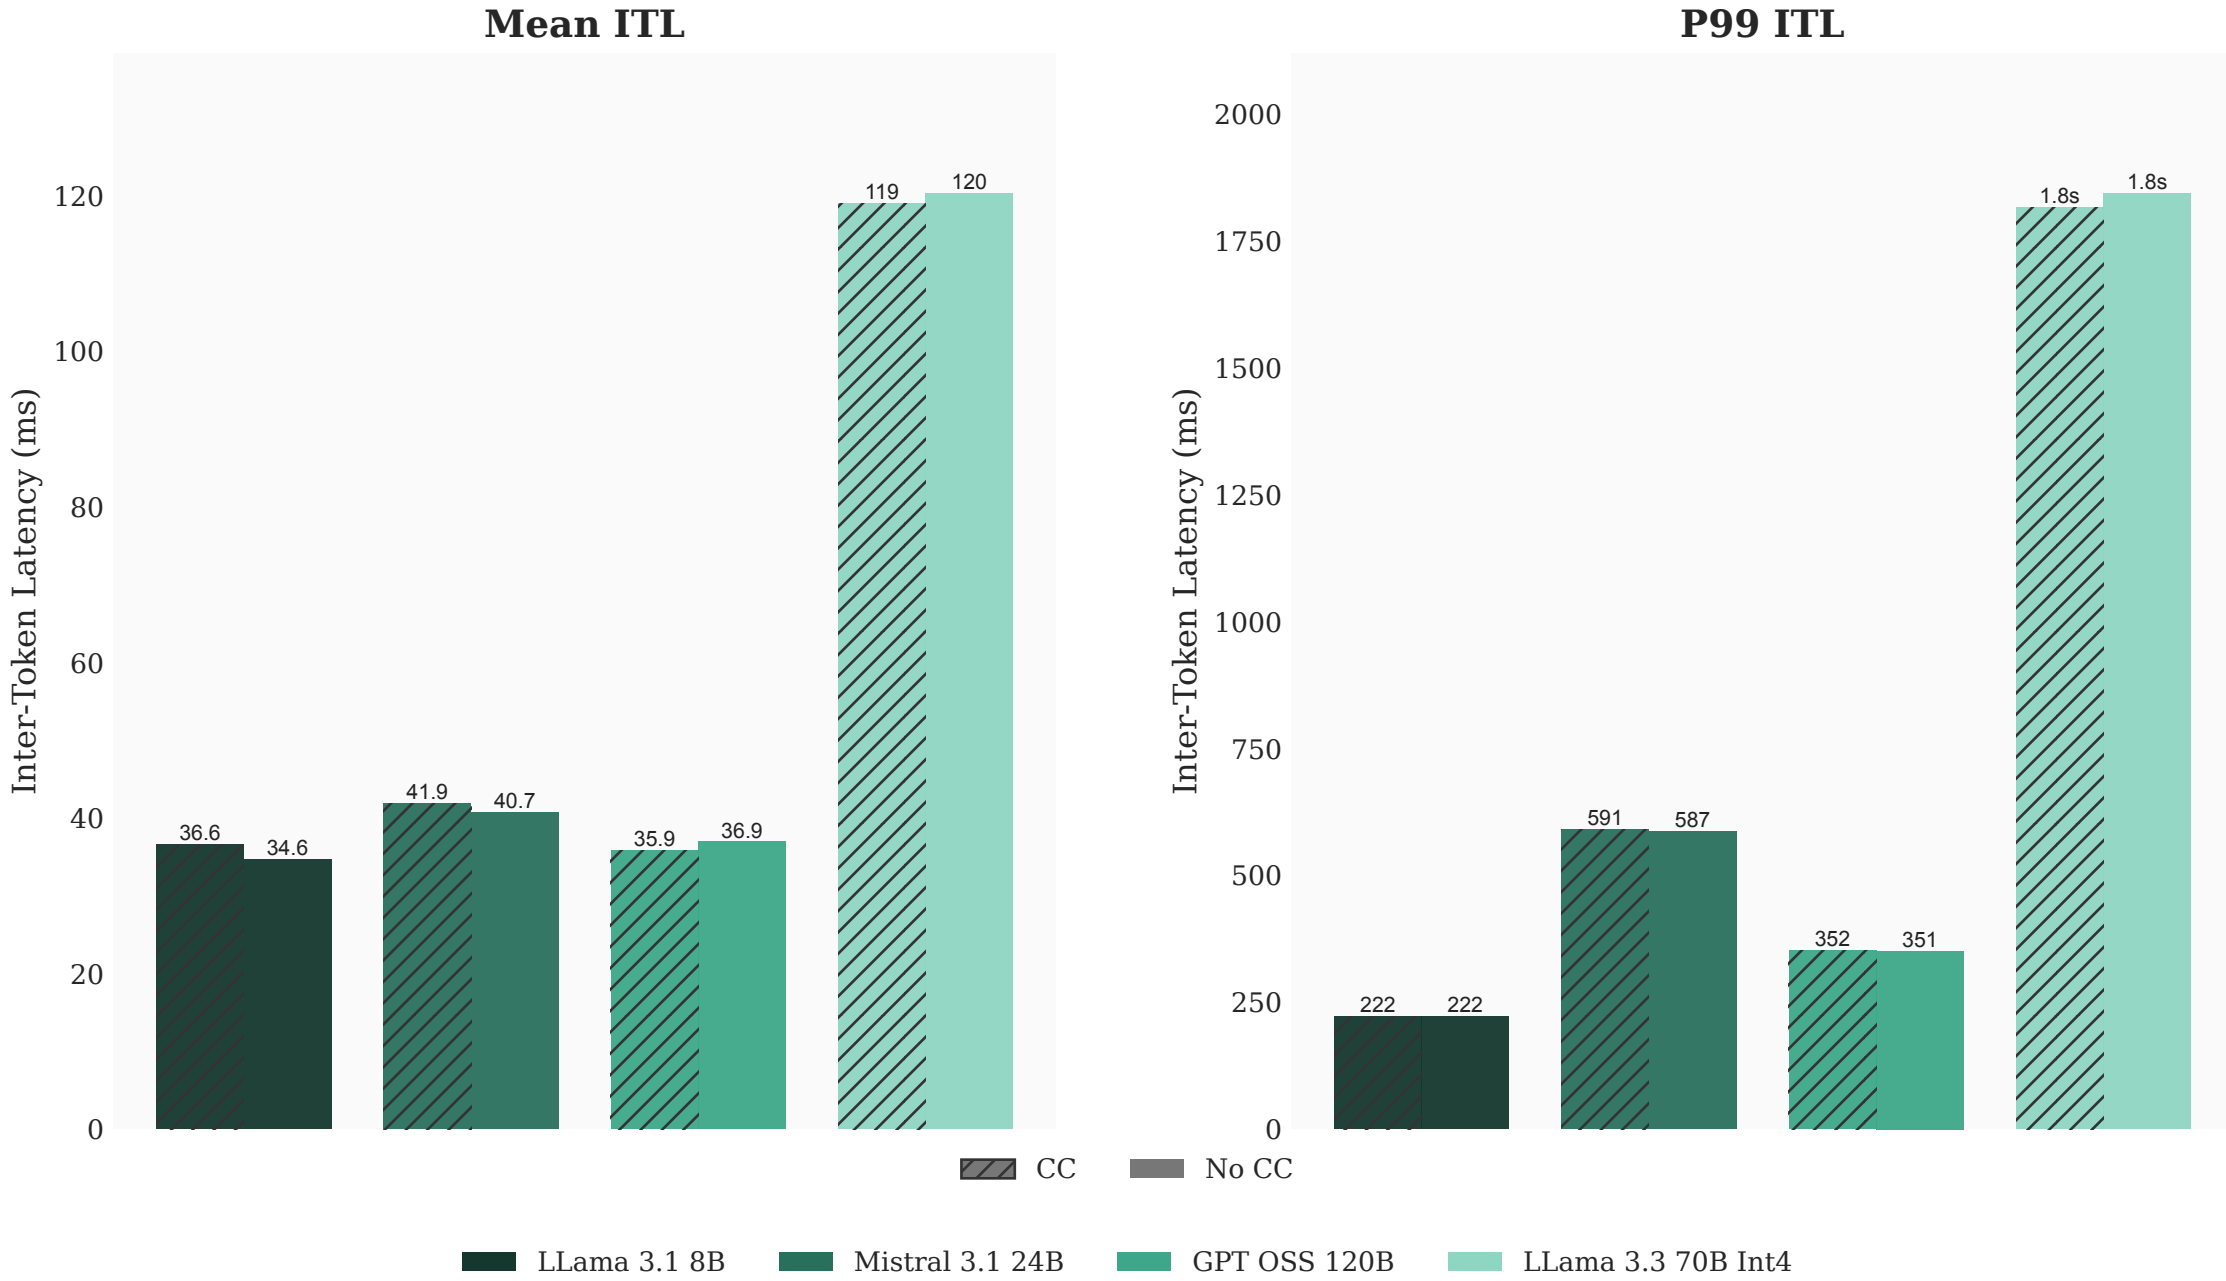
Random (4000 ⇒ 1000) (50 Request Rate)
Mean ITL
P99 ITL
Inter-Token Latency (ms)
0
20
40
60
80
100
120
250
500
750
1000
1250
1500
1750
2000
36.6
34.6
41.9
40.7
35.9
36.9
119
120
222
222
591
587
352
351
1.8s
1.8s
CC
No CC
LLama 3.1 8B
Mistral 3.1 24B
GPT OSS 120B
LLama 3.3 70B Int4

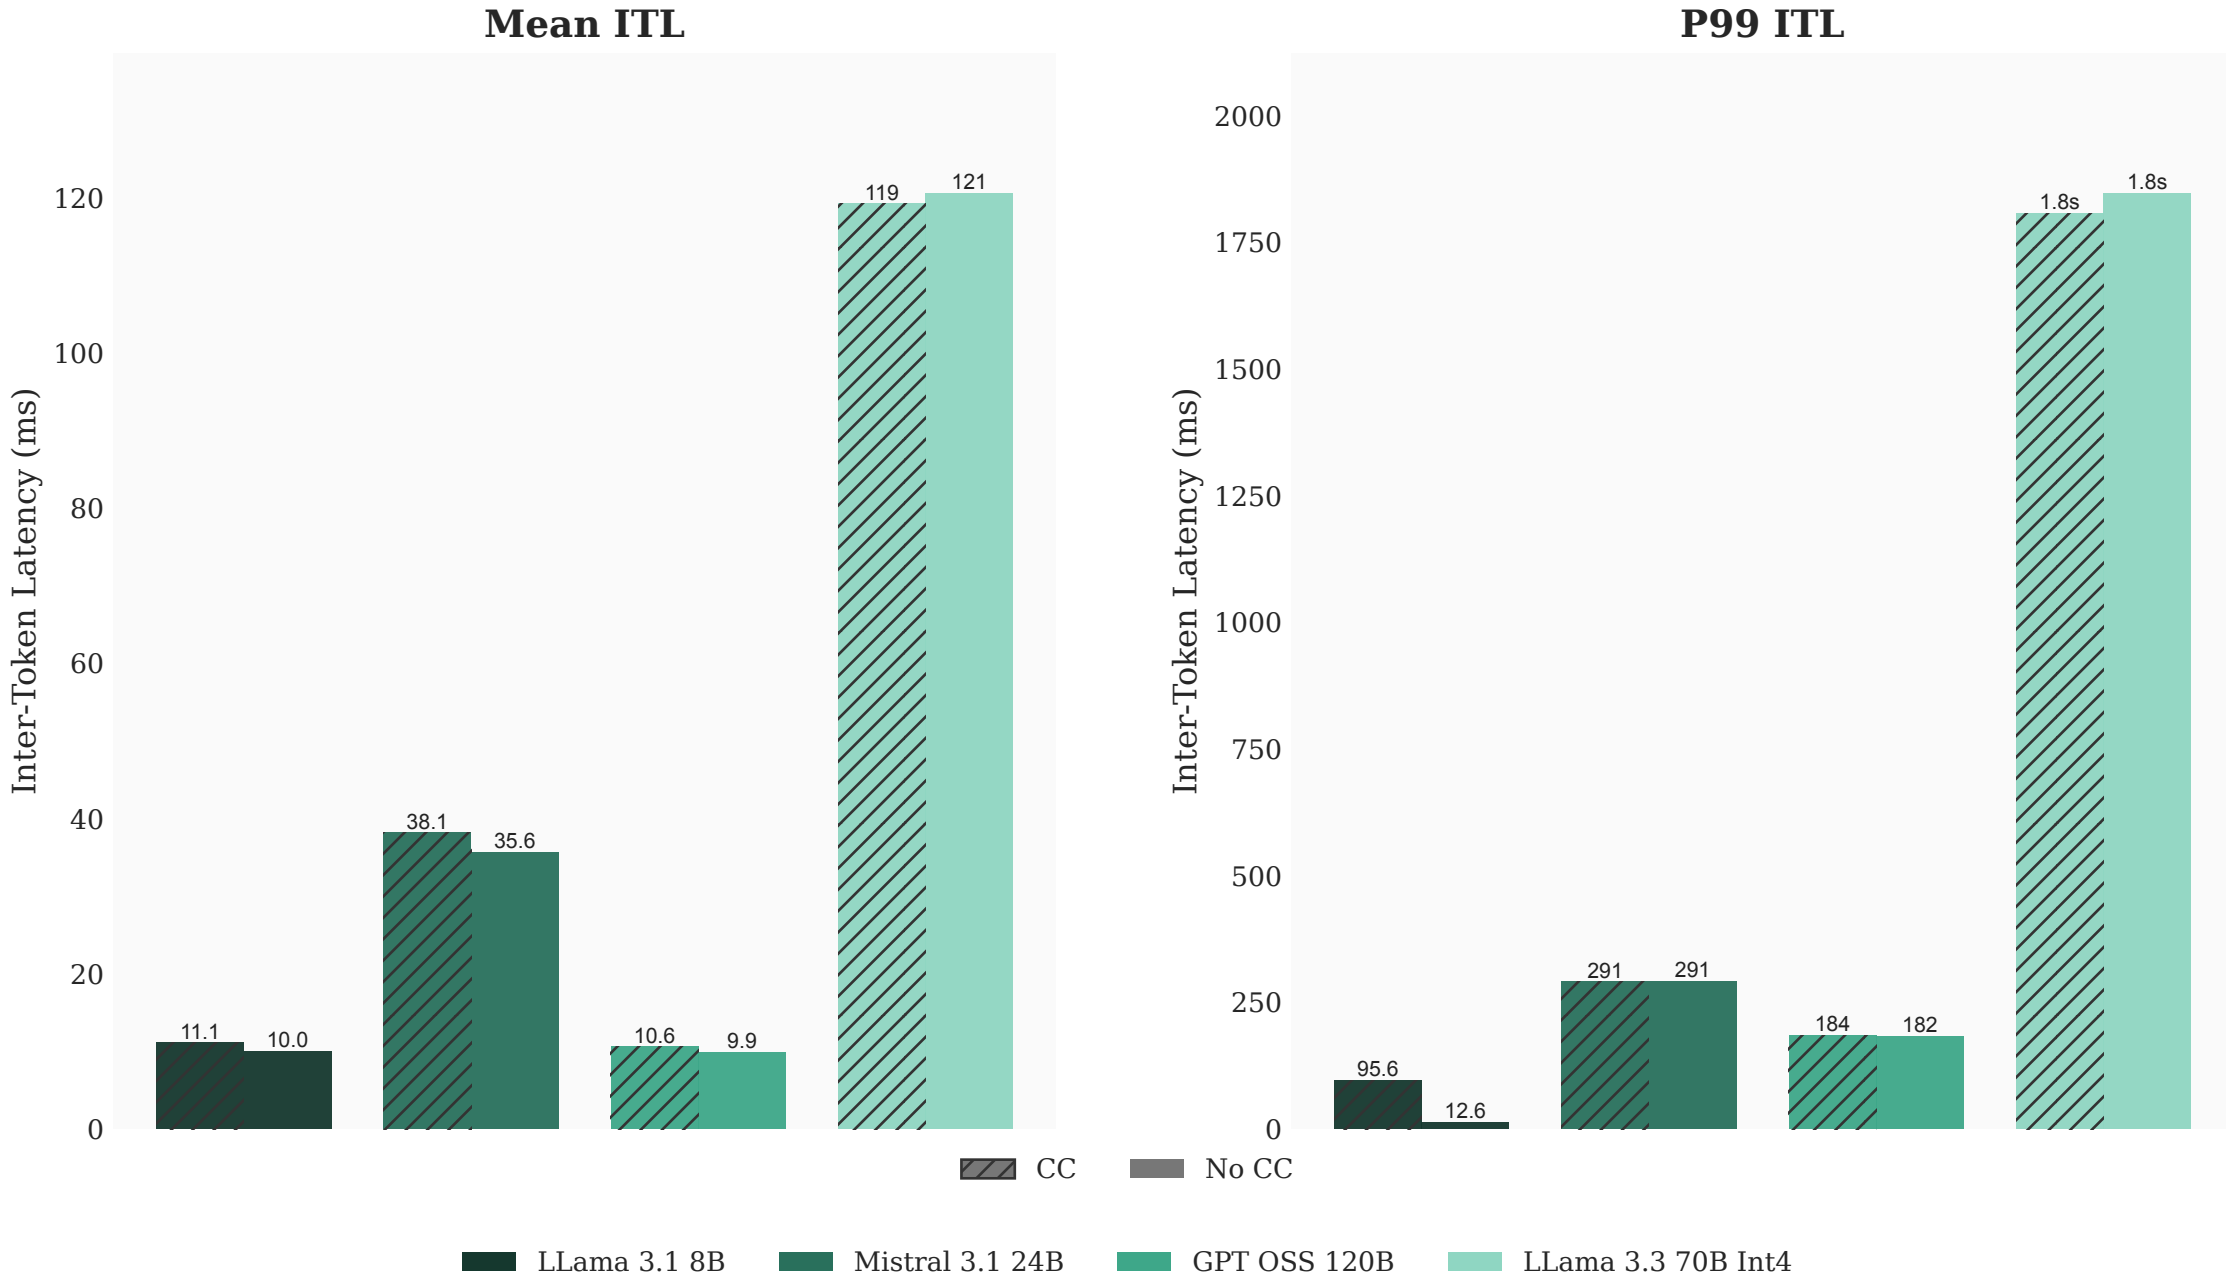
Random (4000 ⇒ 1000) (Single Request)
Mean ITL
P99 ITL
Inter-Token Latency (ms)
Inter-Token Latency (ms)
0
20
40
60
80
100
120
0
250
500
750
1000
1250
1500
1750
2000
11.1
10.0
38.1
35.6
10.6
9.9
119
121
95.6
12.6
291
291
184
182
1.8s
1.8s
CC
No CC
LLama 3.1 8B
Mistral 3.1 24B
GPT OSS 120B
LLama 3.3 70B Int4

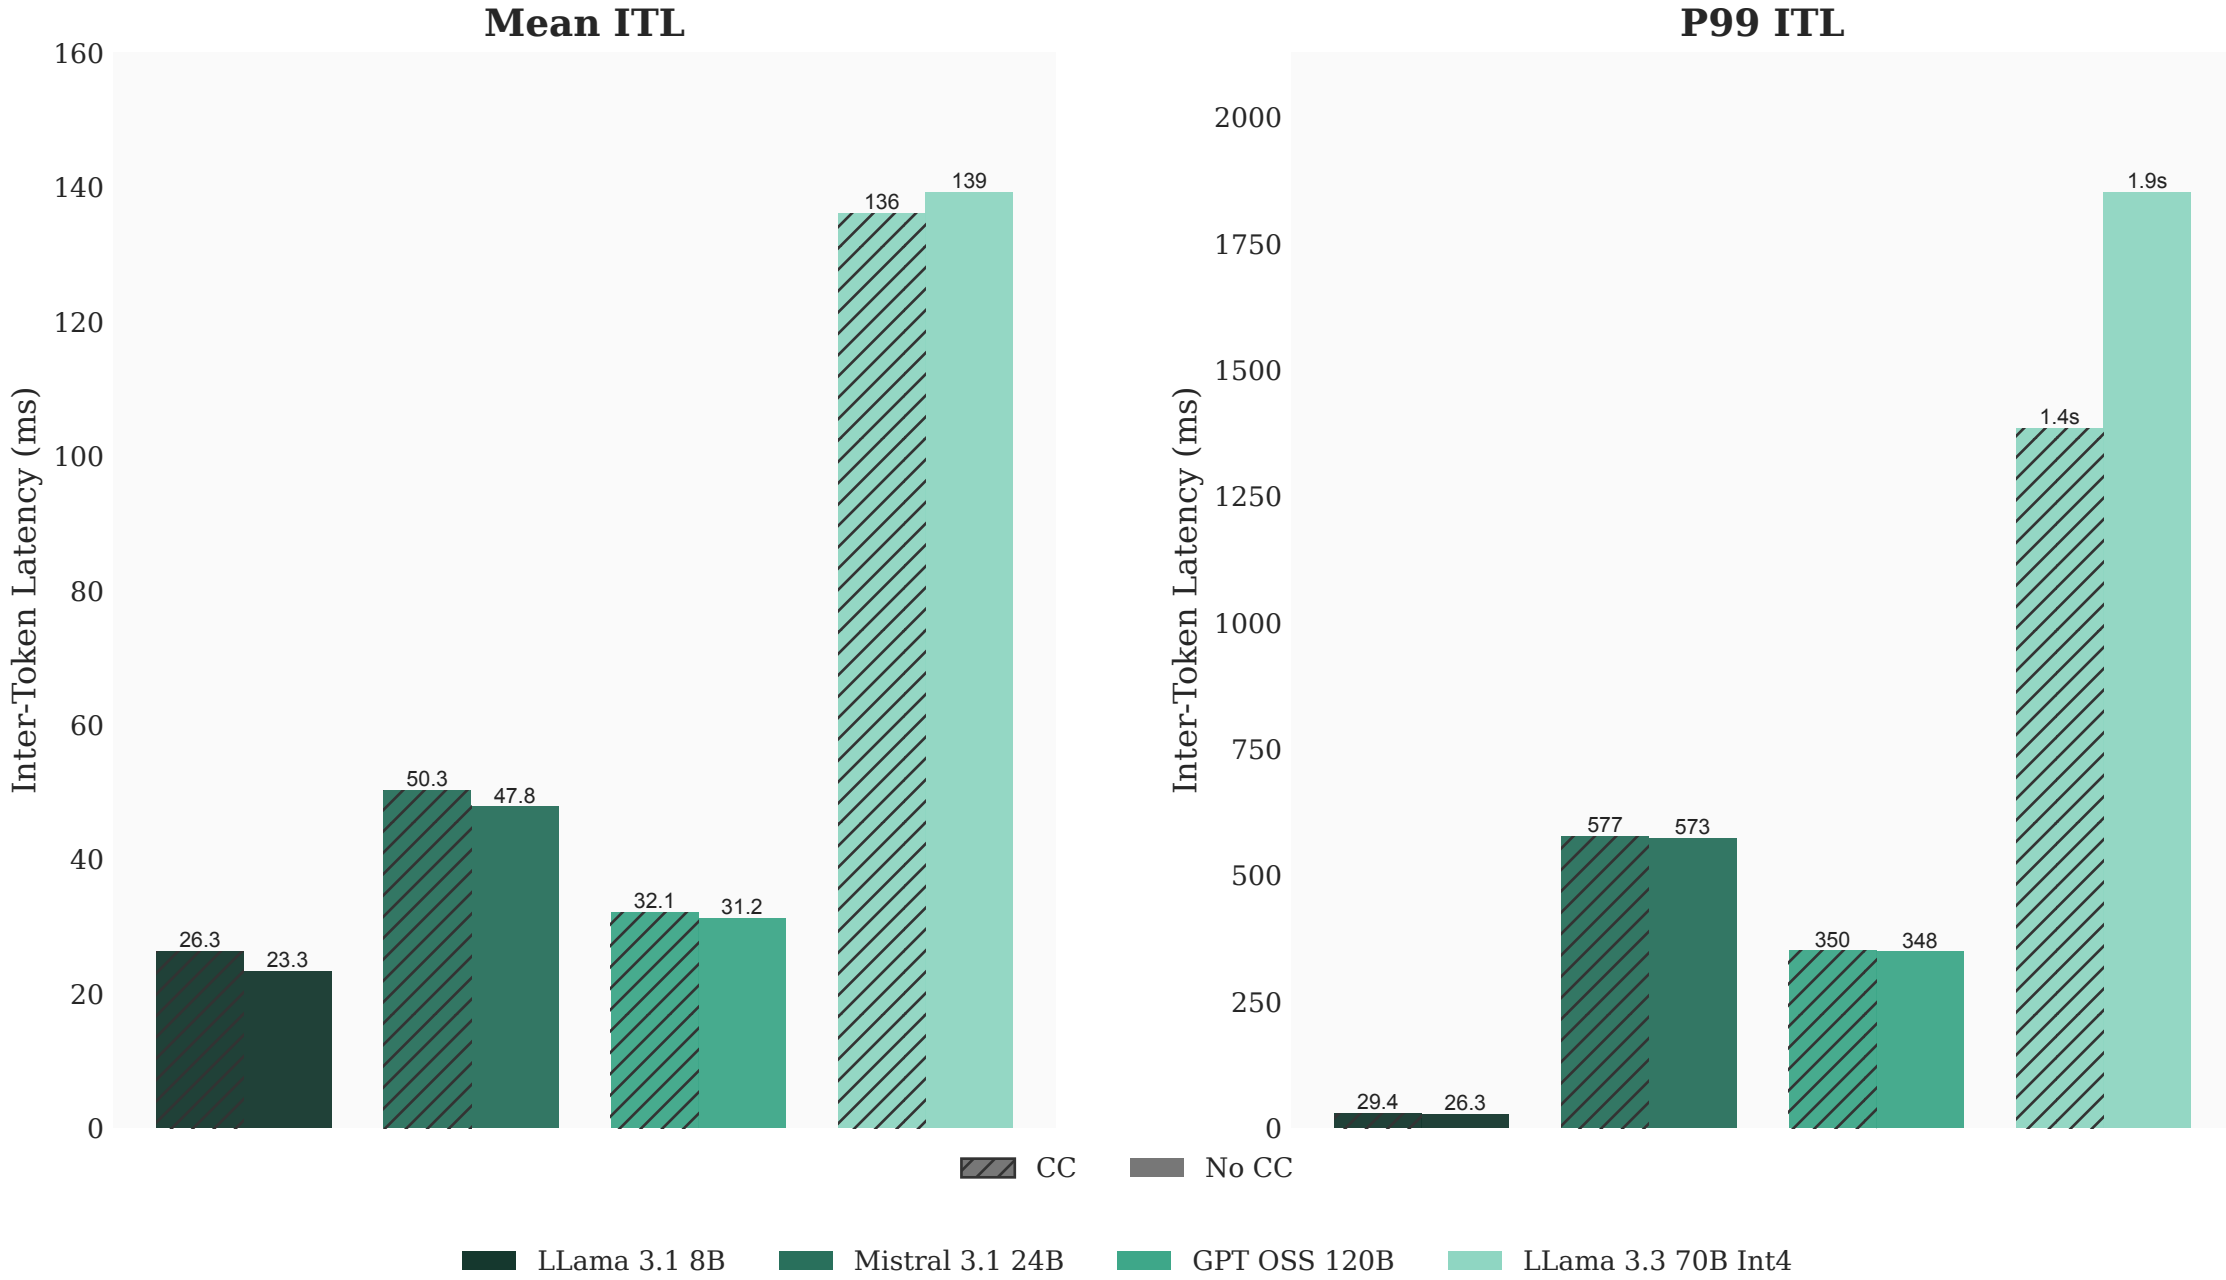
Random (1000 ⇒ 1000) (100 Request Rate)
Mean ITL
P99 ITL
Inter-Token Latency (ms)
0
20
40
60
80
100
120
140
160
250
500
750
1000
1250
1500
1750
2000
26.3
23.3
50.3
47.8
32.1
31.2
136
139
29.4
26.3
577
573
350
348
1.4s
1.9s
CC
No CC
LLama 3.1 8B
Mistral 3.1 24B
GPT OSS 120B
LLama 3.3 70B Int4

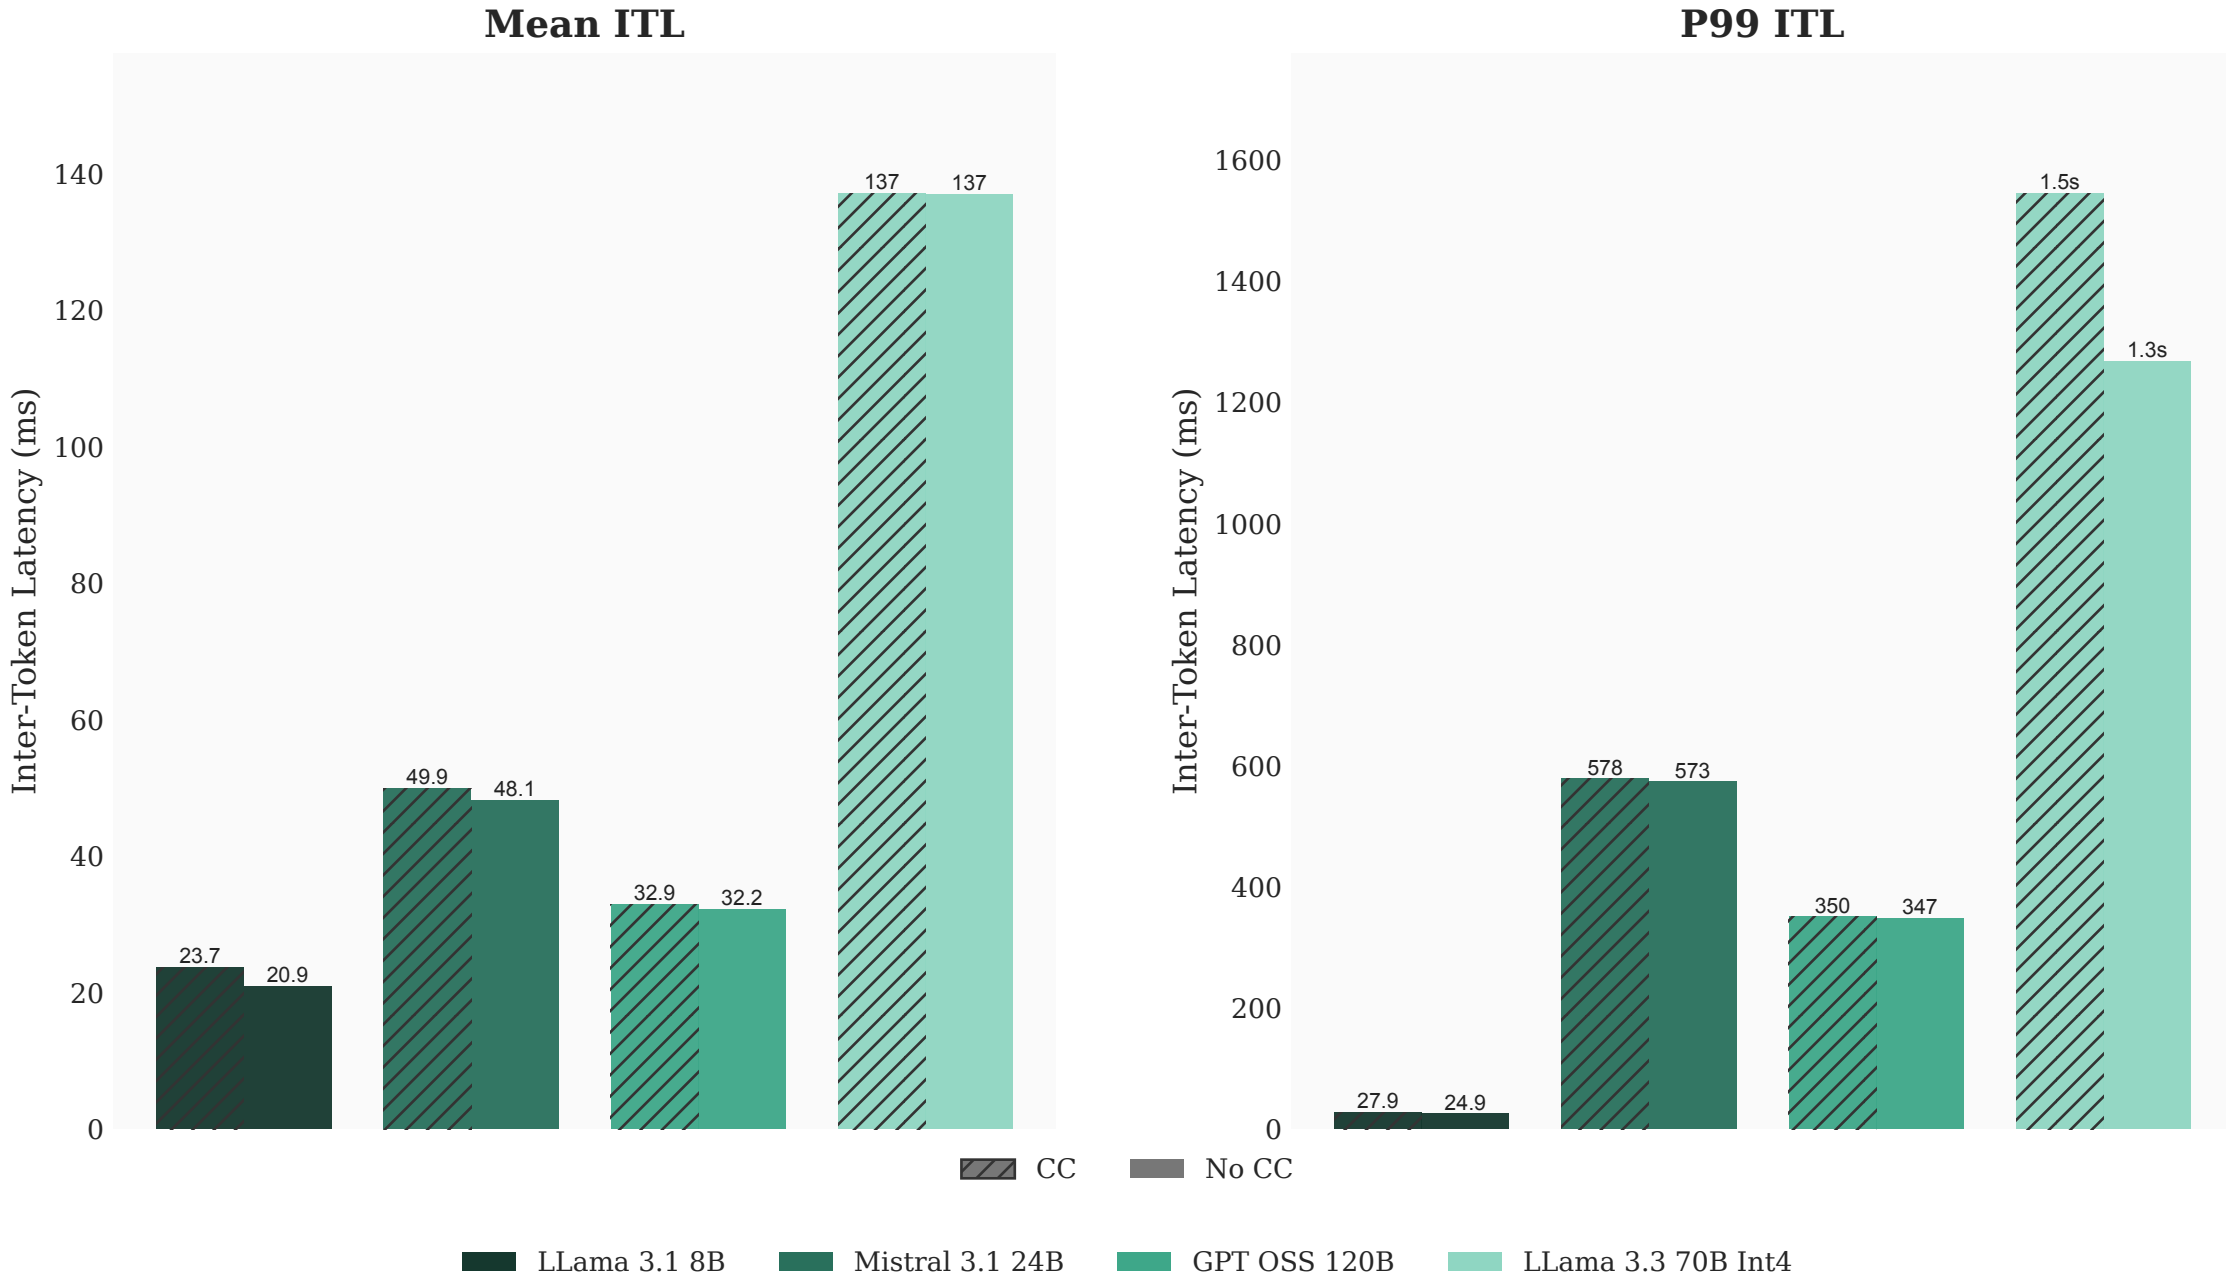
Random (1000 ⇒ 1000) (50 Request Rate)
Mean ITL
P99 ITL
Inter-Token Latency (ms)
0
20
40
60
80
100
120
140
200
400
600
800
1000
1200
1400
1600
23.7
20.9
49.9
48.1
32.9
32.2
137
137
27.9
24.9
578
573
350
347
1.5s
1.3s
CC
No CC
LLama 3.1 8B
Mistral 3.1 24B
GPT OSS 120B
LLama 3.3 70B Int4

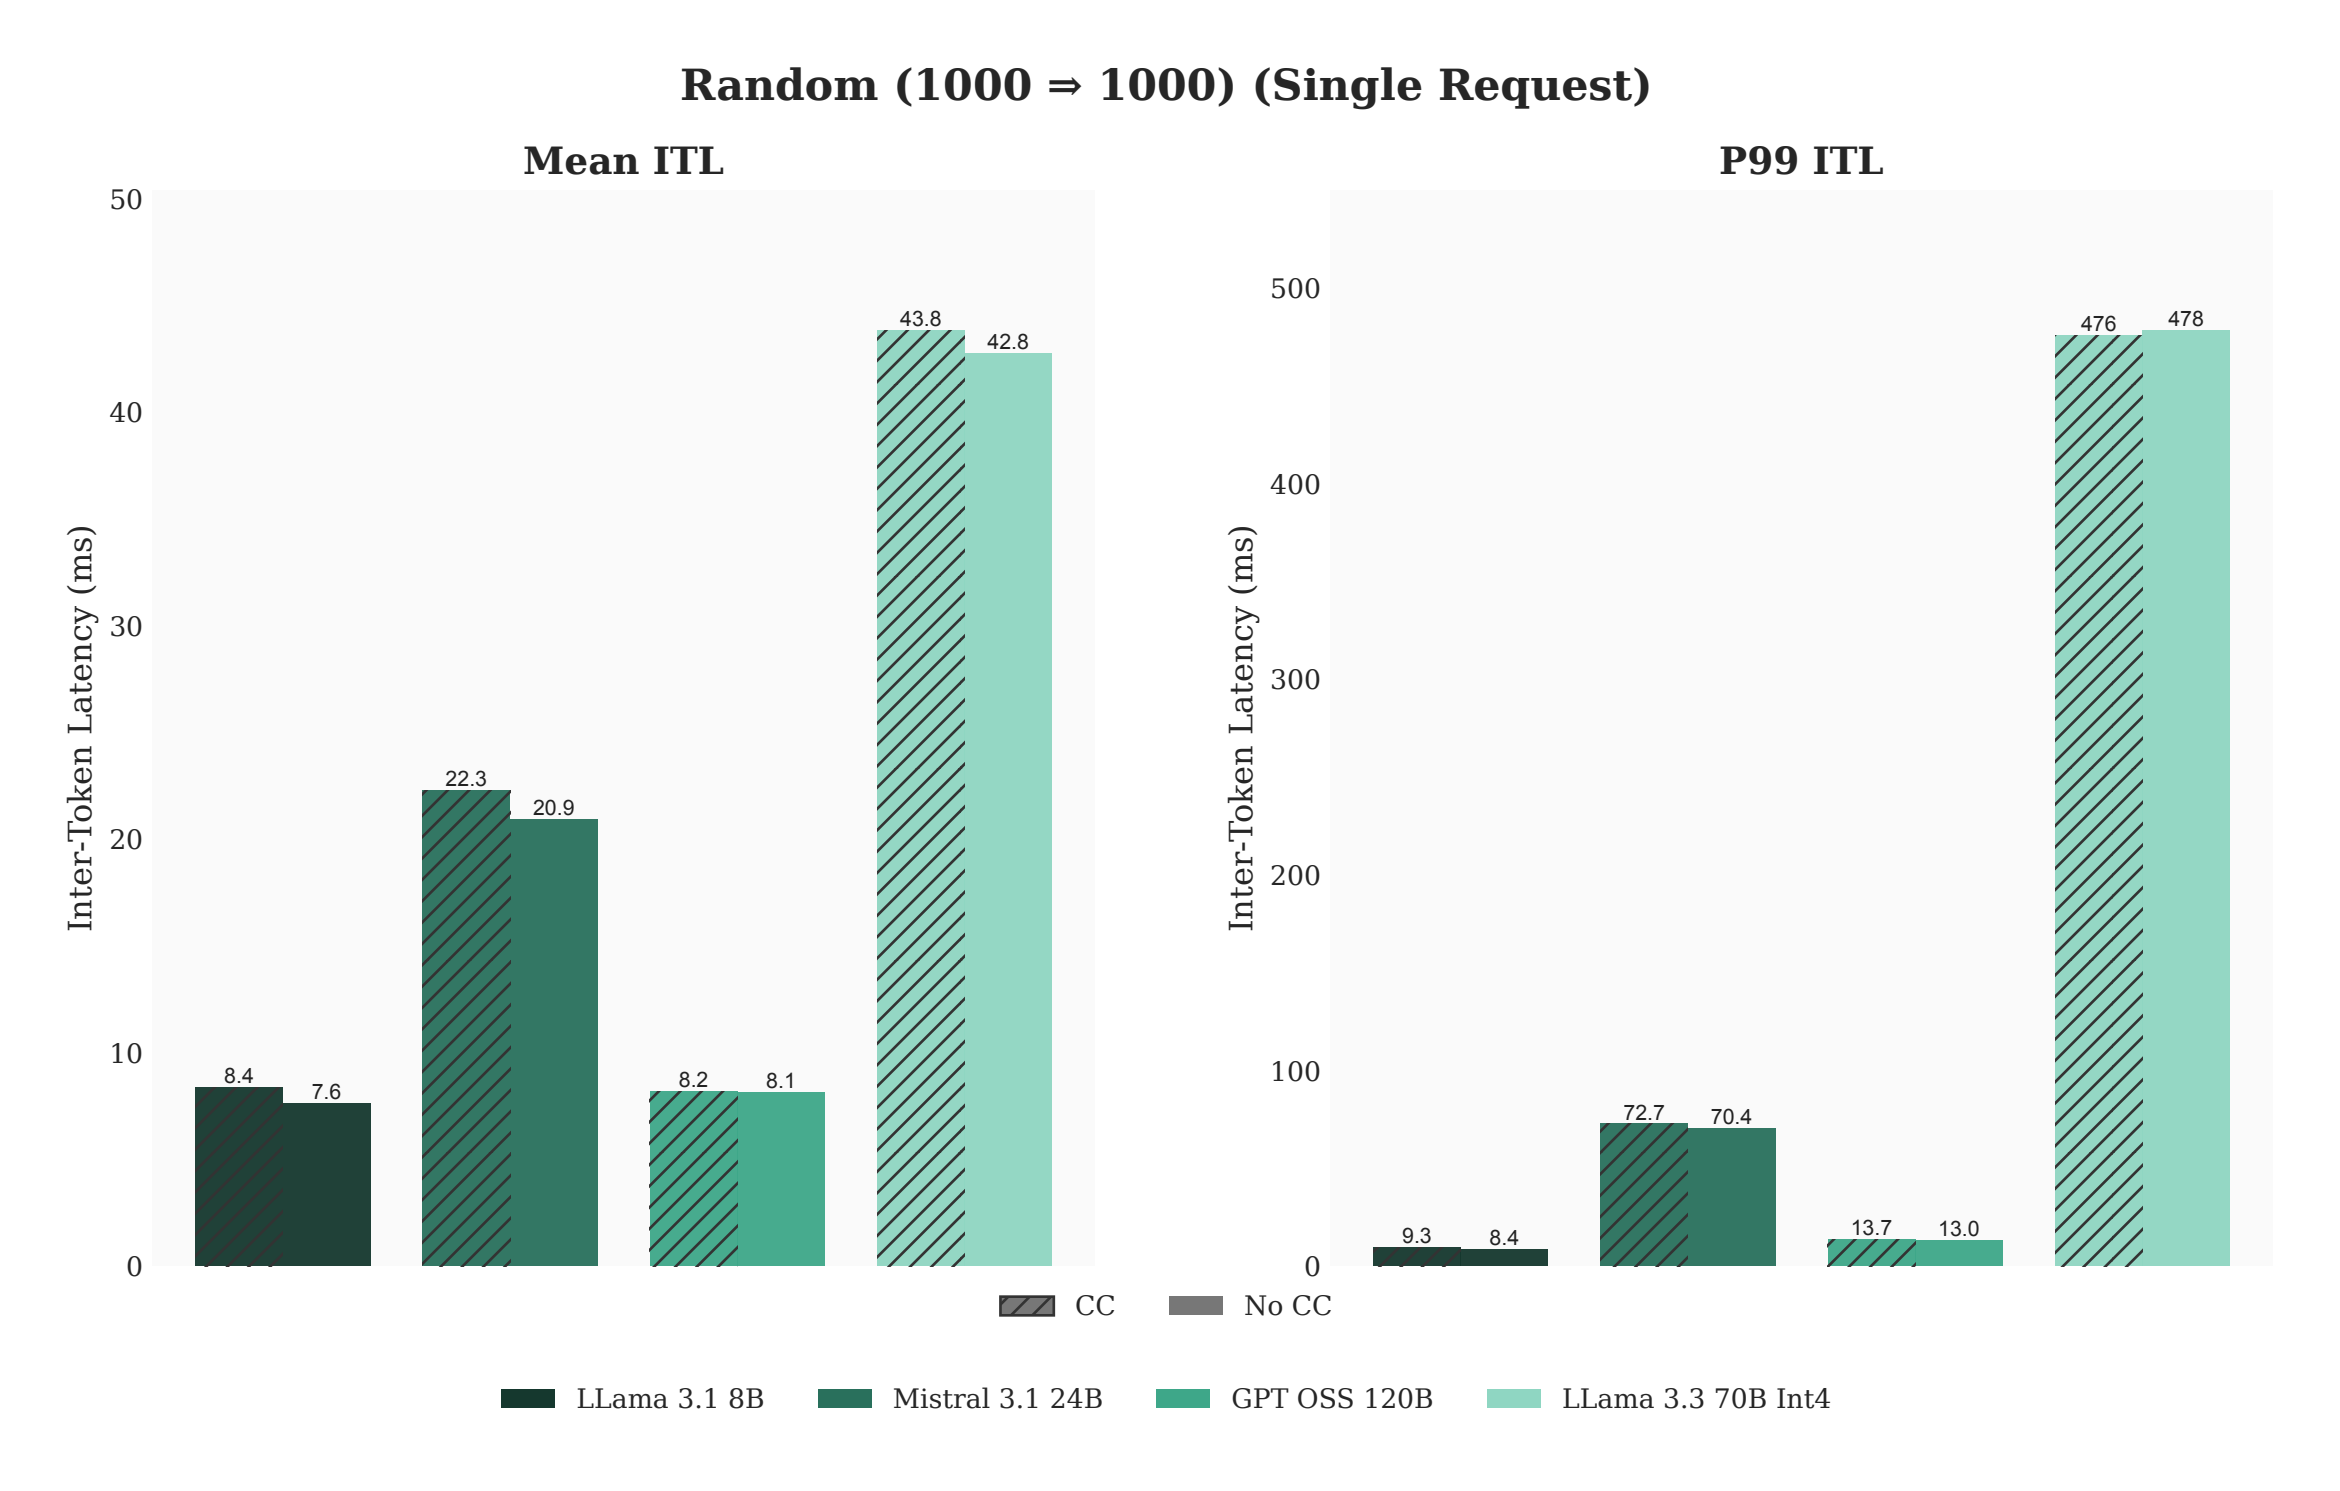
Random (1000 ⇒ 1000) (Single Request)
Mean ITL
P99 ITL
Inter-Token Latency (ms)
Inter-Token Latency (ms)
0
10
20
30
40
50
0
100
200
300
400
500
8.4
7.6
22.3
20.9
8.2
8.1
43.8
42.8
9.3
8.4
72.7
70.4
13.7
13.0
476
478
CC
No CC
LLama 3.1 8B
Mistral 3.1 24B
GPT OSS 120B
LLama 3.3 70B Int4

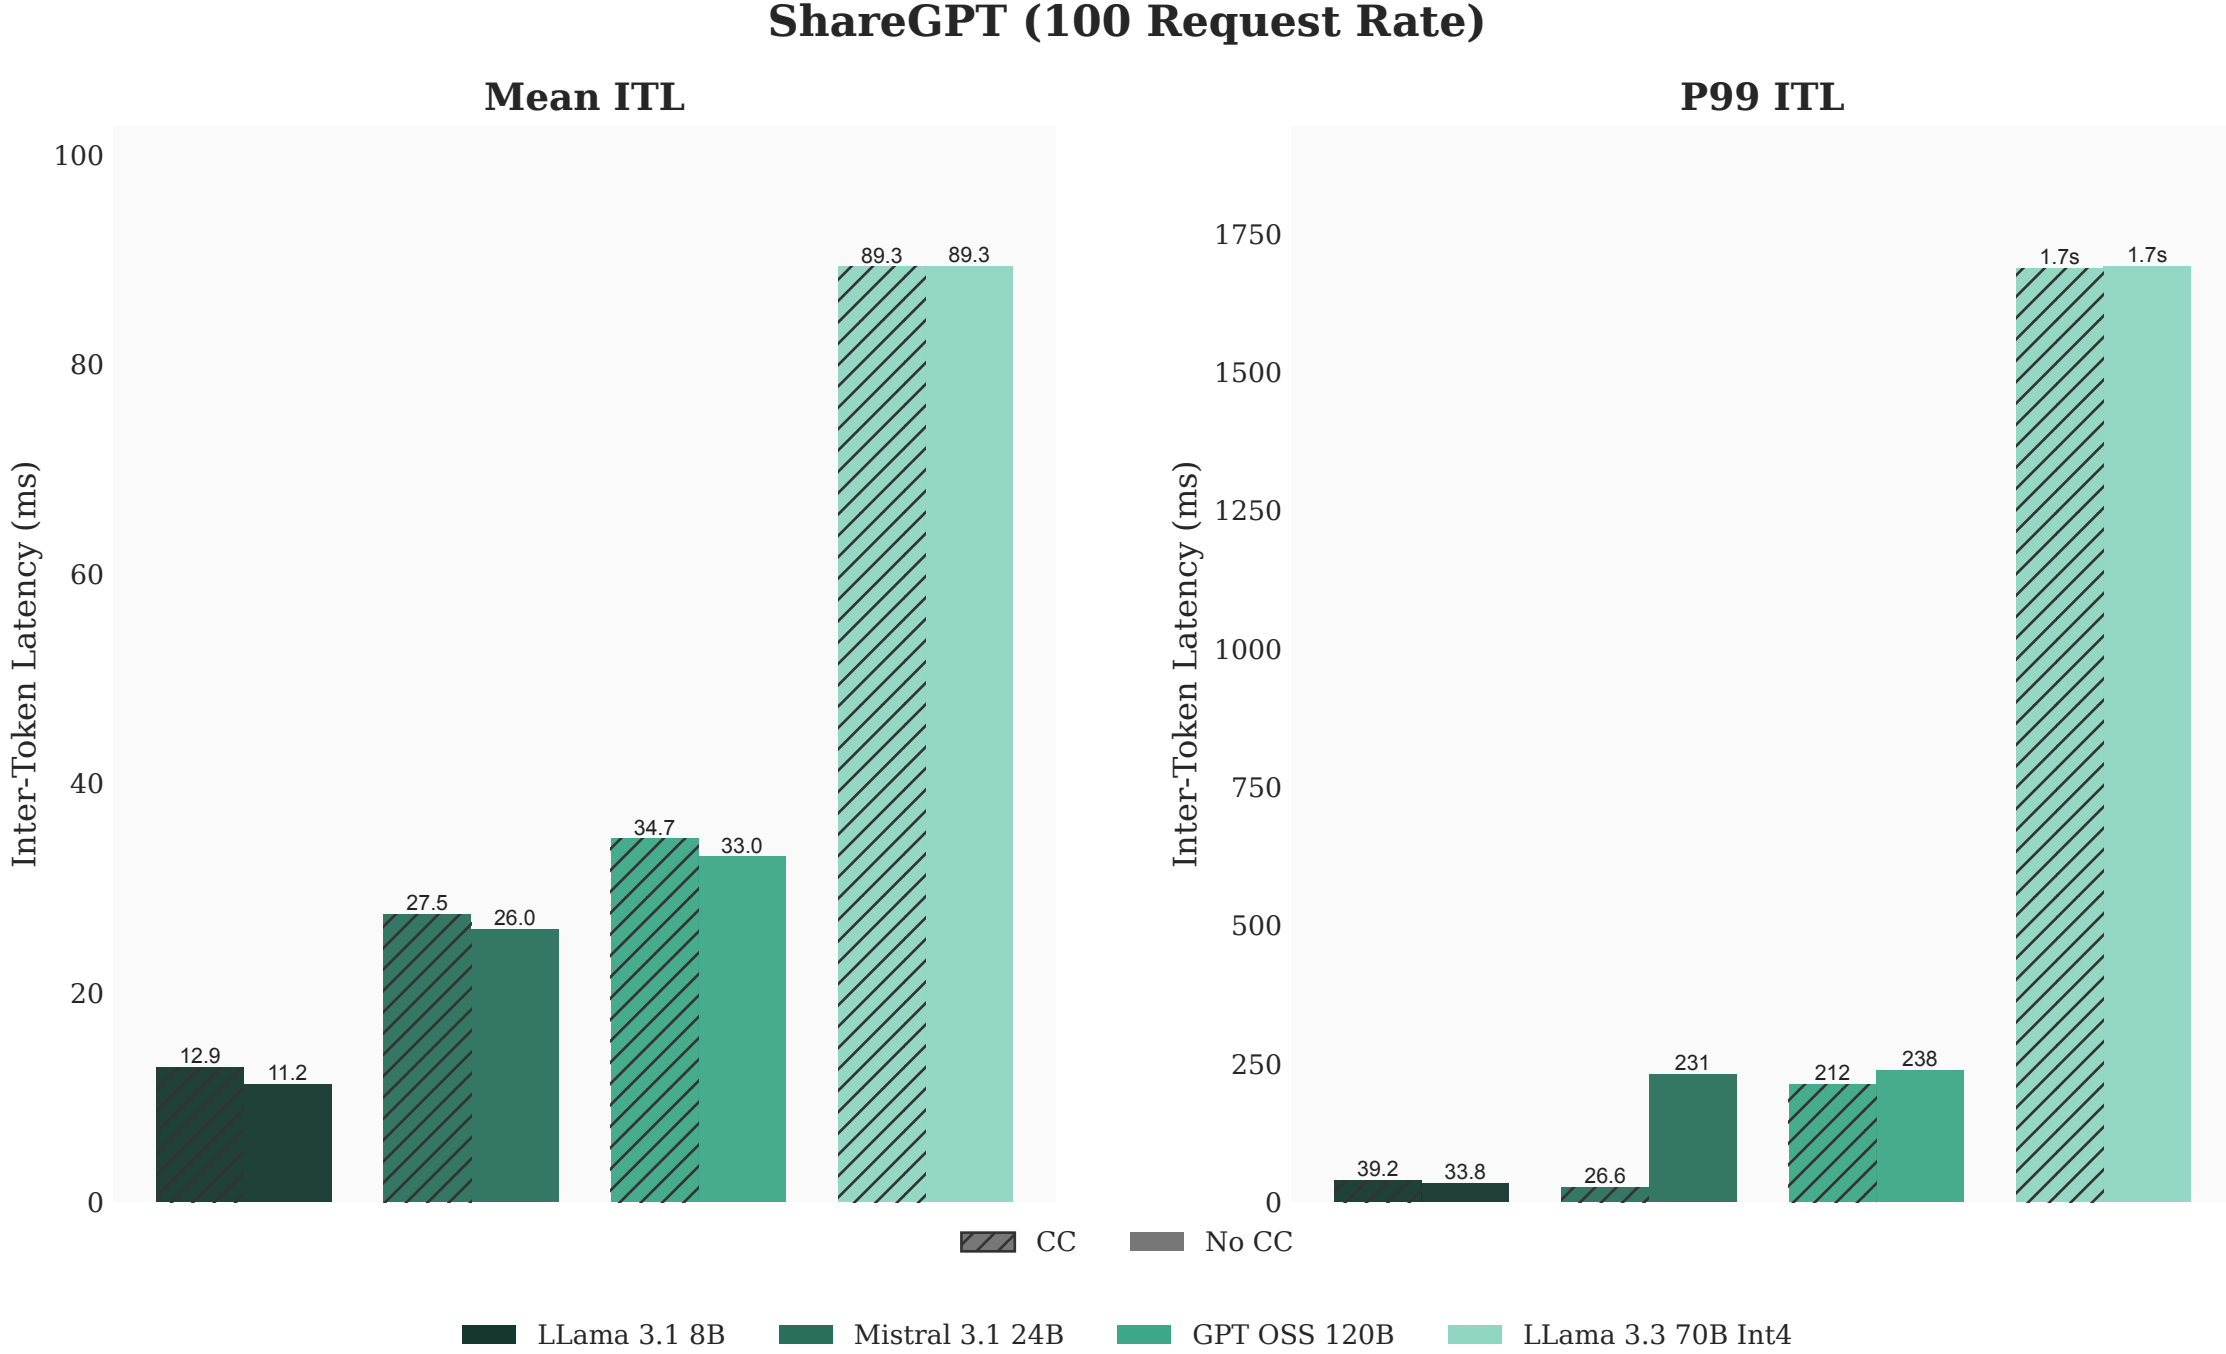
ShareGPT (100 Request Rate)
Mean ITL
P99 ITL
Inter-Token Latency (ms)
Inter-Token Latency (ms)
0
20
40
60
80
100
0
250
500
750
1000
1250
1500
1750
12.9
11.2
27.5
26.0
34.7
33.0
89.3
89.3
39.2
33.8
26.6
231
212
238
1.7s
1.7s
CC
No CC
LLama 3.1 8B
Mistral 3.1 24B
GPT OSS 120B
LLama 3.3 70B Int4

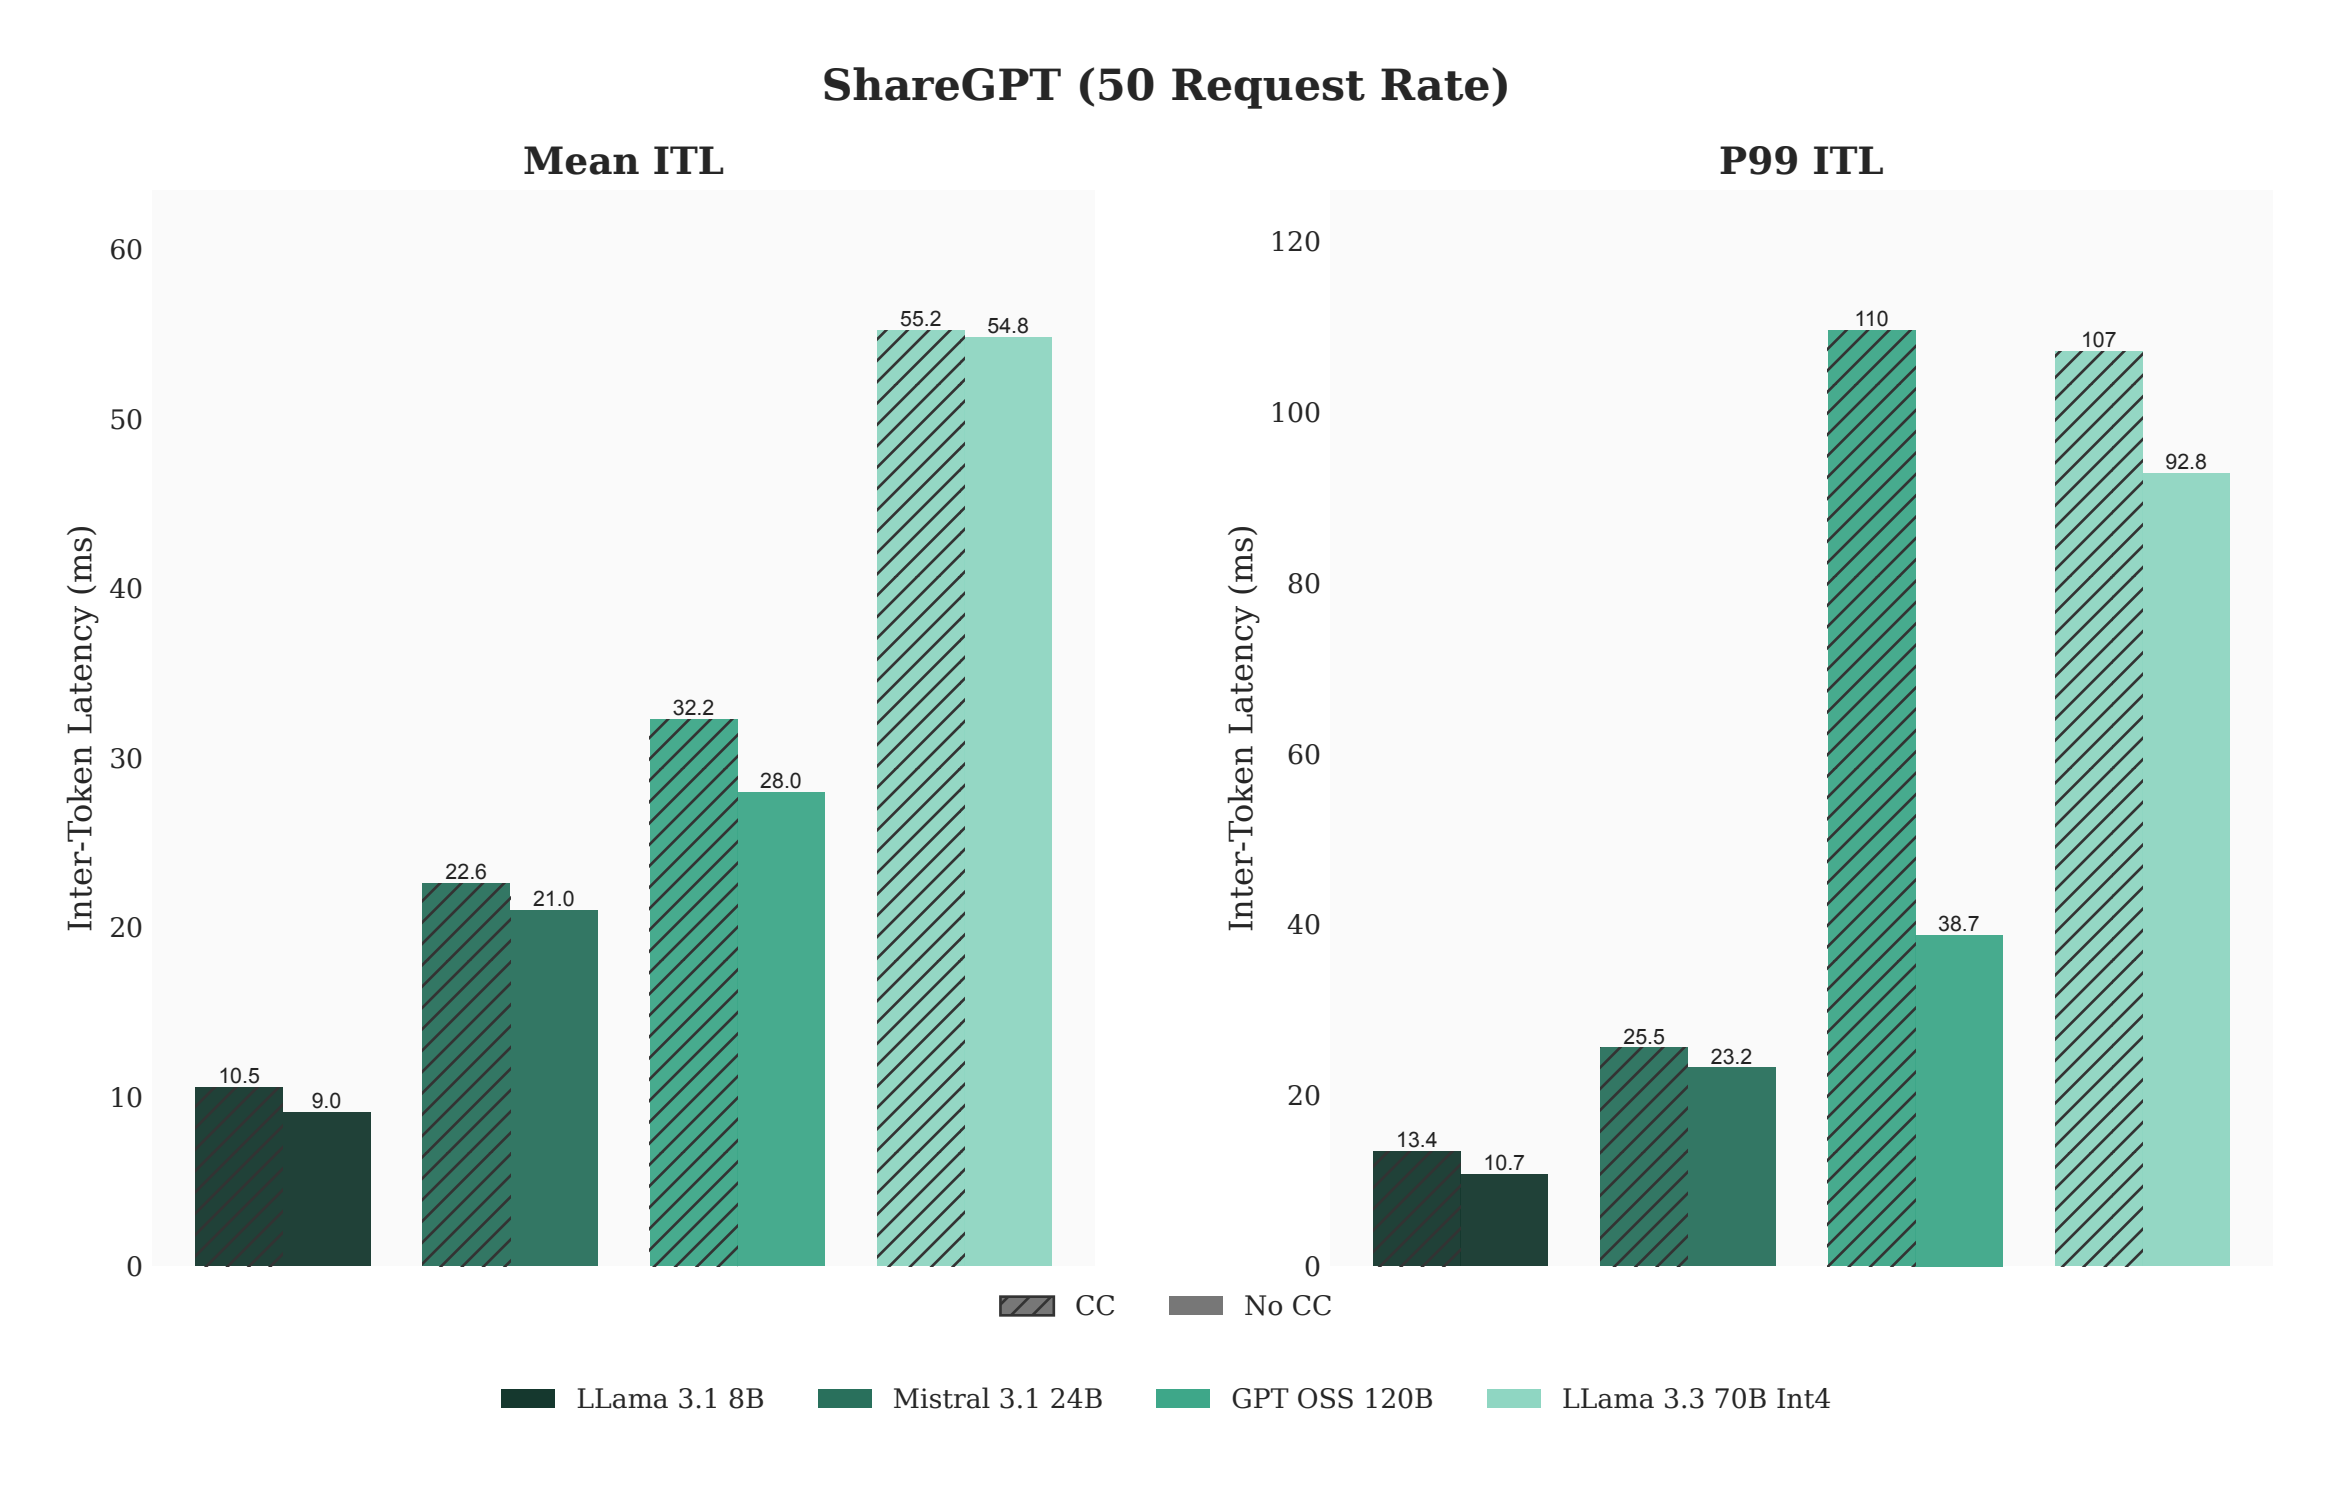
ShareGPT (50 Request Rate)
Mean ITL
P99 ITL
Inter-Token Latency (ms)
Inter-Token Latency (ms)
0
10
20
30
40
50
60
0
20
40
60
80
100
120
10.5
9.0
22.6
21.0
32.2
28.0
55.2
54.8
13.4
10.7
25.5
23.2
110
38.7
107
92.8
CC
No CC
LLama 3.1 8B
Mistral 3.1 24B
GPT OSS 120B
LLama 3.3 70B Int4

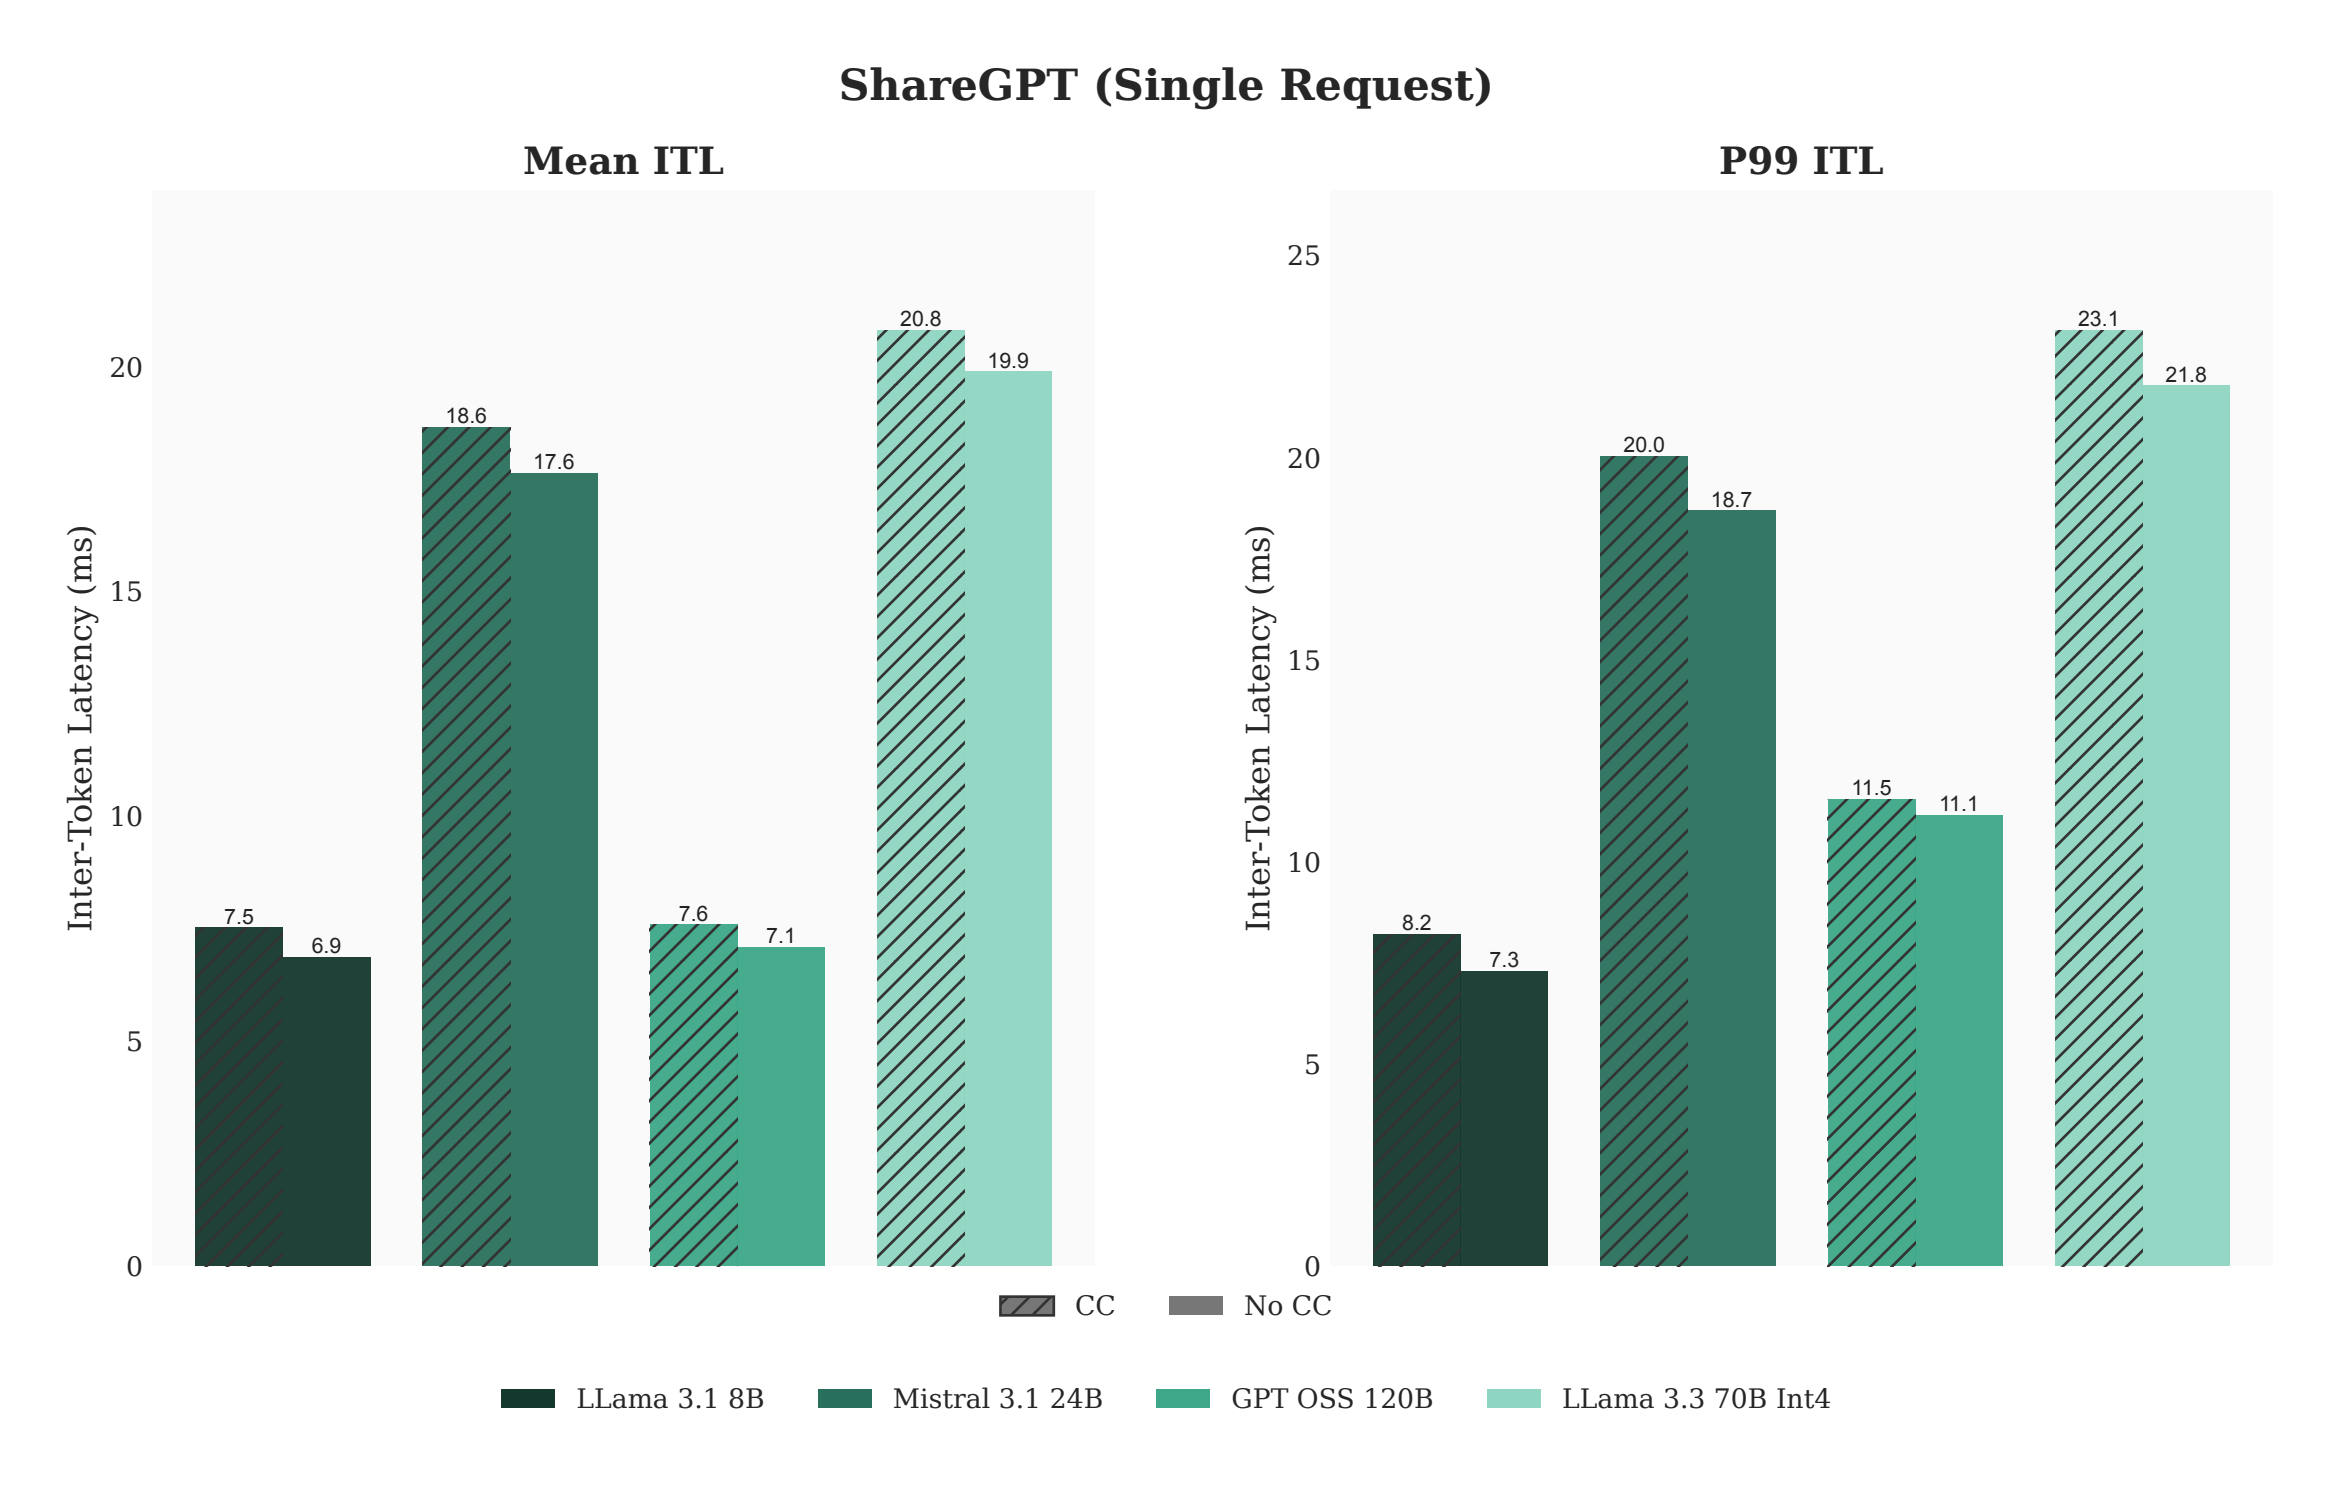
ShareGPT (Single Request)
Mean ITL
P99 ITL
Inter-Token Latency (ms)
0
5
10
15
20
25
7.5
6.9
18.6
17.6
7.6
7.1
20.8
19.9
8.2
7.3
20.0
18.7
11.5
11.1
23.1
21.8
CC
No CC
LLama 3.1 8B
Mistral 3.1 24B
GPT OSS 120B
LLama 3.3 70B Int4


Edit 10K Characters (100 Request Rate)
Mean ITL
P99 ITL
Inter-Token Latency (ms)
Inter-Token Latency (ms)
0
10
20
30
40
50
60
0
50
100
150
200
34.0
31.2
38.0
36.7
28.0
28.6
57.0
57.0
57.5
54.6
202
206
136
139
41.0
40.8
CC
No CC
LLama 3.1 8B
Mistral 3.1 24B
GPT OSS 120B
LLama 3.3 70B Int4

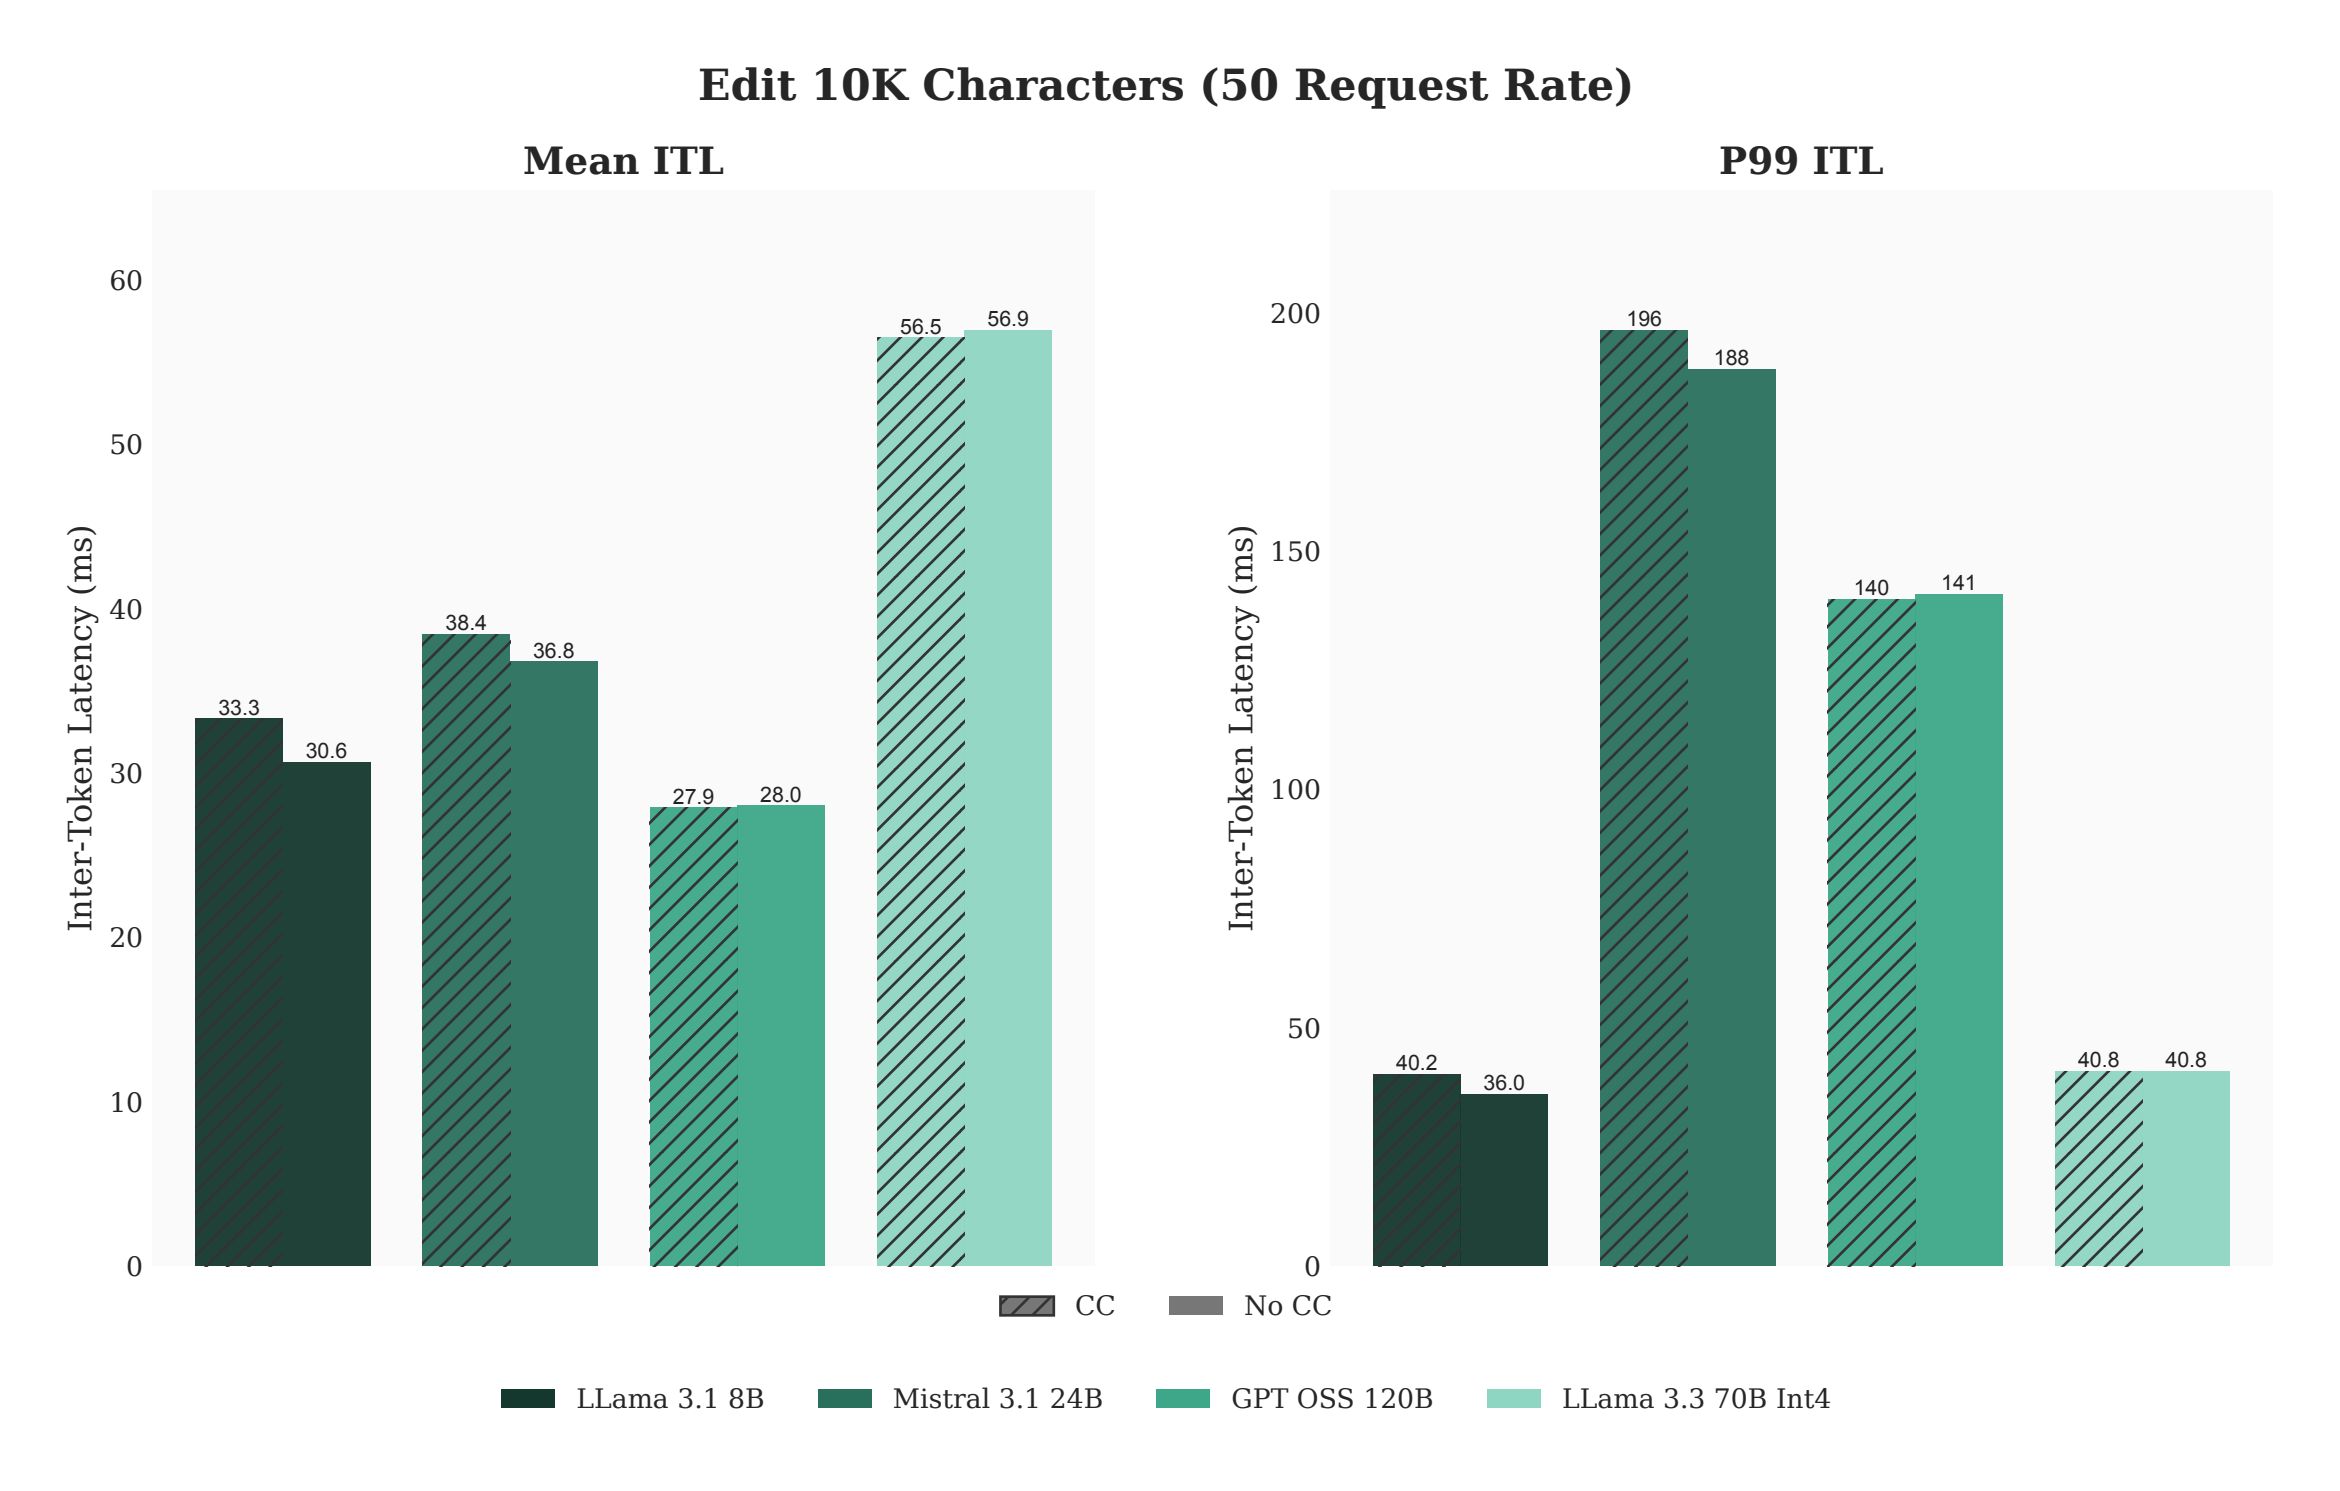
Edit 10K Characters (50 Request Rate)
Mean ITL
P99 ITL
Inter-Token Latency (ms)
Inter-Token Latency (ms)
0
10
20
30
40
50
60
0
50
100
150
200
33.3
30.6
38.4
36.8
27.9
28.0
56.5
56.9
40.2
36.0
196
188
140
141
40.8
40.8
CC
No CC
LLama 3.1 8B
Mistral 3.1 24B
GPT OSS 120B
LLama 3.3 70B Int4

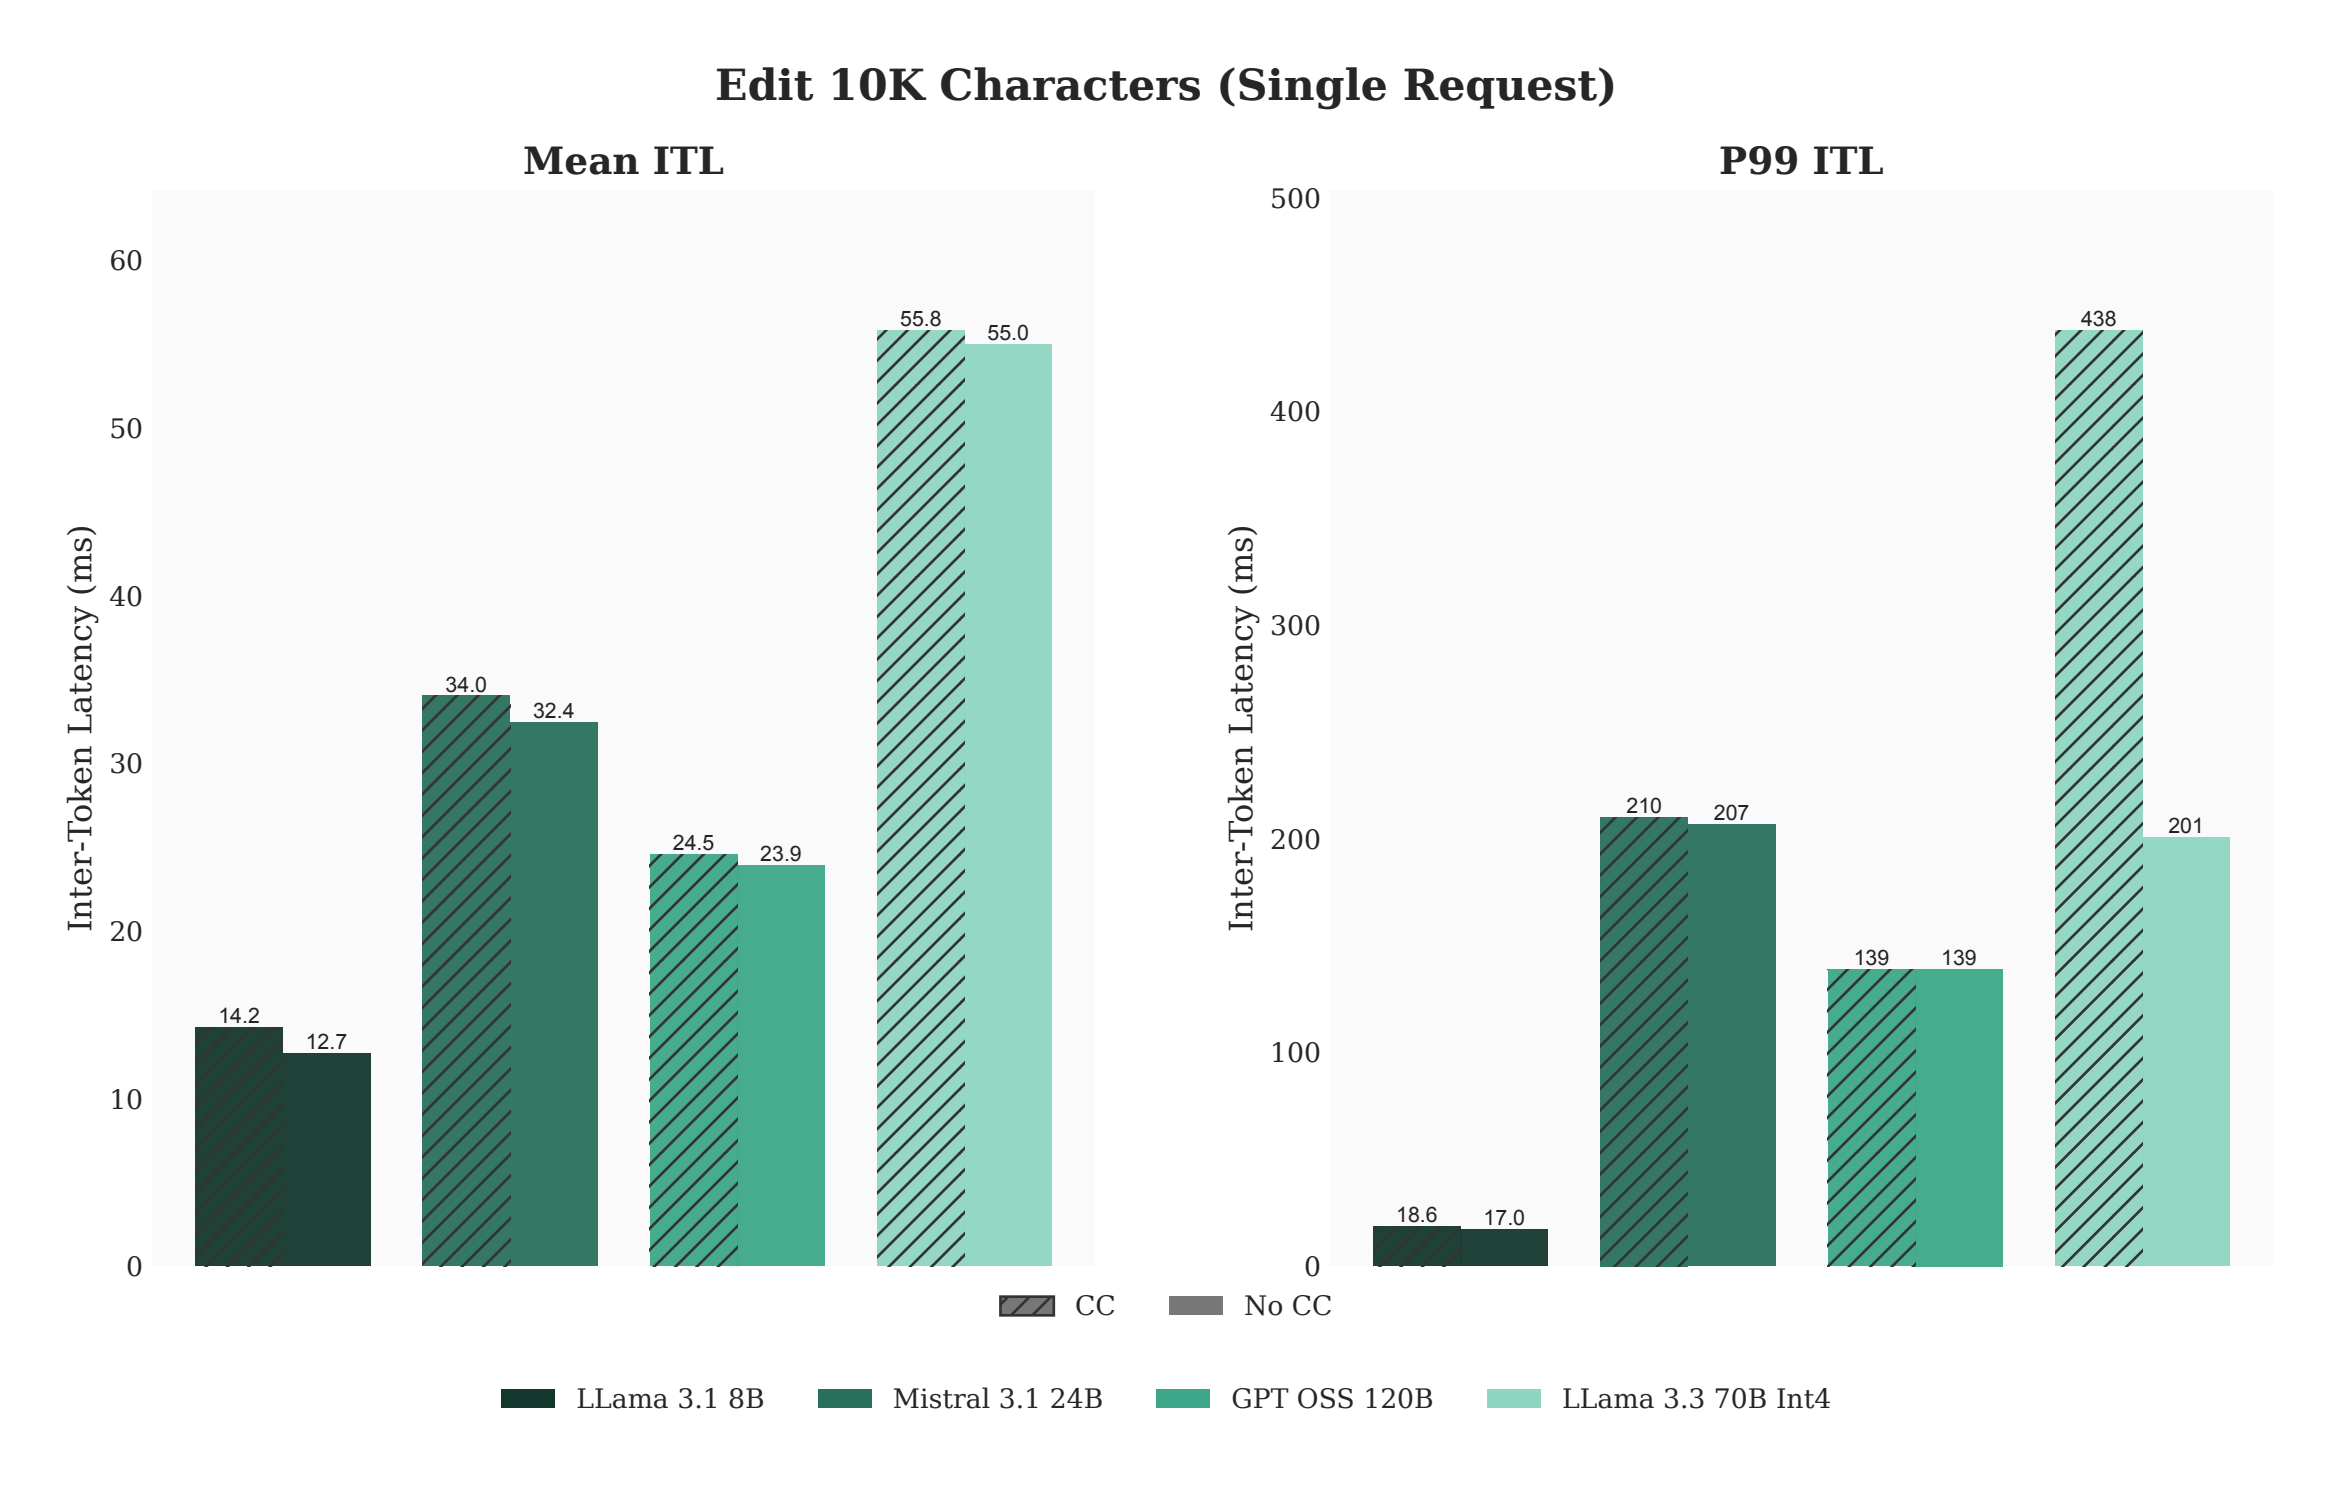
Edit 10K Characters (Single Request)
Mean ITL
P99 ITL
Inter-Token Latency (ms)
Inter-Token Latency (ms)
0
10
20
30
40
50
60
0
100
200
300
400
500
14.2
12.7
34.0
32.4
24.5
23.9
55.8
55.0
18.6
17.0
210
207
139
139
438
201
CC
No CC
LLama 3.1 8B
Mistral 3.1 24B
GPT OSS 120B
LLama 3.3 70B Int4

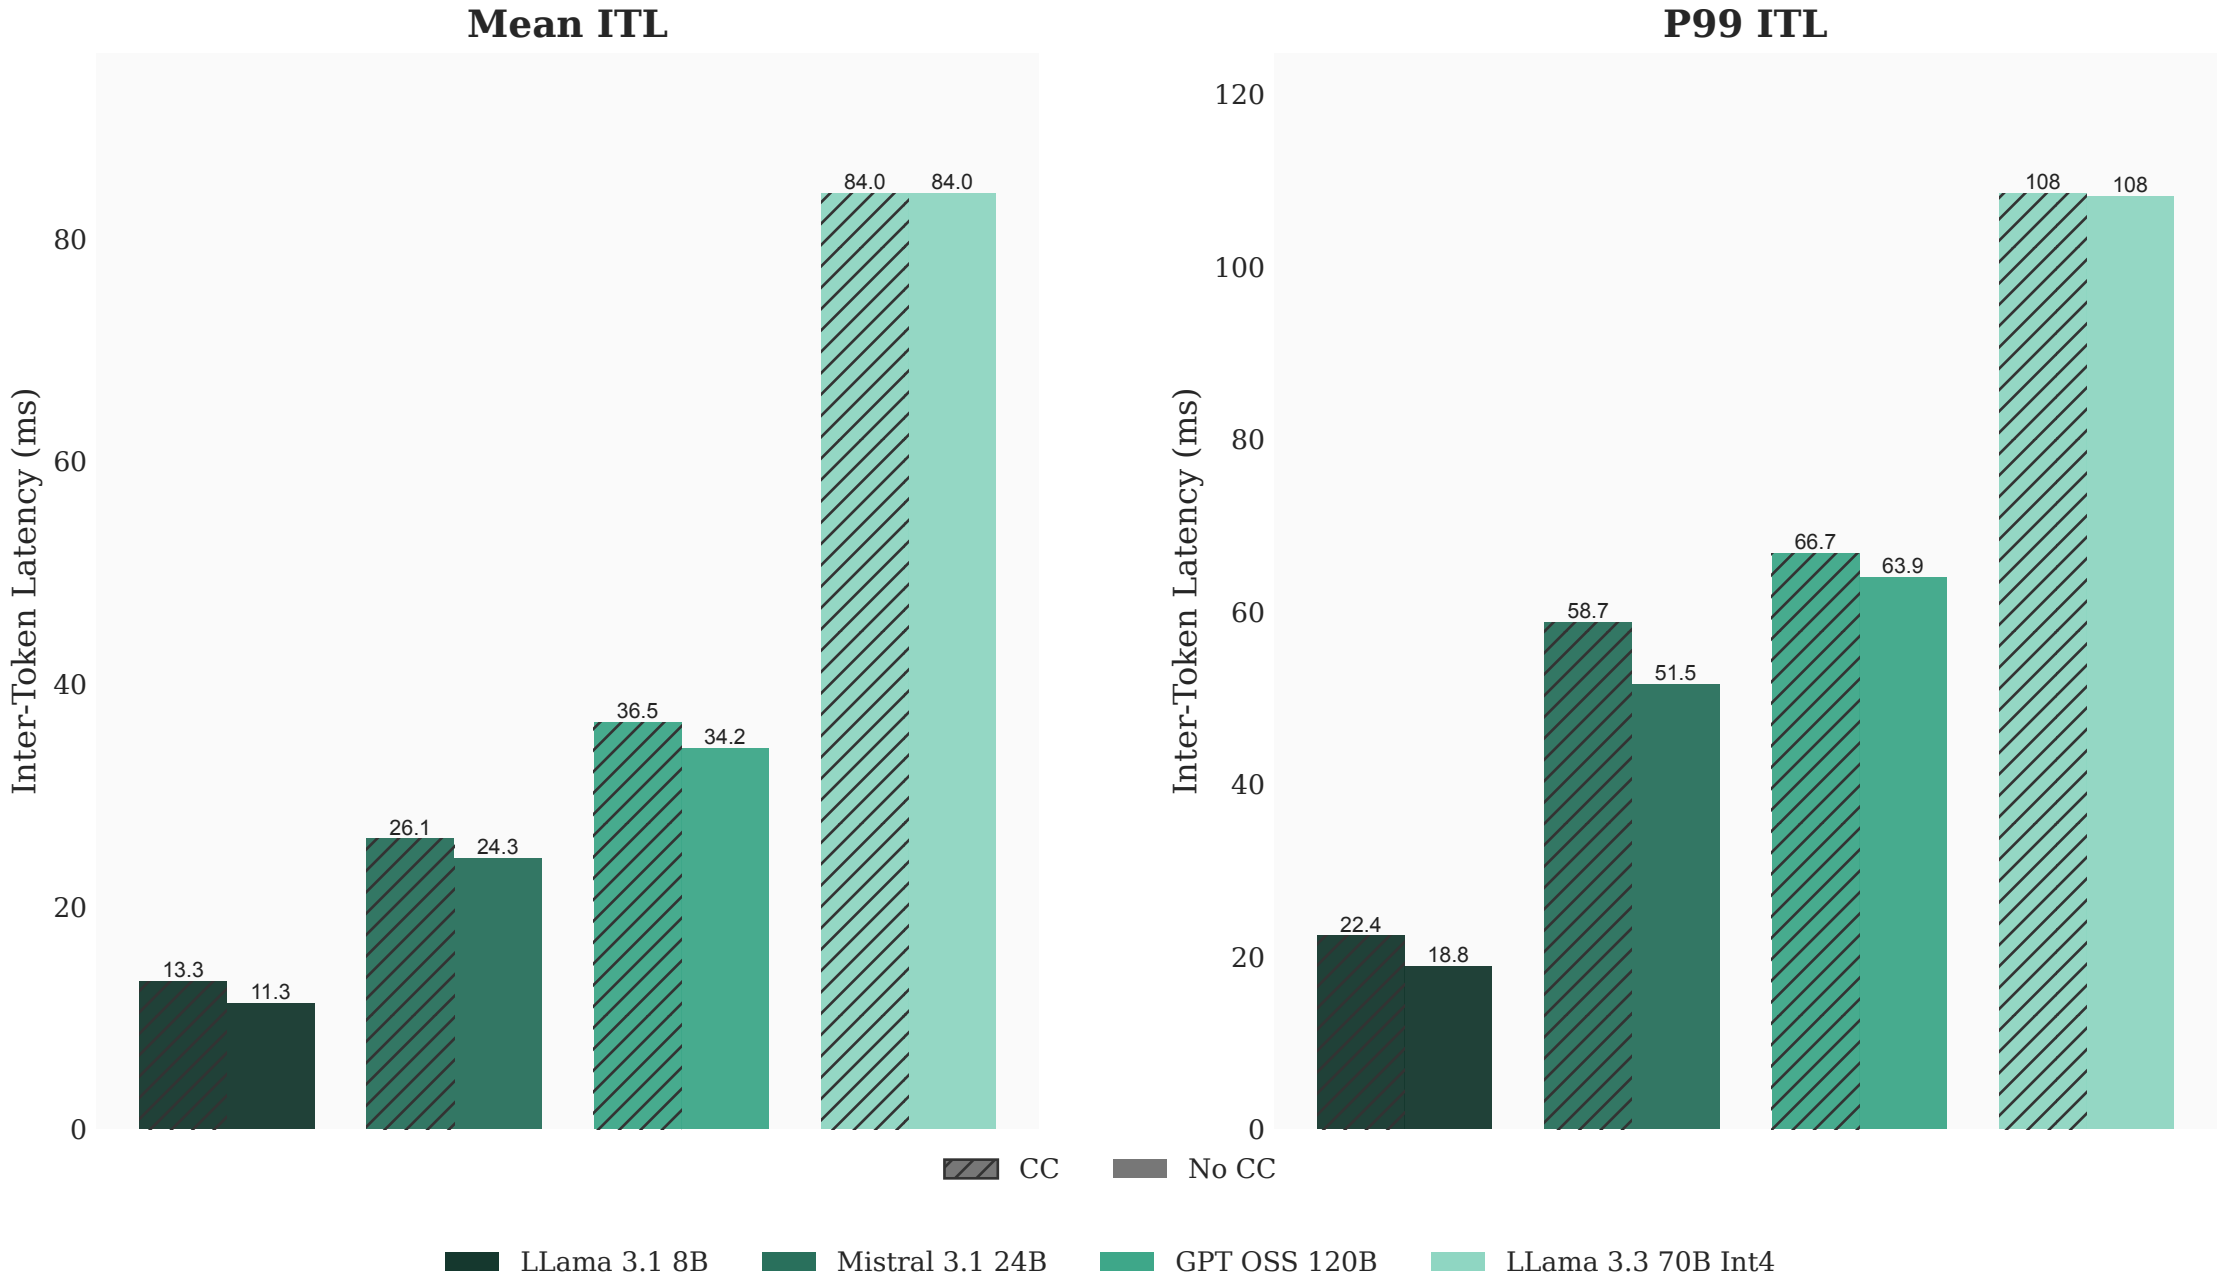
Numina Math (100 Request Rate)
Mean ITL
P99 ITL
Inter-Token Latency (ms)
0
20
40
60
80
100
120
13.3
11.3
26.1
24.3
36.5
34.2
84.0
84.0
22.4
18.8
58.7
51.5
66.7
63.9
108
108
CC
No CC
LLama 3.1 8B
Mistral 3.1 24B
GPT OSS 120B
LLama 3.3 70B Int4

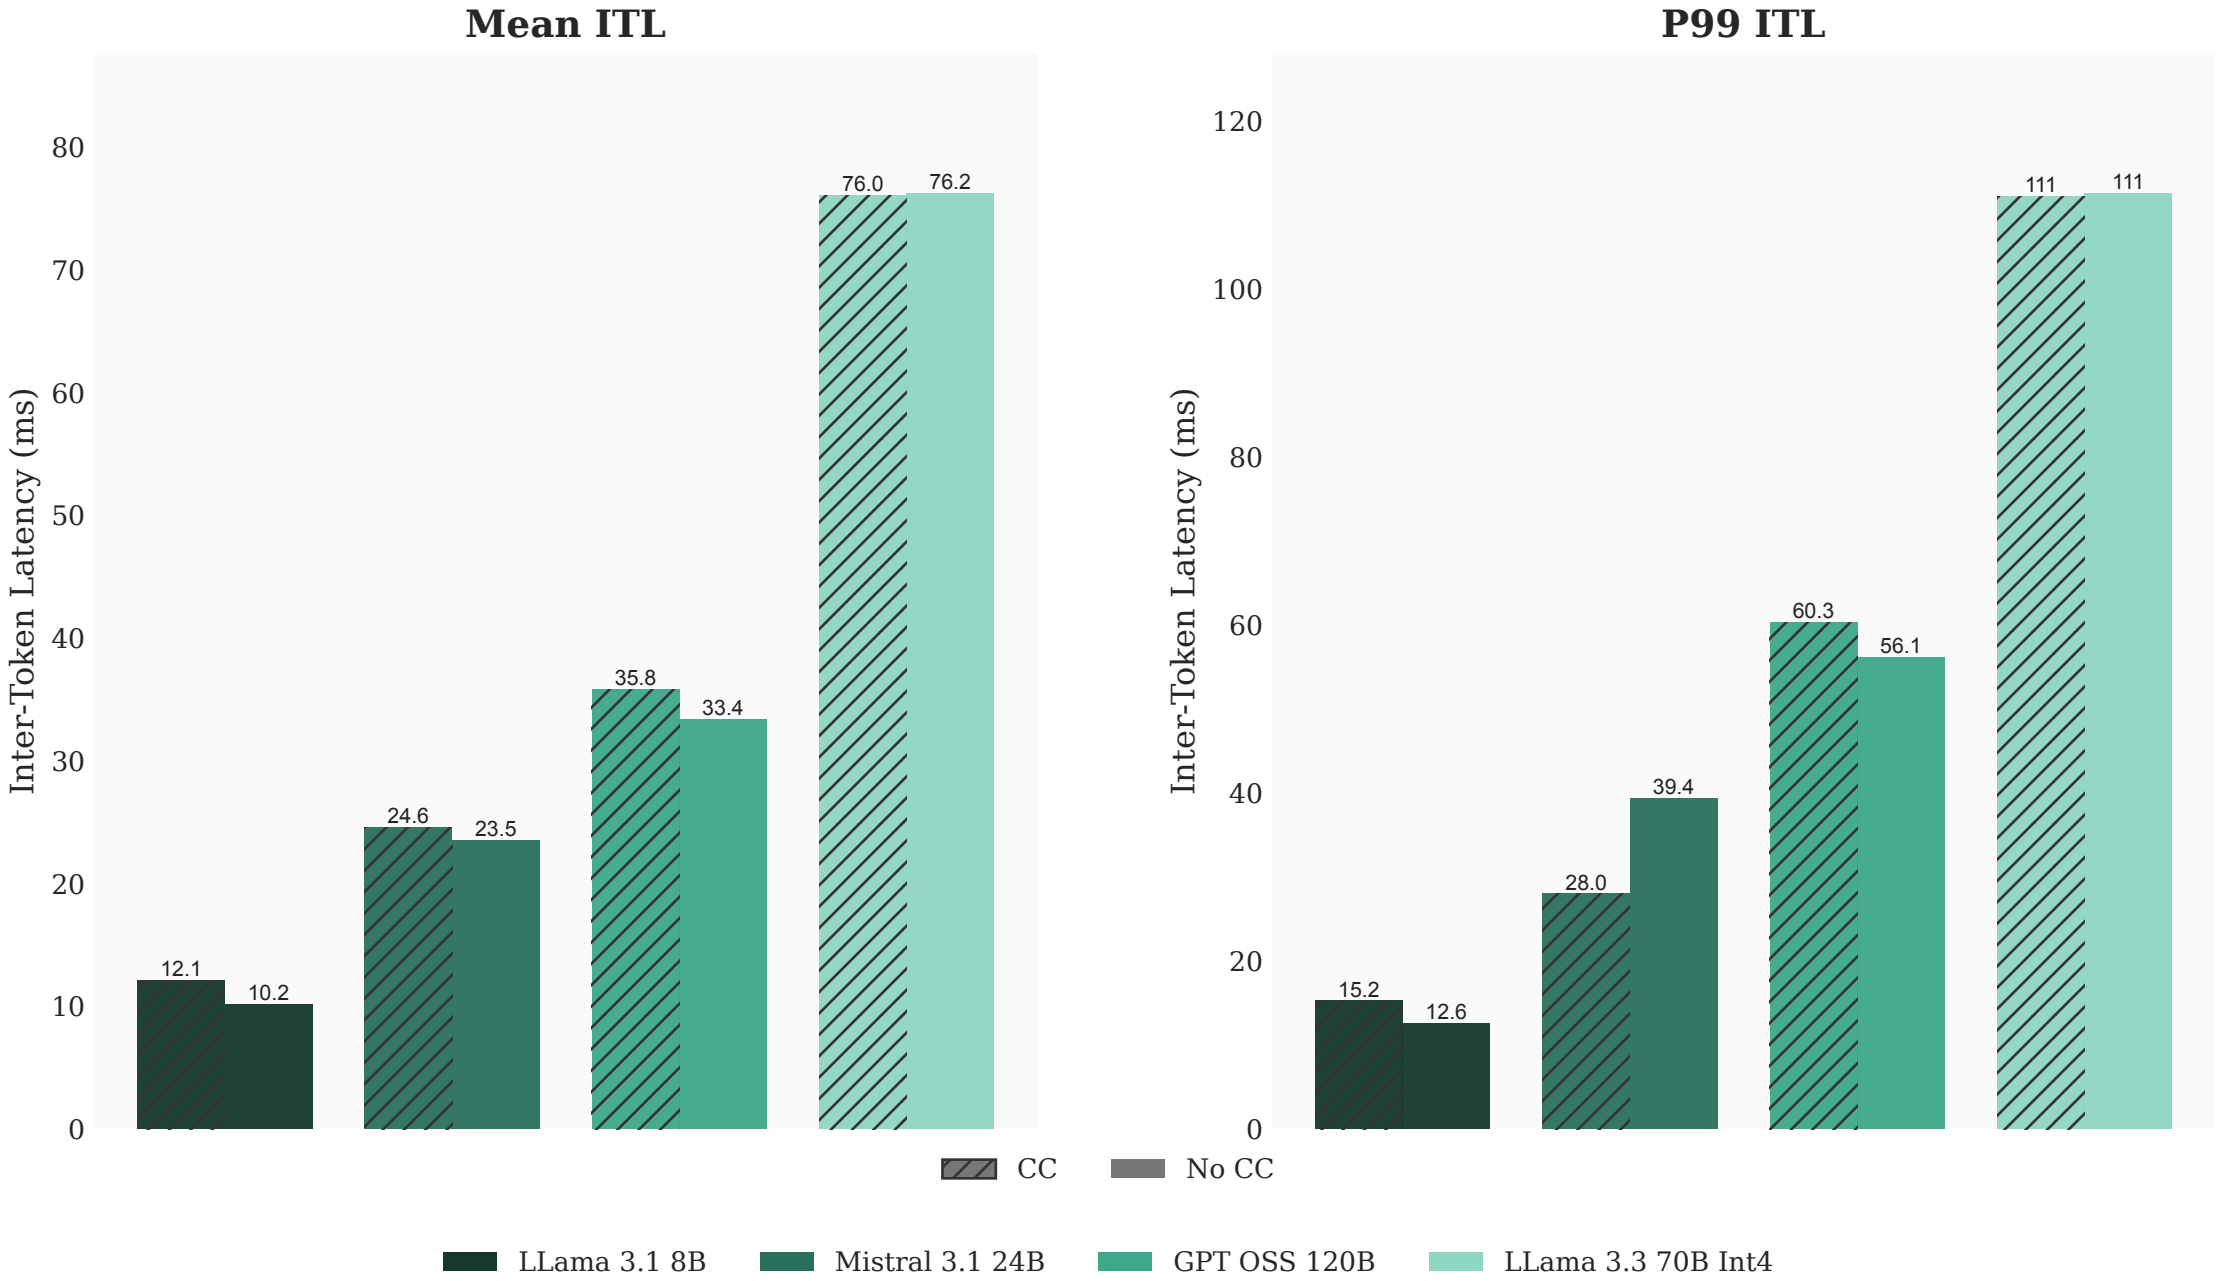
Numina Math (50 Request Rate)
Mean ITL
P99 ITL
Inter-Token Latency (ms)
0
10
20
30
40
50
60
70
80
20
40
60
80
100
120
12.1
10.2
24.6
23.5
35.8
33.4
76.0
76.2
15.2
12.6
28.0
39.4
60.3
56.1
111
111
CC
No CC
LLama 3.1 8B
Mistral 3.1 24B
GPT OSS 120B
LLama 3.3 70B Int4

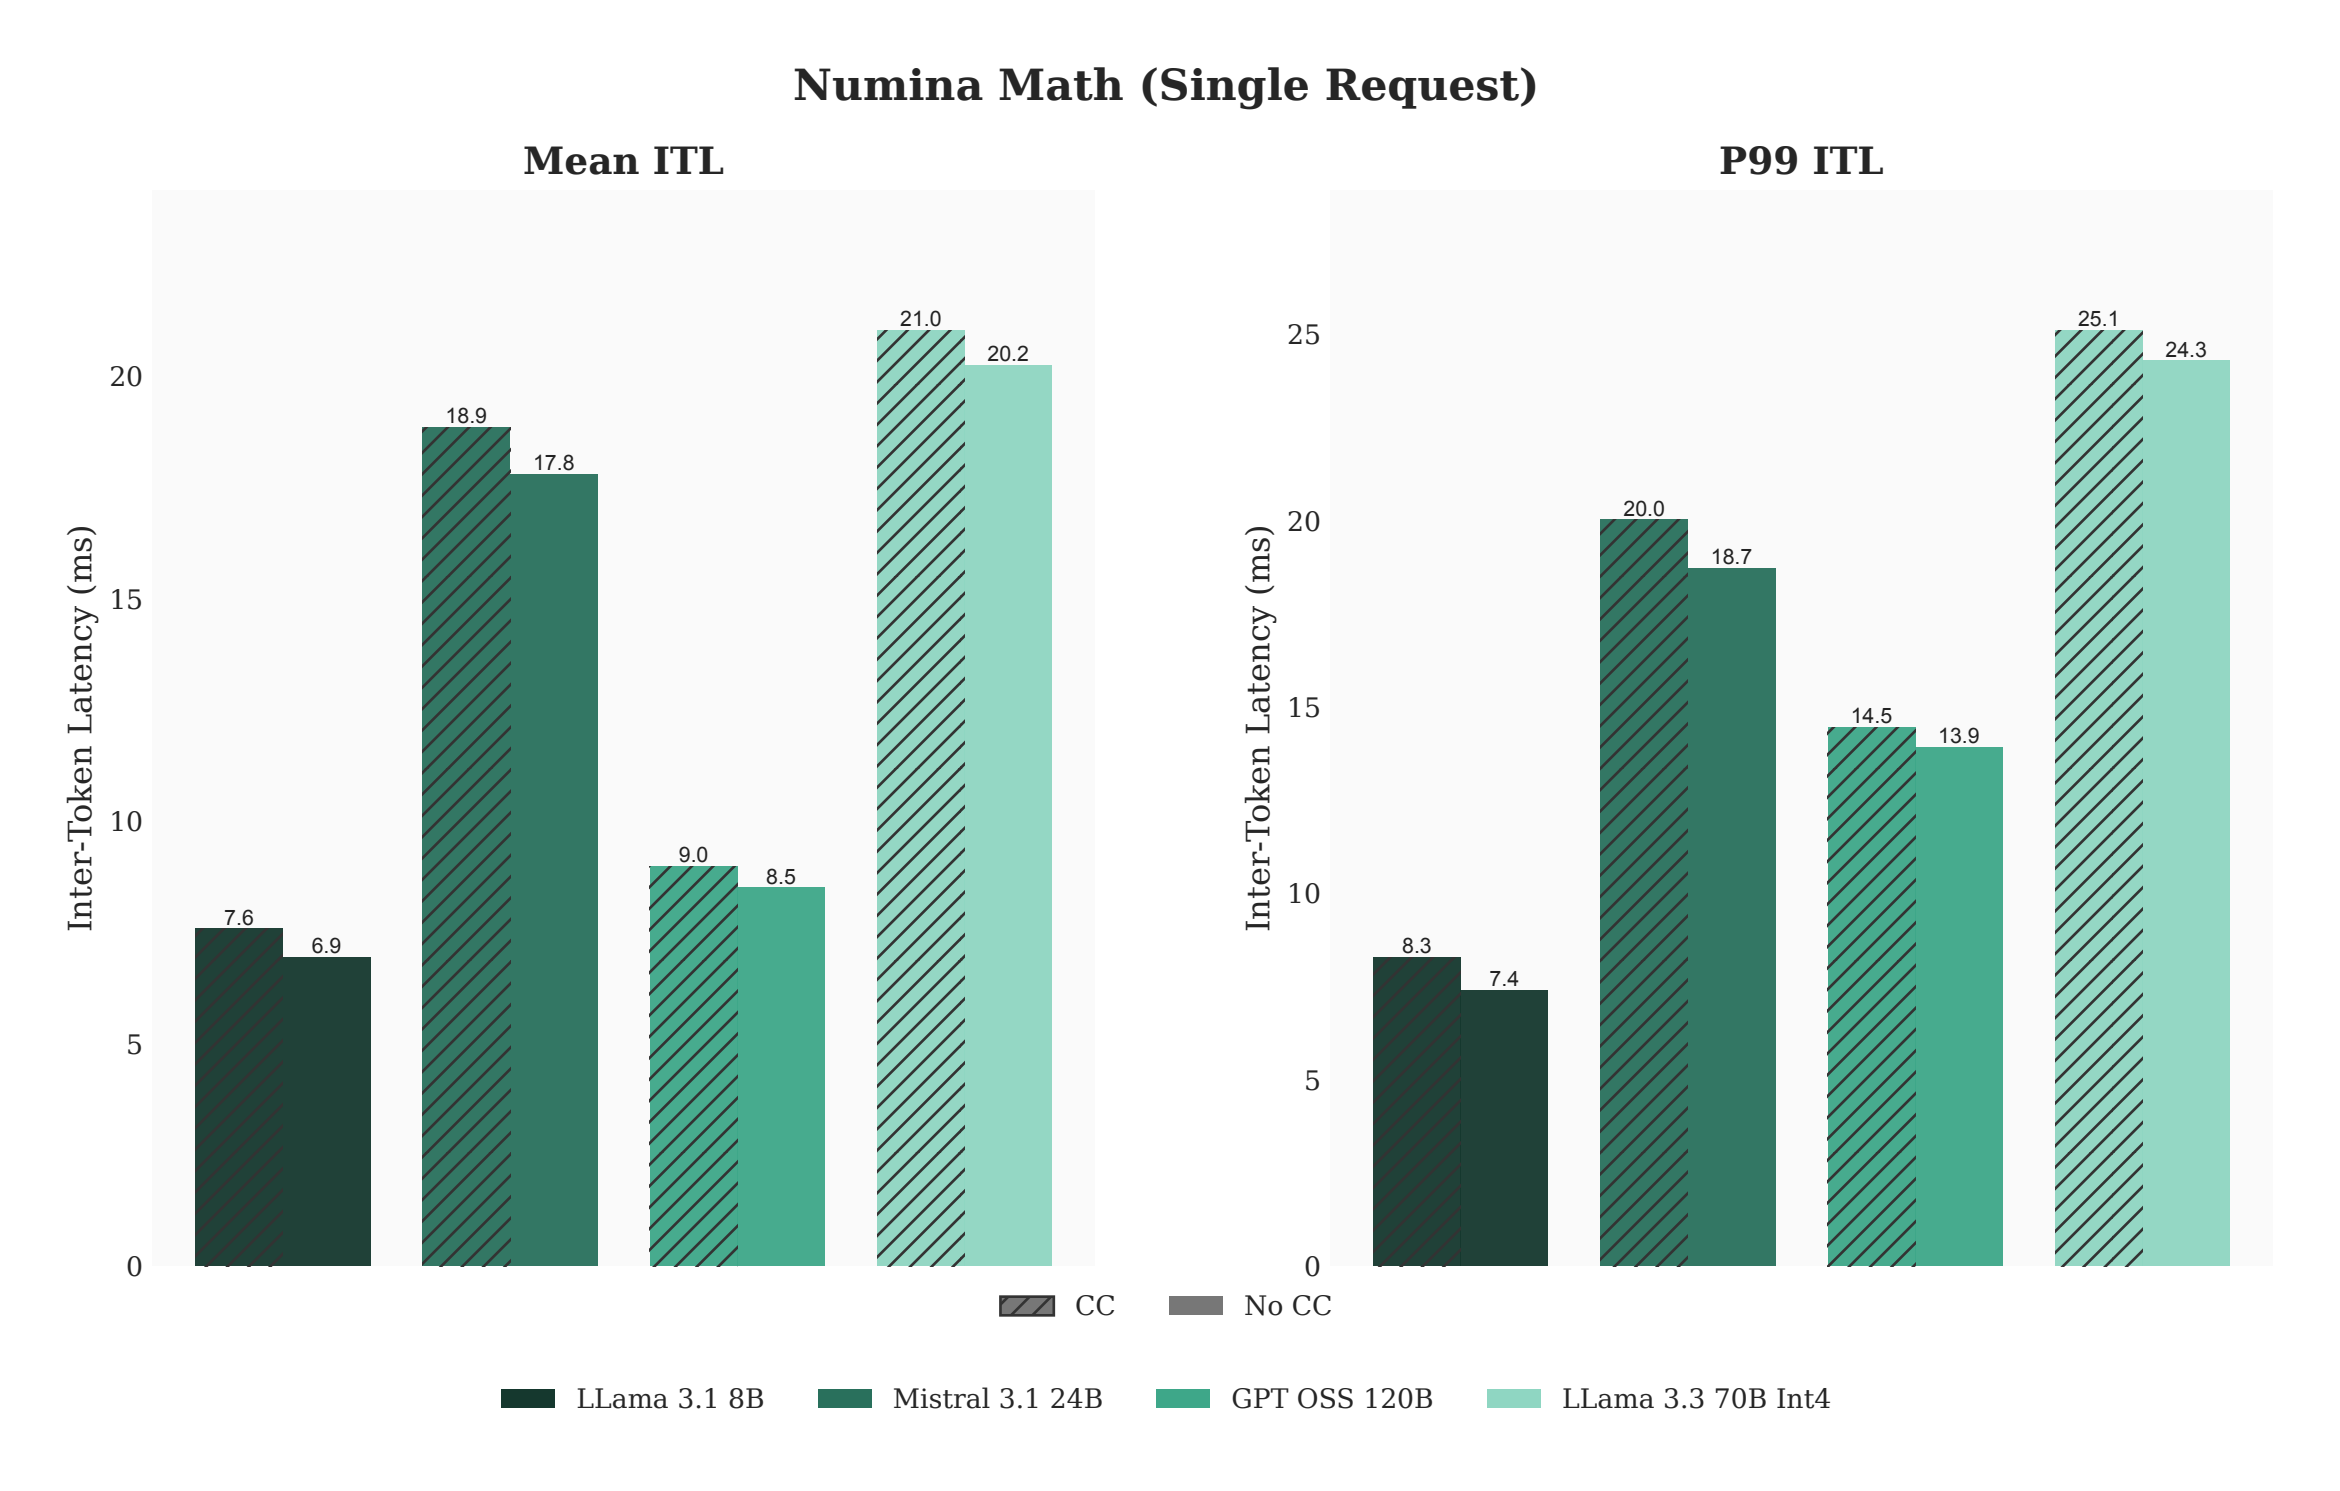
Numina Math (Single Request)
Mean ITL
P99 ITL
Inter-Token Latency (ms)
0
5
10
15
20
25
7.6
6.9
18.9
17.8
9.0
8.5
21.0
20.2
8.3
7.4
20.0
18.7
14.5
13.9
25.1
24.3
CC
No CC
LLama 3.1 8B
Mistral 3.1 24B
GPT OSS 120B
LLama 3.3 70B Int4